AI Regulations Worldwide

ENSURING A NATIONAL POLICY FRAMEWORK FOR ARTIFICIAL INTELLIGENCE

By the authority vested in me as President by the Constitution and the laws of the United States of America, it is hereby ordered:

Section 1. Purpose. United States leadership in Artificial Intelligence (AI) will promote United States national and economic security and dominance across many domains. Pursuant to Executive Order 14179 of January 23, 2025 (Removing Barriers to American Leadership in Artificial Intelligence), I revoked my predecessor's attempt to paralyze this industry and directed my Administration to remove barriers to United States AI leadership. My Administration has already done tremendous work to advance that objective, including by updating existing Federal regulatory frameworks to remove barriers to and encourage adoption of AI applications across sectors. These efforts have already delivered tremendous benefits to the American people and led to trillions of dollars of investments across the country. But we remain in the earliest days of this technological revolution and are in a race with adversaries for supremacy within it.

To win, United States AI companies must be free to innovate without cumbersome regulation. But excessive State regulation thwarts this imperative. First, State-by-State regulation by definition creates a patchwork of 50 different regulatory regimes that makes compliance more challenging, particularly for start-ups. Second, State laws are increasingly responsible for requiring entities to embed ideological bias within models. For example, a new Colorado law banning "algorithmic discrimination" may even force AI models to produce false results in order to avoid a "differential treatment or impact" on protected groups. Third, State laws sometimes impermissibly regulate beyond State borders, impinging on interstate commerce.

My Administration must act with the Congress to ensure that there is a minimally burdensome national standard — not 50 discordant State ones. The resulting framework must forbid State laws that conflict with the policy set forth in this order. That framework should also ensure that children are protected, censorship is prevented, copyrights are respected, and communities are safeguarded. A carefully crafted national framework can ensure that the United States wins the AI race, as we must.

Until such a national standard exists, however, it is imperative that my Administration takes action to check the most onerous and excessive laws emerging from the States that threaten to stymie innovation.

Sec. 2. Policy. It is the policy of the United States to sustain and enhance the United States' global AI dominance through a minimally burdensome national policy framework for AI.

Sec. 3. AI Litigation Task Force. Within 30 days of the date of this order, the Attorney General shall establish an AI Litigation Task Force (Task Force) whose sole responsibility shall be to challenge State AI laws inconsistent with the policy set forth in section 2 of this order, including on grounds that such laws unconstitutionally regulate interstate commerce, are preempted by existing Federal regulations, or are otherwise unlawful in the Attorney General's judgment, including, if appropriate, those laws identified pursuant to section 4 of this order. The Task Force shall consult from time to time with the Special Advisor for AI and Crypto, the Assistant to the President for Science and Technology, the Assistant to the President for Economic Policy, and the Assistant to the President

and Counsel to the President regarding the emergence of specific State AI laws that warrant challenge.

Sec. 4. Evaluation of State AI Laws. Within 90 days of the date of this order, the Secretary of Commerce, consistent with the Secretary's authorities under 47 U.S.C. 902(b), shall, in consultation with the Special Advisor for AI and Crypto, the Assistant to the President for Economic Policy, the Assistant to the President for Science and Technology, and the Assistant to the President and Counsel to the President, publish an evaluation of existing State AI laws that identifies onerous laws that conflict with the policy set forth in section 2 of this order, as well as laws that should be referred to the Task Force established pursuant to section 3 of this order. That evaluation of State AI laws shall, at a minimum, identify laws that require AI models to alter their truthful outputs, or that may compel AI developers or deployers to disclose or report information in a manner that would violate the First Amendment or any other provision of the Constitution. The evaluation may additionally identify State laws that promote AI innovation consistent with the policy set forth in section 2 of this order.

Sec. 5. Restrictions on State Funding. (a) Within 90 days of the date of this order, the Secretary of Commerce, through the Assistant Secretary of Commerce for Communications and Information, shall issue a Policy Notice specifying the conditions under which States may be eligible for remaining funding under the Broadband Equity Access and Deployment (BEAD) Program that was saved through my Administration's "Benefit of the Bargain" reforms, consistent with 47 U.S.C. 1702(e)-(f). That Policy Notice must provide that States with onerous AI laws identified pursuant to section 4 of this order are ineligible for non-deployment funds, to the maximum extent allowed by Federal law. The Policy Notice must also describe how a fragmented State regulatory landscape for AI threatens to undermine BEAD-funded deployments, the growth of AI applications reliant on high-speed networks, and BEAD's mission of delivering universal, high-speed connectivity.

(b) Executive departments and agencies (agencies) shall assess their discretionary grant programs in consultation with the Special Advisor for AI and Crypto and determine whether agencies may condition such grants on States either not enacting an AI law that conflicts with the policy of this order, including any AI law identified pursuant to section 4 or challenged pursuant to section 3 of this order, or, for those States that have enacted such laws, on those States entering into a binding agreement with the relevant agency not to enforce any such laws during the performance period in which it receives the discretionary funding.

Sec. 6. Federal Reporting and Disclosure Standard. Within 90 days of the publication of the identification specified in section 4 of this order, the Chairman of the Federal Communications Commission shall, in consultation with the Special Advisor for AI and Crypto, initiate a proceeding to determine whether to adopt a Federal reporting and disclosure standard for AI models that preempts conflicting State laws.

Sec. 7. Preemption of State Laws Mandating Deceptive Conduct in AI Models. Within 90 days of the date of this order, the Chairman of the Federal Trade Commission shall, in consultation with the Special Advisor for AI and Crypto, issue a policy statement on the application of the Federal Trade Commission Act's prohibition on unfair and deceptive acts or practices under 15 U.S.C. 45 to AI models. That policy statement must explain the circumstances under which State laws that require alterations to the truthful outputs of AI models are preempted by the Federal Trade Commission Act's prohibition on engaging in deceptive acts or practices affecting commerce.

Sec. 8. Legislation. (a) The Special Advisor for AI and Crypto and the Assistant to the President for Science and Technology shall jointly prepare a legislative recommendation establishing a uniform Federal policy framework for AI that preempts State AI laws that conflict with the policy set forth in this order.

(b) The legislative recommendation called for in subsection (a) of this section shall not propose preempting otherwise lawful State AI laws relating to:

(ii) the functions of the Director of the Office of Management and Budget relating to budgetary, administrative, or legislative proposals.

(c) This order is not intended to, and does not, create any right or benefit, substantive or procedural, enforceable at law or in equity by any party against the United States, its departments, agencies, or entities, its officers, employees, or agents, or any other person.

The Updated State of AI Regulations for 2025

In January of last year, we discussed some of the current AI regulations that exist, as well as resources to stay current as the AI regulation landscape changes.

Of course, as Ferris Bueller would say: “Life moves pretty fast. If you don’t stop and look around once in a while, you could miss it.” That statement could certainly be applied to AI regulations, as a lot of changes have happened over the past 15 months. With that in mind, we have developed an update to our article on the landscape of global AI legislation that was published last year in order to reflect changes that have happened since our original post.

As we did last year, we have organized the update by jurisdiction to illustrate the current global landscape, as follows:

The U.S. still lacks a comprehensive federal AI law, but 2024 saw a surge of proposals and state-level actions. Congress introduced numerous AI bills focused on accountability and transparency, though none have become law yet■. However, at the state level, at least 45 states proposed AI-related bills in 2024, and 31 states and territories enacted AI laws or resolutions■. Examples include:

Colorado passed the first broad AI law requiring developers of “high-risk” AI to use reasonable care to prevent algorithmic bias and to disclose AI use to consumers.

Tennessee passed the Ensuring Likeness, Voice, and Image Security Act (ELVIS) Act barring unauthorized AI simulations of a person’s likeness or voice.

California enacted a package of AI laws in September 2024 on topics from deepfakes to transparency■. These include the Defending Democracy from Deepfake Deception Act (AB 2655) mandating large online platforms to detect and label materially deceptive AI-generated election content■, and the AI Transparency Act (SB 942, effective Jan. 2026) requiring AI services with 1M+ users to disclose AI-generated content and implement content detection measures.

Similar to data privacy laws, the patchwork of state AI laws is rapidly expanding in the absence of overarching federal regulation.

On the federal level, executive actions have shaped AI policy. In October 2023, President Biden issued Executive Order 14110 on “Safe, Secure, and Trustworthy AI”, directing a broad range of AI oversight measures (from safety standards to equity and civil rights protections). However, in January 2025 the new administration took a different course: President Trump rescinded Biden’s

order and replaced it with a new order, “Removing Barriers to American Leadership in AI,” which shifts toward deregulation and explicitly prioritizes AI innovation and U.S. competitiveness, instructing agencies to eliminate policies that might “hinder American AI dominance”.

As we discussed last year, the European Union achieved a landmark milestone by enacting the world’s first comprehensive AI law. The EU AI Act was formally adopted in mid-2024 after intensive inter-institutional negotiations. On 12 July 2024, the final Artificial Intelligence Act (Regulation (EU) 2024/1689) was published in the EU’s Official Journal■. This regulation, effective August 1, 2024, establishes harmonized rules for AI across all 27 EU states. Most provisions, however, will not be enforced until August 2, 2026, allowing a two-year phase-in period for compliance.

The EU AI Act takes a risk-based approach: it prohibits a few unacceptable AI practices outright and imposes escalating obligations on other AI systems based on their risk level. For example, AI systems deemed “high-risk” (such as AI in medical devices, recruiting, or credit scoring) must meet strict requirements for risk management, data governance, transparency, human oversight, and conformity assessments before being deployed. Providers of high-risk AI will have to register their systems in an EU database and obtain CE certification. The Act also bans uses of AI that are seen as contrary to EU values – notably real-time biometric identification in public spaces for law enforcement, “social scoring” of citizens by governments, and AI that exploits vulnerable groups in harmful ways (with limited exceptions).

Outside the EU, the United Kingdom has adopted a different approach, preferring guidance over enacting new laws so far. In March 2023, the U.K. government published a white paper outlining a “pro-innovation” AI regulatory framework. Instead of a single AI statute, the U.K. is empowering sectoral regulators (like the health, financial, or transportation regulators) to issue AI rules tailored to their domains■. The strategy centers on five principles (safety, transparency, fairness, accountability, and contestability), which regulators are expected to enforce using existing powers. Companies in the U.K. are advised to follow the white paper principles and monitor regulatory guidance. The landscape may shift in 2025 if the U.K. moves forward with an AI Act of its own, potentially influenced by international developments like the EU AI Act and US policies.

Canada’s effort to pass a federal AI law, the Artificial Intelligence and Data Act (AIDA), has seen twists and delays since 2024. AIDA was proposed as part of Bill C-27 (an omnibus bill that also updates privacy law)■. The act would establish rules for “high-impact” AI systems – requiring impact assessments, mitigation of biases, and registration of such systems – and prohibit reckless AI deployments that could cause serious harm. However, as of early 2025, AIDA has not yet been enacted.

Throughout 2024, Bill C-27 progressed through Parliament but faced extensive review and debate■. In January 2025, the Canadian Parliament was prorogued, effectively killing Bill C-27 on the order paper. This means AIDA in its current form has “died” in committee and would need to be reintroduced in a new parliamentary session.

China has been rapidly expanding its AI regulatory framework, focusing on guiding the development of AI in a safe and government-supervised manner. In mid-2023, China implemented pioneering rules on generative AI services, and it has built on those in 2024–2025. The Interim Measures for Generative AI Services (effective August 15, 2023) require providers of generative AI that is accessible to the public to ensure content is lawful, truthful, and labeled if AI-generated, and to register their algorithms with regulators.

Since January 2024, Chinese authorities have released further refinements: for example, in May 2024 the government’s standards body (NISSTC) issued draft Security Requirements for Generative AI, detailing technical measures to secure training data and models. Also, one notable new regulation is China’s mandatory labeling rule for AI-generated content. In March 2025, the

Cyberspace Administration of China (CAC) issued the final “Measures for Labeling AI-Generated Content”, which take effect on September 1, 2025. These rules compel all online services that create or distribute AI-generated content to clearly label such content.

Beyond content rules, China continues to refine its AI governance frameworks at a high level. In September 2024, China’s national committee on AI governance released an AI Safety Governance Framework aligned with Beijing’s Global AI Governance Initiative■. This framework lays out broad principles – a “people-centered approach” and AI “for good” – and classifies AI risks to guide policymakers, emphasizing ethics, transparency, and continuous risk monitoring, reflecting concerns about bias and security in AI.

In December 2024, Brazil’s Senate approved a comprehensive AI Bill (No. 2338/2023), which adopts an EU-like risk-based model. The bill defines categories of AI systems and corresponding obligations, and it now awaits further approval. If enacted, Brazil’s law would be one of the first in the region to specifically regulate AI development and use.

Many other countries around the world have published or proposed national guidance, but (to our knowledge) none of them have binding laws – yet.

In addition to the two key resources we discussed last year – OECD AI Policy Observatory (OCED.AI) and Global AI Legislation Tracker – here are two additional resources for keeping current on AI policies and regulations in the EU and US:

EU AI Act Resources : In addition to the European Commission’s EU AI Act page and the official Eur-Lex site for Regulation (EU) 2024/1689 contain the text of the law and updates on its implementation timeline, eu is an unofficial tracker that posts the Act’s legislative history, key documents, and timelines■.

: In addition to the European Commission’s EU AI Act page and the official Eur-Lex site for Regulation (EU) 2024/1689 contain the text of the law and updates on its implementation timeline, is an unofficial tracker that posts the Act’s legislative history, key documents, and timelines■. NCSL Artificial Intelligence Legislation Tracker: The US National Conference of State Legislatures maintains an updated tracker summarizing AI-related bills in all U.S. states, including enacted laws (like those in Colorado, California, etc.) and pending bills.

There have been many changes to the AI regulation landscape in the last 15 months, and the expectation is that this space will continue to evolve for the foreseeable future. Cimplifi will continue to revisit the AI regulation landscape periodically to ensure you’re up on the latest global AI guidance and regulations that may impact you.

AI Regulations around the World - 2026

Around the world, at least 72 countries have proposed over 1000 AI-related policy initiatives and legal frameworks to address public concerns around AI safety and governance. Here’s what those AI regulations look like at the beginning of 2026...

(This article will continue to evolve and change as new details and milestones around AI Regulations emerge. A history of these edits can be found at the end of the piece.)

In the last decade, we have witnessed the most significant progress in the history of machine learning. The last five years, in particular, have seen AI transition from being positioned merely as a research or innovation project to a game-changing technology with countless real-world applications. This rapid shift has moved topics such as model governance, safety, and risk from the periphery of academic discussions to become central concerns in the field. Governments and industries worldwide have responded with proposals for regulations, and although the current regulatory landscape is nascent, diverse, and subject to impending change, the direction of travel is clear.

In March of 2023, the Government at the time published its framework for “A pro-innovation approach to AI regulation”. Rather than having concrete overarching regulation governing all AI development, the paper proposed a more activity-based approach that left individual regulatory bodies responsible for AI governance in their respective domains, with “central AI regulatory functions” to help regulators achieve this.

The UK continued to articulate its approach to AI in November 2023 when it hosted the International AI Safety Summit, the first international government-led conference on this subject. The global interest in the event and the attendance by representatives from around the world demonstrated how far the conversation around AI governance had come in recent years. The summit moved the topic of AI Safety from the fringe to centre stage, paving the way for regulations that would be the main story for the next year. 2023 also marked the foundation of the AI Safety Institute (AISI), which evolved into the AI Security Institute in February 2025.

If 2023 was about talking the talk on AI Safety, 2024 was about walking the walk. In February, the Government published its response to public consultation on its original white paper, pledging over £100 million to support regulators and advance research and innovation. Whilst promising to provide regulators with the “skills and tools they need to address the risks and opportunities of AI”, they also announced a deadline of April 30 for these regulators to publish their own approaches to AI management, which many of them did, including the Financial Conduct Authority (FCA), the Information Commissioner's Office (ICO), and the Equality and Human Rights Commission (EHRC).

The most notable impact of these new approaches came on July 31 when the FCA, a financial regulating authority in the UK, set out its Consumer Duty to ensure that all organisations within its remit offer services that represent fair value for consumers. These rules placed increased importance on effective AI governance, especially in competitive, customer-centric markets like insurance pricing.

That same July, the UK’s political landscape also underwent a significant shift when the Labour Party, led by Keir Starmer, won a landslide victory in the general election, unseating the incumbent Conservative government. This change introduced uncertainty regarding the UK’s approach to artificial intelligence (AI) regulation, as the two parties had differing strategies.

Prior to the election, the Conservative government advocated for a ‘pro-innovation’ and non-binding framework for AI regulation. They emphasised flexibility, allowing existing sectoral regulators to apply cross-sectoral principles to AI within their domains without immediate legislative intervention. The focus was on maintaining agility to foster innovation while monitoring the need for future targeted regulations, particularly concerning advanced general-purpose AI systems. Upon taking office, the Labour government signalled a shift towards more structured AI regulation. Their manifesto outlined plans to:

Establish a Regulatory Innovation Office: This office aims to assist existing regulators in updating AI-related regulations, streamlining approval processes, and managing cross-sectoral issues.

Implement Binding Regulations: Labour proposed introducing binding regulations specifically targeting companies developing the most powerful AI models, though details on which companies would be affected were not specified.

Labour proposed introducing binding regulations specifically targeting companies developing the most powerful AI models, though details on which companies would be affected were not specified. Ban Harmful AI Applications: The manifesto included a commitment to prohibit the creation of sexually explicit deepfakes.

The Department of Science, Innovation, and Technology has since published the AI Opportunities Action Plan in January 2025. The plan features “Recommendations for the government to capture the opportunities of AI to enhance growth and productivity and create tangible benefits for UK citizens.” The stated goal of these 50 recommendations is to:

Ministers have since signalled a Frontier AI Bill to give the AI Security Institute statutory powers, including potential pre■deployment model testing authority for the most capable systems. The AISI continues government-led evaluations; its first Frontier AI Trends Report was published on December 18, 2025.

The Labour Party’s initial commitment to a more interventionist approach in regulating artificial intelligence (AI) has already faced significant challenges due to recent international developments, particularly the election of Donald Trump as US President and the outcomes of the AI Action Summit in Paris in February 2025. This event, the 3rd edition of the AI Safety Summits that began at Bletchley Park in Nov 2023, shifted the focus from theoretical safety concerns to implementation, hence the replacement of the word “Safety” with “Action” in its title.

At the Summit, the U.S. and the UK declined to sign a declaration promoting “inclusive and sustainable” AI, which was endorsed by 60 other countries and emphasised ethical and responsible AI development. The UK’s refusal was attributed to concerns over national security and a perceived lack of clarity in global governance frameworks. This decision has drawn criticism from various groups, who argue that it could undermine the UK’s position as a leader in ethical AI innovation.

As of the start of 2026, a Private Member’s Artificial Intelligence (Regulation) Bill was also reintroduced, having failed to progress into law in 2024, and is progressing in the House of Lords, representing a renewed attempt to introduce legislation that is specific to AI.

On 21 April 2021, Europe’s landmark AI Act was officially proposed and seemed destined to be a first-in-the-world attempt to enshrine a set of AI regulations into law. The underlying premise of the bill was simple: regulate AI based on its capacity to cause harm. Drafters of the bill outlined various AI use cases and applications and then classified them with an appropriate degree of AI risk from minimal to high. Limited-risk AI systems would need to comply with minimal requirements for transparency. In contrast, high-risk AI systems, such as those related to aviation, education, and biometric surveillance, would have higher requirements and need to be registered in an EU database.

Additionally, some AI systems were deemed to have an unacceptable level of risk and would be banned outright, save for a few exceptions for law enforcement. These included:

Just as the Act was nearing its final stages at the end of 2022, ChatGPT took the world by storm, attracting 100 million users in its first two months and becoming one of the fastest-growing software applications ever. The new wave of Generative AI tools and general-purpose AI systems (GPAI) did not fit neatly into the EU’s original risk categories, and revisions to the act were required.

Nearly a year later, and after extensive lobbying attempts by Silicon Valley tech giants to water down the language on foundation models, a provisional agreement was reached. Despite this, there were further concerns that the bill would fail at the 11th hour due to reservations from certain countries around data protection and that the Act's regulations governing advanced AI systems would hamstring young national startups like France's Mistral AI and Germany's Aleph Alpha. But with the Parliament issuing assurances that these fears would be allayed with formal declarations in the future, and with the creation of the EU's Artificial Intelligence Office to enforce the AI Act, a deal was reached unanimously in February 2024 to proceed with the legislation, and it received final approval from the 27-nation bloc in a plenary vote on March 13, 2024.

The EU AI Act officially became law on the 1st of August 2024, with implementation staggered from early 2025 onwards. 'Unacceptable■risk' bans and AI literacy requirements have applied since Feb 2, 2025, while general■purpose AI (GPAI) rules, governance, penalties, and notified bodies began applying Aug 2, 2025. High■risk obligations broadly apply from Aug 2, 2026, with legacy GPAI models required to comply by Aug 2, 2027. In Ireland, where companies such as Meta, TikTok, and Google have their EU headquarters, the Government designated national competent authorities via S.I. No. 366/2025 and set up its national coordination mechanisms in Sept 2025.

For all EU member states, within 3 years of the AI Act coming into force, sometime around the middle of 2027, all other rules will be in place, including requirements for AI that gets classified as high-risk, which will require third-party conformity assessment under other EU rules. Non-compliance with the rules will lead to fines of up to €35 million or 7% of global turnover, depending on the infringement and the company's size.

As of this blog's most recent update in January 2026, despite initial uncertainty around the AI Act's ban on AI systems posing "unacceptable risk" and which systems actually fall under this umbrella, the EU has since provided greater clarity in the form of the GPAI Code of Practice and guidance timelines, with enforcement continuing to ramp up in the coming years. Nevertheless, there continues to be rising concern from business leaders around how the Act could stifle innovation, discourage investment, and cause the EU to fall behind in the global AI arms race. This lack of clear direction has led many large international companies to adopt a more reticent approach to deploying AI more widely throughout the European Union.

Like the UK, the US does not yet have a comprehensive AI regulation, but numerous frameworks and guidelines exist at both the federal and state levels. However, a dramatic shift in federal AI policy occurred in January 2025 with the transition from the Biden administration to the Trump administration.

In October 2023, President Joe Biden signed Executive Order 14110 on the Safe, Secure, and Trustworthy Development and Use of Artificial Intelligence. This order sought to impose reporting and safety obligations on AI companies, requiring them to test and disclose their models' capabilities and risks. However, upon assuming office in January 2025, President Donald Trump revoked this executive order within days, signalling a shift toward AI deregulation and industry-led innovation.

On January 23, 2025, President Trump signed Executive Order 14179, "Removing Barriers to American Leadership in Artificial Intelligence." The Order redirected policy and initiated a cross■government AI Action Plan whilst revoking or directing reviews of actions under Biden's EO 14110. Trump's new order eliminated key federal AI oversight policies established under the previous administration, signalling a departure from the regulatory framework put in place by President Biden.

It also sought to encourage AI development by reducing regulatory barriers, allowing companies greater freedom to innovate without stringent federal oversight. Additionally, the order emphasised

an industry-driven approach to AI governance, positioning the private sector as the primary driver of AI advancements. Trump framed this shift as necessary to ensure that AI development remains free from what he referred to as “ideological bias. The Trump Administration later released ‘Winning the Race: America’s AI Action Plan’ on Jul 31, 2025, detailing deliverables across its agencies.

While the Trump administration’s latest moves represent a decisive push toward AI deregulation, it does not mean that all federal AI-related frameworks have been abandoned. Several non-legislative federal initiatives remain in place, reflecting a bipartisan acknowledgement of AI’s growing role in national security and economic competitiveness. The Department of Homeland Security’s “Roles and Responsibilities Framework for Artificial Intelligence in Critical Infrastructure” continues to provide guidance for AI deployment in key sectors such as energy, water, and transportation, though it remains voluntary under Trump’s governance. Similarly, the bipartisan SAFE Innovation AI Framework, which serves as a set of recommendations for AI developers, companies, and policymakers, has remained in discussion as an industry-led effort rather than a federally enforced mandate.

At the same time, agencies such as the Securities and Exchange Commission (SEC) have taken independent steps to address AI-related risks. In September 2024, the SEC published an AI Compliance Plan to promote innovation while managing risks associated with AI in financial markets. However, under Trump’s administration, the SEC’s enforcement approach may be softened in favour of a more industry-friendly stance, prioritising economic growth over regulatory intervention.

Beyond regulatory rollback, the Trump administration has launched Project Stargate, a $500 billion AI infrastructure initiative backed by private sector investment from OpenAI, Oracle, Japan’s SoftBank, and MGX, the UAE government's tech investment arm. This signals a major push toward AI commercialisation and global leadership, reinforcing Trump’s belief that AI innovation should be driven by market forces rather than government mandates.

In short, the recent political and economic turbulence means the future of AI regulation in the US in the coming years is far from certain.

While the federal government has pivoted to deregulation, states are actively passing their own AI laws, creating a fractured regulatory landscape. In the 2025 legislative session, all 50 states, Puerto Rico, the Virgin Islands, and Washington, D.C., have introduced legislation, while thirty-eight states adopted or enacted around 100 measures.

On May 17, 2024, Colorado became the first U.S. state to pass a comprehensive AI law known as the Colorado AI Act. The legislation requires AI developers and deployers of high-risk AI systems to exercise reasonable care to prevent algorithmic discrimination. It also mandates clear disclosures to consumers when interacting with AI systems. This law is widely expected to serve as a model for future AI regulations in other states, particularly as concerns over AI bias and transparency continue to grow. However, the bill was amended in August 2025 to delay implementation to June 2026, and further changes may be considered over the course of the year.

California had proposed the Safe and Secure Innovation for Frontier Artificial Intelligence Models Act, a bill designed to mandate safety tests for powerful AI models and establish a publicly funded cloud computing cluster for AI research and development. However, in September 2024, Governor Gavin Newsom vetoed the bill, arguing that it could stifle AI innovation and harm California’s competitive edge in the technology sector. The decision reflects a wider debate over how far state regulations should go before they begin to limit industry growth and global competitiveness.

Instead of passing sweeping AI regulations, many states are focusing on targeted laws addressing specific AI applications. In March 2024, Utah enacted the Artificial Intelligence Policy Act, which

establishes liability for companies that fail to disclose the use of generative AI when required. Similarly, New York and Florida have introduced legislation aimed at regulating AI-driven hiring algorithms, bias mitigation, and consumer protection. These state-level efforts reflect a growing concern about how AI affects fairness in hiring practices and corporate accountability.

The proliferation of deepfake content has led to a surge in state-level legislation aimed at AI-generated impersonation and non-consensual deepfake pornography. Following the Taylor Swift deepfake scandal in January 2024, several states, including California, New York, and Texas, introduced laws criminalising non-consensual deepfake pornography and placing stricter regulations on AI-generated impersonation. By the end of 2024, hundreds of bills had been introduced, and at least 40 new laws had been enacted across more than 25 states, with a primary focus on criminalising AI-generated explicit content and protecting individuals from digital exploitation. California has implemented laws to protect minors from AI-generated harmful imagery, making it clear that such content is illegal even if produced by AI

With each state pursuing its own AI policies, businesses that operate across multiple states must now navigate a complex patchwork of AI regulations with different compliance requirements, liability risks, and enforcement mechanisms. As AI adoption continues to grow, this fragmented approach to regulation is likely to create new challenges for companies while also shaping the national conversation on AI governance in the absence of comprehensive federal legislation.

China was one of the first countries to establish a national AI strategy, publishing the New Generation Artificial Intelligence Development Plan in July 2017. This blueprint set ambitious goals for AI development, with targets extending to 2030. Over the years, China has introduced sector-specific AI regulations, including rules for managing AI algorithms, generative AI services, and deep synthesis technologies. In May 2024, scholars proposed a draft of the Artificial Intelligence Law of the People’s Republic of China, which—if enacted—would impose legal requirements on AI developers and deployers, particularly those working on high-risk or “critical” AI-based systems.

In March 2025, authorities issued the final Measures for Labelling AI-Generated and Synthetic Content, which came into effect in September and requires labels for AI-generated content and detection mechanisms on platforms. Additionally, the State Administration for Market Regulation and the Standardisation Administration of China jointly released three national standards aimed at enhancing the security and governance of generative AI, standards which officially took effect on November 1, 2025.

While China has embraced AI innovation, the government remains highly invested in controlling its development and deployment. According to recent reports, Beijing has adopted a dual approach, allowing companies and research labs to push the boundaries of AI development while maintaining strict regulatory oversight on public-facing AI services.

In September 2024, China’s National Technical Committee 260 on Cybersecurity released the AI Safety Governance Framework, aligning with the Global AI Governance Initiative. This framework introduced guidelines for the ethical and secure development of AI technologies. At the same time, the Cyberspace Administration of China (CAC) began implementing stricter rules on AI-generated content, including a mandatory watermarking system to combat misinformation and enhance public trust in digital media. These measures include audio Morse codes, encrypted metadata, and VR-based labelling systems designed to clearly distinguish AI-generated content from authentic material.

One of the most significant AI developments in China has been the meteoric rise of DeepSeek, a Hangzhou-based AI startup that has quickly positioned itself as a global competitor to OpenAI and Google’s Gemini. In February 2025, DeepSeek accelerated the launch of its R2 model, which

boasts comparable capabilities to Western AI models but at a fraction of the training costs. This development sent shockwaves through the global AI industry, contributing to one of the largest single-day drops in US tech stock values in history.

DeepSeek's rapid ascent has also drawn the attention of policymakers and defence agencies in the West. The US Navy has officially banned the use of DeepSeek AI among its personnel, citing security concerns over potential data vulnerabilities and foreign influence. Meanwhile, discussions within the US and EU regulatory circles have intensified over whether stricter export controls should be imposed on AI technologies to limit the transfer of cutting-edge capabilities to China.

The political and economic turbulence surrounding AI competition—particularly the growing divide between US and Chinese AI ecosystems—suggests that the future of AI regulation will be shaped not just by domestic concerns but also by global technological rivalries.

In 2022, Japan released its National AI Strategy, promoting the notion of "agile governance," whereby the government provides non-binding guidance and defers to the private sector's voluntary efforts to self-regulate. In 2023, to coincide with the G7 Summit, Japan published the Hiroshima International Guiding Principles for Organisations Developing Advanced AI Systems, which aim to establish and promote guidelines worldwide for safe, secure, and trustworthy AI.

In 2024, the Government ramped up its efforts to legislate for AI's use in Japan. In February, the government disclosed the rough draft of the "Basic Law for the Promotion of Responsible AI" (AI Act). In April 2024, the Ministry of Economy, Trade and Industry issued the AI Guidelines for Business Ver 1.0, intended to offer guidance to all entities (including public organisations, such as governments and local governments) involved in developing, providing, and using AI. The guidelines encourage all entities involved in AI to follow 10 principles: safety, fairness, privacy protection, data security, transparency, accountability, education and literacy, fair competition, innovation, and a human-centric approach that "enables diverse people to seek diverse well-being."

Japan subsequently enacted the Act on the Promotion of Research and Development and Utilisation of AI-Related Technologies (the AI Promotion Act) which was passed in May 2025, and entered into effect in June. The Act represents a non-binding national framework that sets strategic direction, coordination, and transparency goals without imposing penalties.

Around the world, at least 69 countries have proposed over 1000 AI-related policy initiatives and legal frameworks to address public concerns around AI safety and governance.

In Australia, the government has highlighted the application of existing regulatory frameworks for AI to compensate for the current laws and policies governing the technology, which experts warn may leave the country "at the back of the pack". Since then, the Australian Department of Industry, Science and Resources released the Voluntary AI Safety Standard in August 2024, and in September, Australia's Digital Transformation Agency released its policy for the responsible use of AI in government. Finally, in December 2025, Australia published its National AI Plan, setting out how the Government would "support Australia to build an AI-enabled economy that is more competitive, productive and resilient".

In India, a task force has been established to make recommendations on ethical, legal and societal issues related to AI and to establish an AI regulatory authority. According to the country's National Strategy for AI, India hopes to become an "AI garage" for emerging and developing economies. Furthermore, in late 2025, the Government unveiled the India AI Governance Guidelines to steer safe, inclusive, and responsible adoption.

In the UAE, President Sheikh Mohamed Bin Zayed Al Nahyan announced the creation of the Artificial Intelligence and Advanced Technology Council (AIATC) in January 2024.

In Saudi Arabia, the first regulatory framework in relation to AI, the AI Ethics Principles, was published in 2023. Saudi Arabia is actively advancing its AI and data-driven innovation as part of its Vision 2030 economic diversification strategy. The Saudi Data and Artificial Intelligence Authority (SDAIA) plays a pivotal role in this transformation, overseeing initiatives like the National Data Bank, which serves as a centralised repository facilitating AI applications across sectors such as healthcare, education, finance, and smart city development. While specific AI regulations are still under development, the nation relies on broad, technology-specific guidelines to steer AI practices.

This is by no means an exhaustive list, but it all speaks to a regulatory landscape that has taken significant strides in the last 24 months to become more defined while varying greatly across borders and jurisdictions. What this means practically for those currently using AI is that specific considerations will be required for each jurisdiction in order to maintain compliance.

In data-rich, high-stakes sectors like defence and national security, abiding by the latest AI regulations will soon evolve from recommendations and frameworks to set-in-stone laws and rules. As such, they will carry with them the potential for significant fines, penalties, censure, and damage to reputations. As 2025 progresses, organisations need to monitor regulatory developments, be aware of the requirements on the horizon, and invest significant time and capital into proper AI governance to ensure they stay ahead of the game.

Global AI Law and Policy Tracker: Highlights and takeaways

The race to dominate the data and infrastructure that powers artificial intelligence is paired with a race to create law and policy that facilitates control over this technological revolution. The passage of the EU AI Act in 2024 signaled a high-water mark in comprehensive legislation governing AI, though a more recent trend has been to temper regulatory limits on the technology in the name of competition and innovation. As the updated Global AI Law and Policy Tracker demonstrates, many nations continue to debate impactful policies, testing new governance models as the risks and rewards of AI investments are revealed daily.

While the EU considers putting a pause on the implementation of part of its AI Act, other countries are progressing efforts to pass new AI legislation. For example, South Korea finalized its AI Framework Act in January 2025, which strengthens transparency and safety requirements and offers various promotional measures, such as support for research and development and for AI adoption and workforce preparation. Likewise, Japan enacted the AI Promotion Act in May 2025, a light touch regulation that encourages companies to cooperate with government safety measures and empowers the government to publicly disclose the names of companies that use AI to violate human rights. Furthermore, China promulgated its AI Labeling Rules, which generally require service providers, as defined by China's existing suite of AI regulations, to add both explicit and implicit labels to AI-generated content.

The list of draft AI legislation is even more formidable. It includes Argentina's Bill on Personal Data Protection in AI Systems, which seeks to regulate the use of personal data used to develop AI systems beyond the scope of Argentina's existing data protection law, and India's proposed Digital India Act, which aims to update India's regulatory regime vis a vis cyberspace and provide provisions that govern AI-generated content. Proposals also include Brazil's Bill no. 2338/2023, which would create a risk-based framework that imposes risk assessments and appeal mechanisms on high-risk systems, as well as Vietnam's Draft Law on AI, which emphasizes human-centrism, risk-based management, and regulatory distinctions depending on an entity's place in the AI supply chain.

There are also new laws that indirectly impact AI development. Australia made amendments to its Privacy Act, which regulates disclosures around automated decision making, and the U.K.'s passage of the Data (Use and Access) Act amends the U.K. General Data Protection Regulation in ways intended to promote innovation and economic growth, such as clarifying when personal data can be used for scientific research and liberalizing the lawful bases available for automated decision-making.

While nations erect regulatory guardrails, many are simultaneously extending policy that attracts investment in AI development and infrastructure. Chile is one such illustration, ranking first in one Latin American AI Index as it expands construction of data centers, lays additional subsea fiber-optic cables, and promotes local AI startups. On the opposite side of the continent, Brazil plans to invest USD4 billion in AI business projects, infrastructure, training initiatives, public service improvements and regulatory frameworks.

Many of the Gulf states are likewise intent on becoming AI hubs. For example, the United Arab Emirates hosts a growing startup and research community, as well as state-of-the-art supercomputing resources, including Stargate UAE, a partnership between the Emirates, the U.S. and OpenAI to build frontier-scale compute capacity in the UAE and enable AI tools across critical sectors. Additionally, Saudi Arabia intends to leverage its young and vibrant population and centralized governance ecosystem, hosting events and attracting investment to become a leading exporter of data and AI by 2030.

South Korea also seeks to be a center of AI innovation. The nation has launched the AI Open Innovation Hub, a national AI development support platform, and plans to build the world's highest-capacity AI data center.

Of course, deregulatory signals are also on the horizon. Top of mind is talk that the EU may postpone implementation of the AI Act, encapsulated in a Digital Omnibus on AI Regulation Proposal released by the European Commission in mid-November 2025. The proposal noted several implementation challenges, including delays in designating competent authorities, as well as a lack of harmonized standards for high-risk AI requirements and necessary guidance tools. Several changes to the act were proposed, including delayed entry into application of provisions governing high-risk AI systems to align enforcement with the availability of compliance tools. Other amendments include reduced documentation requirements for small to medium-sized enterprises and reinforced oversight powers for the AI Office regarding general-purpose AI models. While this proposal is just that, it does foreshadow future negotiations over the AI Act's rollout and regulatory bite.

Furthermore, Australia's Productivity Commission issued a report that, in part, hedges against the over-regulation of AI. The report notes the chilling effect that burdensome regulation may have on investment and highlights the importance of pursuing regulatory goals at the lowest possible cost to innovation. Similarly, Canada's Competition Bureau published a report indicating that regulation specific to the AI sector can hinder innovation, impose burdens on growth and create barriers to entry for startups.

The U.S. has also been at the forefront of the deregulatory trend, with one executive order from President Donald Trump that seeks to remove barriers to AI development and stimulate innovation, and another that works to unleash prosperity through deregulation. These executive orders foretold the administration's AI Action Plan, which aims to accelerate innovation and build infrastructure by, for example, removing red tape and onerous regulation. Moreover, a December 2025 executive orderoutlines further efforts by the White House to limit U.S. state AI laws.

Where binding legislation is lacking, operational and technical standards continue to fill the void, not to mention, avoid the pitfalls of traditional lawmaking. For example, Canada's government established the AI and Data Standardization Collaborative to develop standards based on tested, multistakeholder needs and ensure consistency across domestic and international frameworks. The collaborative advances the idea that standards make products and services safer while also fostering innovation.

The Australian Department of Industry, Science and Resources released the Voluntary AI Safety Standard, which comprises 10 guardrails for developing safe and responsible AI, including testing, transparency and accountability requirements. Balancing risk and reward, the standard aims to ensure reliable AI in high-risk settings as well as enable the flourishing of AI in low-risk settings.

Other nations working in this space include China, whose Standardization Administration released three standards to improve generative AI security, and India, whose Ministry of Electronics and Information Technology is working with industry stakeholders to develop various standards for such metrics as reliability, explainability and privacy. Likewise, Kenya's Bureau of Standards released a Draft IT AI Code of Practice to provide guidance for AI applications based on a common framework, and the U.K.'s AI Standards Hub is an initiative dedicated to global AI standardization.

Use of copyrighted data to train AI systems remains a contentious legal issue. Recent developments on this front include a public consultation process initiated by Hong Kong's Commerce and Economic Development Bureau and Intellectual Property Department to assess potential updates to copyright laws to create an exception for computational data analysis and processing. Along these lines, the U.S District Court for the Northern District of California found that training an AI model on copyrighted works likely qualifies as fair use. However, the court also found that storage of the same works in a central library only constitutes fair use if those works were legally obtained.

Less recently but nonetheless relevant, Japan amended its Copyright Act, which now permits the use of copyrighted works for AI development and training purposes, so long as the use is not intended to replicate the work's expressive content. Additionally, Israel's Ministry of Justice issued an opinion holding the use of copyrighted material is permitted for machine learning purposes.

Amid all this policymaking, international cooperation over AI governance continues. Singapore stands out as a leader in diplomacy, unveiling an initiative with the U.S. to create interoperability between the two nations' governance frameworks, and signing agreements with Australia and the EU AI Office to cooperate on AI safety and innovation. Similarly, Brazil's data protection authority, the Autoridade Nacional de Proteção de Dados, met with France's DPA, the Commission nationale de l'informatique et des libertés, to strengthen cooperation on AI, data protection and digital education; in addition, Brazil and Nigeria signed a memorandum to strengthen their collaboration on AI development and technology transfer. Furthermore, Canada joined several other G20 nations in drafting a set of principles to guide adoption of AI in the telecommunications industry, focusing on growth, security and societal benefits. Finally, the U.K. and Qatar committed to increased collaboration on AI research.

Even as deregulatory trends arise, law and policy governing AI nonetheless continues to proliferate, if only in new and creative expressions. Some jurisdictions continue to debate comprehensive laws, others consider the efficacy of standards, while most remain convinced that international collaboration and diplomacy is integral to the successful governance of this disruptive technology. Follow the IAPP's Global AI Law and Policy Tracker to stay up to date with these evolving policy positions.

The New Rules of AI: A Global Legal Overview

As AI continues to transform industries, businesses must navigate a complex regulatory environment while balancing innovation with compliance. In this report we highlight current legal frameworks and offer strategic insights for companies seeking to leverage AI responsibly. By understanding the regulatory landscape and implementing robust compliance measures, organizations can harness the potential of AI while safeguarding against legal and ethical pitfalls.

The EU AI Act has entered into force as the world's first comprehensive AI-focused law. It is intended to apply to companies with a physical presence in the European Economic Area (EEA) and notably to companies without a physical presence in the EEA in certain circumstances. The act regulates two types of AI: foundational models (general-purpose AI models, or GPAI models) and applications built on GPAI models (AI systems):

The AI Act imposes a tiered set of obligations based on the risk associated with different GPAI models and AI systems, ranging from minimal-risk to high-risk and prohibited categories:

Importantly, the European Commission on November 19, 2025 published its "digital omnibus" legislative proposals, which seek to amend multiple major EU digital regulatory laws, notably the AI Act and the GDPR. These proposals, if enacted into law, are intended to ease AI-related compliance obligations, including by (for example):

These proposals, which are welcome news for companies active in the EU, will be considered in more detail early in 2026, and there is no guarantee that these will be enacted into law.

The United Kingdom, which has for several years adopted (relative to the EU) a compliance-lite approach to AI regulation, published its AI Opportunities Action Plan. The plan seeks to position the nation as a global leader in AI technology, leveraging both private and public sector solutions to enhance public services and drive economic growth. The plan emphasizes the creation of data centers and technology hubs, with a focus on AI safety and regulation that aligns with pro-growth ambitions.

The United States has taken a bold step with the release of "America's AI Action Plan," which outlines a strategic framework for securing the nation's dominance in AI through innovation, infrastructure development, and international diplomacy. The plan marks a significant shift toward a deregulated environment, encouraging private sector–led innovation while emphasizing the importance of AI systems being free from ideological bias. The plan also highlights the need for rapid AI skill development and the establishment of regulatory sandboxes to accelerate AI deployment in critical sectors.

State-level initiatives in the United States, such as California's AI transparency law and Texas's Responsible AI Governance Act, reflect a growing recognition of the need for tailored AI regulations that address specific risks while fostering innovation. These efforts are complemented by the Federal Trade Commission's Operation AI Comply, which aims to address deceptive AI practices and ensure consumer protection. The US Securities and Exchange Commission's focus on cybersecurity and AI underscores the importance of regulatory vigilance in safeguarding against potential risks.

The White House's December 2025 executive order marks a major shift toward a unified national policy framework for AI, with broad implications for technology companies, state governments, and regulated industries. It aims to establish a minimally burdensome national standard for AI policy, limiting state-level regulatory divergence.

In Asia, Singapore’s Infocomm Media Development Authority launched the Model AI Governance Framework for Generative AI to address concerns and facilitate innovation in generative AI (GenAI). India is actively shaping its AI regulatory landscape with initiatives and guidelines for responsible AI development and deployment, but lacks specific AI laws. India is also considering AI-related laws to serve as a companion to its Digital Personal Data Protection Act 2023.

AI Watch: Global regulatory tracker - United States

Laws/Regulations directly regulating AI (the “AI Regulations”) Currently, there is no comprehensive federal legislation or regulations in the US that regulate the development of AI or specifically prohibit or restrict their use. President Trump has signaled a permissive approach to AI regulation, issuing an Executive Order for Removing Barriers to American Leadership in AI ("Removing Barriers EO") in January 2025, that rescinds President Biden's Executive Order for the Safe, Secure, and Trustworthy Development and Use of AI ("Biden EO").1 The Removing Barriers EO calls for federal departments and agencies to revise or rescind all policies, directives, regulations, and other actions taken by the Biden administration that are "inconsistent" with "enhanc[ing] America's global AI dominance." Many policies were already in place from the Biden EO and it remains to be seen what the extent of the changes will be. In July 2025, the Trump administration published the America's AI Action Plan2 ("the Plan"), which identifies more than 90 federal policy actions, with an aim to secure US AI leadership in artificial intelligence and place innovation at the core of US AI policy. This approach contrasts with the more risk-focused approaches adopted by the European Union and certain state-level initiatives such as the Colorado AI Act. While the Plan prescribes various proposed incentives to businesses, the practical impact in states with existing or emerging AI regulatory frameworks remains uncertain. The US Congress has been considering numerous AI bills covering a wide range of issues. It is unclear if the Republican-held Congress will use this as an opportunity to enact AI legislation or focus on other priorities. That said, many of the proposed bills emphasize the development of voluntary guidelines and best practices for AI systems, reflecting a cautious approach to regulation aimed at fostering innovation without imposing strict mandates. This approach is influenced by concerns over stifling technological progress and maintaining competitiveness, particularly against countries like China (which produces approximately four STEM graduates for every STEM graduate in the US). Existing US federal laws have limited application to AI. A non-exhaustive list of key examples includes: Federal Aviation Administration Reauthorization Act, which includes language requiring review of AI in aviation. 3

National Defense Authorization Act for Fiscal Year 2019, which directed the Department of Defense to undertake various AI-related activities, including appointing a coordinator to oversee AI activities. 4

National AI Initiative Act of 2020, which focused on expanding AI research and development and created the National Artificial Intelligence Initiative Office that is responsible for "overseeing and implementing the US national AI strategy." 5

Nevertheless, various frameworks and guidelines exist to guide the regulation of AI, including: The AI Action Plan has a deregulation and pro-innovation agenda. It recommends that the Office of Management and Budget work with federal agencies to assess states' AI regulatory environments when making federal funding decisions, ensuring resources are not provided to states with restrictive legal frameworks. However, it is unclear how much impact this recommendation will have or what it will look like in practice. The AI Action Plan also emphasizes the Trump administration's key objective of enhancing the United States' AI infrastructure for geopolitical leadership while protecting against foreign adversary threats. Central to this strategy is the goal of exporting the full

AI technology stack, including hardware, models, software and applications to countries willing to join a proposed "AI Alliance." While this would create opportunities for US businesses to expand into new markets, businesses may also need to reevaluate their supply chains, partnership structures and compliance programs to avoid inadvertently granting adversaries or entities of concern access to controlled AI technologies. Another notable aspect of the AI Action Plan (also set forth in the Executive Order "Preventing Woke AI in the Federal Government")6 is the update to the federal procurement guidelines to ensure that only "unbiased" large language models (i.e., considered free from "ideological dogmas such as DEI" and other "partisan or ideological judgments)" be eligible for government use. Therefore, AI companies engaging in government contracting, or whose products may otherwise be evaluated under the forthcoming guidelines on ideological neutrality, should closely monitor developments in this area. The White House Blueprint for an AI Bill of Rights, issued under Biden, asserts guidance around equitable access and use of AI systems.7 The AI Bill of Rights provides five principles and associated practices to help guide the design, use and deployment of "automated systems" including safe and effective systems; algorithmic discrimination and protection; data privacy; notice and explanation; and human alternatives, consideration and fallbacks. While the Removing Barriers EO did not specifically revoke the AI Bill of Rights, the Trump Administration may be less likely to pursue the development of principles set out in the AI Bill of Rights, to the extent these principles are perceived as "inconsistent" with "enhanc[ing] America's global AI dominance." Nevertheless, AI developers may keep these principles in mind when designing AI systems. Several leading AI companies – including Adobe, Amazon, Anthropic, Cohere, Google, IBM, Inflection, Meta, Microsoft, Nvidia, Open AI, Palantir, Salesforce, Scale AI, Stability AI – have voluntarily committed to "help move toward safe, secure, and transparent development of AI technology."8 These companies committed to internal/external security testing of AI systems before release, sharing information on managing AI risks and investing in safeguards. The Federal Communications Commission issued a declaratory ruling stating that the restrictions on the use of "artificial or pre-recorded voice" messages in the 1990s era Telephone Consumer Protection Act include AI technologies that generate human voices, demonstrating that regulatory agencies will apply existing law to AI.9 The Federal Trade Commission (FTC), under the Biden administration, had signaled an aggressive approach to use its existing authority to regulate AI.)10 The FTC issued a warning to market participants that it may violate the FTC Act to use AI tools that have discriminatory impacts, make claims about AI that are not substantiated, or to deploy AI before taking steps to assess and mitigate risks.11 The FTC has already taken enforcement action against various companies that have deceived or otherwise harmed consumers through AI.12 As discussed below, the FTC has notably banned Rite Aid from using AI facial recognition technology without reasonable safeguards.13 That said, the AI Action Plan directs the FTC to review, and, where appropriate, seek to modify or set aside investigations, orders, consent decrees and injunctions from prior administration that may unduly burden AI innovation. Therefore, it remains to be seen how aggressive the FTC will be on AI under the Trump administration.

On September 12, 2023, the US Senate held public hearings regarding AI,14 which laid out potential forthcoming AI regulations. Possible legislation could include requiring licensing and creating a new federal regulatory agency. Additionally, US lawmakers held closed-door listening sessions with AI developers, technology leaders and civil society groups on September 13, 2023 in a continued push to understand and address AI.15

SANDBOX Act16, which seeks to establish a federal "regulatory sandbox" for AI developers to apply for waivers or modifications on compliance with federal regulations in order to test, experiment with, or temporarily offer AI products and services.

The SAFE Innovation AI Framework,17 which is a bipartisan set of guidelines for AI developers, companies and policymakers. This is not a law, but rather a set of principles to encourage federal law-making on AI.

The Stop Spying Bosses Act,19 which aims to regulate employers surveilling employees with machine learning and AI techniques.

The Draft No FAKES Act,20 which would protect voice and visual likenesses of individuals from unauthorized recreations from Generative AI.

The AI Research Innovation and Accountability Act,21 which calls for greater transparency, accountability and security in AI, while establishing a framework for AI innovation. It would create an enforceable testing and evaluation standard for high-risk AI systems and require companies that use high-risk AI systems to produce transparency reports. It also empowers the National Institute of Standards and Technology to issue sector-specific recommendations to regulate them.

The American Privacy Rights Act, which would create a comprehensive consumer privacy framework.22 The draft bill includes provisions on algorithms, including a right to opt-out of covered algorithms used to make or facilitate consequential decisions.

House Republicans had included a provision in the "One Big Beautiful Bill Act" (enacted July 4, 2025) proposing a 10-year moratorium on state and local AI regulations. This move purportedly aimed to establish uniform federal oversight but sparked bipartisan opposition from many state lawmakers concerned about losing the ability to address AI-related harm locally. Ultimately, on July 1, 2025, the Senate voted 99-1 to remove the proposed moratorium – there was broad agreement among lawmakers and consumer protection advocates that the provision was unreasonably vague and likely to spur a wave of litigation about its scope and impact. Various lawmakers also expressed concerns that the moratorium would undermine efforts to regulate AI for purposes such as improving children's online safety and preventing deceptive trade practices.

Absent federal legislation, we can expect to see more state laws, and state legislatures have already introduced a substantial number of bills aimed at regulating AI, notably:

On May 17, 2024, Colorado enacted the first comprehensive US AI legislation, the Colorado AI Act. The Act creates duties for developers and for those that deploy AI. Unlike certain state privacy laws, there is no revenue threshold for applicability – the Act applies to all developers and deployers of high-risk AI systems in Colorado. The Act focuses on automated decision-making systems and defines a covered high-risk AI system as one that "when deployed, makes, or is a substantial factor in making a consequential decision" that has a material legal or similarly significant effect on the provision or denial to any consumer of, or the cost or terms of: education, employment, essential government services, healthcare, housing, insurance, and legal services. There is a specific focus on bias and discrimination, and developers and deployers must use reasonable care to avoid discrimination via AI systems that make, or are a substantial factor in making a consequential decision in the above enumerated fields. The Act will go into effect in 2026. Before the law takes effect, the state legislature will host special legislative sessions to address the fiscal impact of the Act and practical considerations relating to its implementation. This will offer an opportunity to respond to industry concerns and clarify some of the legislative requirements. 23

The Colorado AI Act is emerging as a key template for state AI regulation. For example, state legislatures in Connecticut, Massachusetts, New Mexico, New York, and Virginia are considering bills that would generally track the Colorado AI Act and impose safeguards against bias by AI systems.

In September 2024, California enacted various AI bills relating to transparency, privacy, entertainment, election integrity, and government accountability. Some of the key laws include: Assembly Bill 2655 (in effect): Defending Democracy from Deepfake Deception Act: 24 requires large online platforms to identify and block the publication of materially deceptive content related to elections in California during specified time periods before and after an election. Additionally, under

this Act, large online platforms must label – within 72 hours of notice – certain content as inauthentic, fake, or false during specified time periods before and after an election in California. Assembly Bill 1836 (in effect): Use of Likeness: Digital Replica Act: 25 establishes a cause of action for beneficiaries of deceased celebrities to recover damages for the unauthorized use of an AI-created digital replica of the celebrity in audiovisual works or sound recordings. This Act requires deployers of AI systems to obtain the consent of a deceased personality's estate before producing, distributing, or making available the digital replica of a deceased personality's voice or likeness in an expressive audiovisual work or sound recording. Note: In March 2024, Tennessee enacted a similar bill entitled the Ensuring Likeness, Voice, and Image Security Act ("ELVIS Act") that expands personality rights by prohibiting the unauthorized use of AI to mimic a person's name, photograph, voice, or likeness, without a license. 26 Senate Bill 942: California AI Transparency Act (effective January 1, 2026): 27 mandates that "Covered Providers" (AI systems that are publicly accessible within California with more than one million monthly visitors or users) implement comprehensive measures to disclose when content has been generated or modified by AI. This Act outlines requirements for AI detection tools and content disclosures, and establishes licensing practices to ensure that only compliant AI systems are permitted for public use. Covered Providers that violate the Act are liable for a penalty of US$5,000 per violation per day. Note: State legislatures in Washington and Virginia are currently considering similar AI transparency bills. Assembly Bill 2013 (effective January 1, 2026): Generative AI: Training Data Transparency Act: 28 mandates that developers of generative AI systems (GenAI) publish a "high-level summary" of the datasets used to develop and train GenAI systems. For example, developers of GenAI systems would need to publish a summary of the following information, which is non-exhaustive: Sources and owners of the datasets Description of how the datasets further the intended purpose of the GenAI system Whether the datasets include any information protected by IP law Whether the datasets include personal information as defined in the CCPA Whether the datasets were purchased or licensed by the developer Assembly Bill 3030 (in effect): Health Care Services: Artificial Intelligence Act: 29 requires health care providers that use GenAI to generate patient communications to (i) disclaim that the communication was generated by a GenAI system, and (ii) provide clear instructions for how the patient can contact a human health care provider for assistance. Where the GenAI communication has been reviewed by a human health care provider, the disclaimer requirements do not apply. Other bills governing AI across a range of fields include: Assembly Bill 2602 (in effect): Contracts against Public Policy: Personal or Professional Services: Digital Replica Act 30 Bill 896 (in effect): Generative Artificial Intelligence Accountability Act 31 Assembly Bill 2885 (in effect): Unified Definition of Artificial Intelligence

In May 2025, the California Privacy Protection Agency (CPPA) Board finalized the long-awaited rules on cybersecurity, risk assessments and automated decision-making technologies (ADMT) ("the Proposed CPPA Regulations" 32). The CPPA must now submit the Proposed CPPA regulations to the Office of Administrative Law, which has 30 working days to review them for compliance with the Administrative Procedure Act. If approved and finalized, the Proposed CCPA regulations would impose new obligations on businesses, including providing consumers with the right to opt out of ADMT in contexts involving "significant decisions" (e.g., in housing, employment, credit, or healthcare) that replace or substantially replace human decision-making. Businesses would also be required to conduct risk assessments when engaging in certain high-risk data practices, such as selling or sharing personal information, processing sensitive personal information, using ADMT for significant decisions, using personal data to train ADMT, or using profiling technologies in education or employment-related contexts. Notably, the CPPA Board members cautioned that the Proposed CCPA Regulations are intended to be responsive to how technology is currently used and could be revisited if they do not meet the moment.

). The CPPA must now submit the Proposed CPPA regulations to the Office of Administrative Law, which has 30 working days to review them for compliance with the Administrative Procedure Act. If approved and finalized, the Proposed CCPA regulations would impose new obligations on businesses, including providing consumers with the right to opt out of ADMT in contexts involving "significant decisions" (e.g., in housing, employment, credit, or healthcare) that replace or substantially replace human decision-making. Businesses would also be required to conduct risk

assessments when engaging in certain high-risk data practices, such as selling or sharing personal information, processing sensitive personal information, using ADMT for significant decisions, using personal data to train ADMT, or using profiling technologies in education or employment-related contexts. Notably, the CPPA Board members cautioned that the Proposed CCPA Regulations are intended to be responsive to how technology is currently used and could be revisited if they do not meet the moment. Since the August 29, 2025 deadline for state legislature fiscal committees to advance bills has passed, the list of active AI-related bills in California has narrowed. Bills that advanced and remain active include: AB 1018: requires developers to conduct performance evaluations of covered automated decision systems (that make consequential decisions) before deployment. SB 53: requires developers of "foundation models" deemed to present a "critical risk" to create, follow, and publish a safety and security protocol, including catastrophic risk testing for foundation models and monitoring for critical safety incidents. SB 243: aims to counter the marketing of chatbots as a recourse to loneliness and/or mental health struggles.

Bills that failed at committee stage include: SB 420: would have required notification when automated systems made certain decisions about individuals, as well as a right to human review in certain instances. AB 412: would have required generative AI developers to disclose copyrighted training data and provide rights owners with a mechanism to request certain information.

In May 2024, the Utah Artificial Intelligence Policy Act 33 went into effect. The Act requires individuals and entities to disclose the use of GenAI in communications with consumers. For individuals and entities that engage in "regulated occupations" (i.e., those who must obtain a license or state certification to practice the occupation, such as lawyers or health care providers), the disclosure must be made "prominently" at the beginning of any communication with the consumer, regardless of whether the consumer asks whether they are dealing with a GenAI system (i.e., a proactive disclosure obligation). For individuals and entities that do not engage in "regulated occupations," the disclosure must be made "clearly and conspicuously" only if the consumer asks whether they are dealing with a GenAI system (i.e., a reactive disclosure obligation). Non-compliant individuals and entities may be fined up to US$2,500 per violation by the Utah Division of Consumer Protection. Enactment of the Utah Artificial Policy Act and the California Health Care Services: Artificial Intelligence Act both underscore legislative concern with ensuring that consumers know when they are dealing with GenAI systems in certain settings.

went into effect. The Act requires individuals and entities to disclose the use of GenAI in communications with consumers. More than 40 state AI bills were introduced in 2023, with Connecticut 34 and Texas 35 actually adopting statutes. Both of those enacted statutes establish state working groups to assess state agencies' use of AI systems to ensure they do not result in unlawful discrimination.

and Texas actually adopting statutes. Both of those enacted statutes establish state working groups to assess state agencies' use of AI systems to ensure they do not result in unlawful discrimination. In Texas, Governor Greg Abbott signed the Texas Responsible AI Governance Act (TRAIGA) into law on June 22, 2025. Although the original bill mirrored the more expansive Colorado and EU AI Acts, Texas lawmakers significantly narrowed its scope during the legislative process. The final version eliminates many private sector obligations (e.g., impact assessments and consumer disclosures) and limits most compliance obligations to government use of AI. However, TRAIGA still imposes categorical restrictions on the development and deployment of AI systems for certain purposes such as behavioral manipulation, unlawful discrimination, and infringement of constitutional rights. TRAIGA also provides that AI systems may not be developed or distributed with the sole intent of producing, assisting or aiding in producing, or distributing child pornography or unlawful deepfake videos or images. Intentionally developing or distributing an AI system that engages in explicit text-based conversations while impersonating a child under the age of 18 is also prohibited. TRAIGA takes effect on January 1, 2026.

As for international commitments, on September 5, 2024, the United States joined Andorra, Georgia, Iceland, Norway, the Republic of Moldova, San Marino, the United Kingdom, Israel, and the European Union to sign the Council of Europe's Framework Convention36 on AI. The treaty will enter into force on the first day of the month following three months after five signatories, including at least three Council of Europe Member States, have ratified it. Countries from all over the world will be eligible to join and commit to its provisions. However, given changes to US AI policy under President Trump, the US may withdraw from the Framework Convention on AI or otherwise decline to adhere to the Framework.

Existing legislation has been the primary way in which the US regulates AI as established law, including privacy and intellectual property laws, which are generally applicable to AI technologies.

Notably, in April 2023, the Federal Trade Commission, Equal Employment Opportunity Commission, Consumer Financial Protection Bureau, and Department of Justice issued a joint statement noting that "existing legal authorities apply to the use of automated systems and innovative new technologies."37 As cited above, in February 2024, the Federal Communications Commission applied restrictions in the Telephone Consumer Protection Act on AI-generated voices.

Several states have enacted comprehensive privacy legislation that can also regulate AI. A non-exhaustive list of notable state legislation includes:

The Minnesota Consumer Data Privacy Act (which took effect on July 31, 2025) grants consumers additional rights compared to other state privacy laws, such as the ability to question profiling decisions and access the data used in those decisions. The consumer may also review their data used in the profiling and correct inaccurate data for reevaluation.

The Biometric Information Privacy Act in Illinois,39 which is very broad and allows for extremely high damages for violations. There is currently pending litigation in the AI context

Existing intellectual property laws also apply to AI, both with respect to the data AI technologies are trained upon and the outputs of such technologies. For example, with respect to outputs, the US District Court has held that human authorship is an essential part of a valid copyright claim, and the Copyright Office will refuse to register a work unless it was created by a human being."40 There are also numerous cases before the courts in the US alleging copyright infringement, among other things, with respect to training data. On the product liability front, there is a growing number of lawsuits against large language model developers on issues such as defective design and deceptive business practices.

The National Artificial Intelligence Initiative and Removing Barriers EO define AI as "a machine-based system that can, for a given set of human-defined objectives, make predictions, recommendations, or decisions influencing real or virtual environments. Artificial intelligence systems use machine- and human-based inputs to perceive real and virtual environments; abstract such perceptions into models through analysis in an automated manner; and use model inference to formulate options for information or action."41

A recent Texas statute establishing an AI advisory council (HB 2060) defines an "automated decision system" as "an algorithm, including an algorithm incorporating machine learning or other artificial intelligence techniques, that uses data-based analytics to make or support governmental decisions, judgments or conclusions" 42

Connecticut's Public Act No. 22-15 defines "profiling" as "any form of automated processing performed on personal data to evaluate, analyze or predict personal aspects related to an identified or identifiable individual's economic situation, health, personal preferences, interests, reliability, behavior, location or movements" 43

The CCPA defines "profiling" as "any form of automated processing of personal information, [...] to evaluate certain personal aspects relating to a natural person and in particular to analyze or predict aspects concerning that natural person's performance at work, economic situation, health, personal preferences, interests, reliability, behavior, location, or movements." 44 Additionally, recently enacted California Assembly Bill 10084543 clarifies that the CCPA applies to consumers' "personal information" regardless of its format. Specifically, AB 1008 45 clarifies that the CCPA encompasses "personal information" contained in "abstract digital formats" (i.e., generative AI systems that are capable of outputting consumers' personal information). Further, recently enacted California Senate Bill 1223 46 clarifies that "sensitive personal information" under the California Privacy Rights Act (CPRA) encompasses consumers' neural data. As with AB 1008, SB 1223 aims to keep pace with emerging technology (in this case, neurotechnology) in an effort to protect information about consumers' brain and nervous system functions. While SB 1223 does not articulate a specific nexus to AI systems, if signed into law, it would constrain developers and deployers from using neural data under the CPRA.

The recently finalized CPPA's rules on cybersecurity, risk assessments and ADMT defines "ADMT" as "any technology that processes personal information and uses computation to replace human decision-making or substantially replace human decision-making" (note: the rules defines "substantially replace human decision-making" as a business using the technology's output to make a decision without human involvement).

As noted above, there are currently no comprehensive federal laws that have been enacted to specifically regulate AI. Accordingly, there is no specific territorial scope of federal legislation. However, many existing statutes regulate activities in which AI can be used, and those federal statutes typically apply nationally and, in some cases, extra-territorially. State legislation regulating AI generally has extra-territorial effect as its application typically extends to entities that target its residents from within or outside the state.

As noted above, there are currently no comprehensive federal laws that directly regulate AI. Accordingly, there is no specific federal sectoral scope at this stage. Nevertheless, there are certain sector-specific frameworks that have been implemented in the US to regulate the use of AI. A non-exhaustive list of key examples includes:

In the insurance sector, the National Association of Insurance Commissioners issued a model bulletin 47 that focuses on governance frameworks, risk management protocols and testing methodologies that insurers should have in place to govern their use of AI systems that impact insurance consumers. Once adopted by the NAIC (expected early 2024), state insurance departments could use the bulletin at their discretion as the bulletin is not new law, but instead enforces the application of current laws to insurers' use of AI and serves as guidance as to regulatory expectations

that focuses on governance frameworks, risk management protocols and testing methodologies that insurers should have in place to govern their use of AI systems that impact insurance consumers. Once adopted by the NAIC (expected early 2024), state insurance departments could use the bulletin at their discretion as the bulletin is not new law, but instead enforces the application of current laws to insurers' use of AI and serves as guidance as to regulatory expectations In the employment sector, the City of New York enacted Local Law 144 of 202148 that "prohibits employers and employment agencies [in the city] from using an automated employment decision tool unless the tool has been subject to a bias audit within one year of the use of the tool, information about the bias audit is publicly available, and certain notices have been provided to employees or job candidates"49

As noted above, there is currently no comprehensive federal legislation in the US that directly regulates AI. Accordingly, there are currently no specific or unique federal obligations imposed on developers, users, operators and/or deployers of AI systems. However, developers, users, operators and deployers of AI systems should anticipate that existing law will apply to any regulated activity that uses AI, and consult legal counsel about the potential liabilities that may arise. While potentially novel, the use of AI does not per se provide a shield from the application of existing law.

As noted above, there is currently no comprehensive legislation in the US that directly regulates AI. However, proposed legislation at the federal and state level generally seeks to address the following issues:

As noted above, there is currently no comprehensive legislation in the US that directly regulates AI. AI is also not generally classified according to risk in the relevant frameworks and principles.

Currently, there is no AI-specific federal regulator in the US. However, in April 2023, the Federal Trade Commission, Equal Employment Opportunity Commission, Consumer Financial Protection Bureau and Department of Justice issued a joint statement clarifying that their authority applies to "software and algorithmic processes, including AI."50

Similarly, state regulators that regulate privacy legislation likely also have the authority to regulate AI vis-à-vis existing privacy provisions. The FTC has been active in this area, and we can expect to see more from them going forward; see discussion of Rite Aid below.

As noted above, there are currently no comprehensive federal laws or regulations in the US that have been enacted specifically to regulate AI. As such, enforcement and penalties relating to the creation, dissemination and/or use of AI are governed by application of existing law to situations involving AI, through regulatory or judicial application of non-AI-specific federal and state statutes or AI-specific state privacy legislation.

In addition, the Federal Trade Commission has evoked an interest in and focus on regulating AI through enforcement. On December 19, 2023, the FTC settled a significant action focused on artificial intelligence bias and discrimination against Rite Aid regarding the company's use of facial recognition technology for retail theft deterrence. This illustrative case provides guidance on the FTC's enforcement on AI systems. For example, the proposed consent order51 between Rite Aid and the FTC:

Specifies that after Rite Aid's ban on using AI facial recognition expires, if Rite Aid operates AI facial recognition technology for surveillance, it must maintain a comprehensive automated biometric security or surveillance system monitoring program that identifies and addresses the risks of such operation and notifies consumers of its use of AI facial recognition. Rite Aid must also provide a means for consumers to lodge complaints, and investigate and respond to all complaints received, among other requirements

With respect to Colorado AI Act, the Colorado Attorney General has rule-making authority to implement, and exclusive authority to enforce, the requirements of the Act.52 A developer or deployer who violates the Act is deemed to engage in unfair or deceptive trade practices.

Enforcement mechanisms and penalties vary under the different California AI bills. Bills that specifically provide for enforcement include:

Senate Bill 942: California AI Transparency Act: provides for penalties of US$5,000 per violation per day, enforceable through civil action by the California Attorney General, city attorneys, or county counsel

Assembly Bill 3030: Health Care Services: Artificial Intelligence Act: enforceable by the Medical Board of California and Osteopathic Medical Board of California, with non-compliance punishable by, inter alia, civil penalties, suspension or revocation of a medical license, and administrative fines as set out in the California Health and Safety Code

Assembly Bill 2655: Defending Democracy from Deepfake Deception Act: the California Attorney General, any district attorney, or any city attorney may seek injunctive relief to compel removal of materially deceptive content

16 See Strengthening Artificial Intelligence Normalization and Diffusion by Oversight and eXperimentation Act ("SANDBOX Act")

White & Case means the international legal practice comprising White & Case LLP, a New York State registered limited liability partnership, White & Case LLP, a limited liability partnership incorporated under English law and all other affiliated partnerships, companies and entities.

This article is prepared for the general information of interested persons. It is not, and does not attempt to be, comprehensive in nature. Due to the general nature of its content, it should not be regarded as legal advice.

European AI Act Updates

EU AI Act 2026 Updates: Compliance Requirements and Business Risks

Following the transitional periods established by the European Parliament and the European Commission, the remaining provisions of the Artificial Intelligence Act will become applicable on 2 August 2026.

Most of them, in turn, are set to apply from this date, and others are being phased in over a period of six to 36 months from the date of entry into force.

A number of additional obligations will take effect ahead of the Regulation's full applicability, triggering new compliance requirements and increasing regulatory exposure for stakeholders.

In light of the associated business risks and the accelerating regulatory timeline, this article has been updated to reflect the latest developments and the obligations that will soon become applicable.

Adopted on 21 May 2024, the EU Artificial Intelligence Act lays down harmonized rules on artificial intelligence and is the first-ever comprehensive legal framework on AI worldwide. It sets out risk-based rules for AI developers and deployers regarding specific uses of AI. Please, find more about AI Act applicability here.

The EU's regulatory model for artificial intelligence is founded on a graduated, risk-oriented structure. Rather than targeting specific technologies, the framework differentiates AI systems by the potential harm they may pose, imposing escalating obligations where risks increase. AI systems which significantly impact fundamental rights, are therefore either prohibited or subject to stricter requirements and human oversight.

At the same time, it can also be misused and provide novel and powerful tools for manipulative, exploitative and social control practices.

Such practices are particularly harmful and abusive and should be prohibited because they contradict Union values of respect for human dignity, freedom, equality, democracy and the rule of law and fundamental rights.

For example, AI systems providing social scoring of natural persons by public or private actors may lead to discriminatory outcomes and the exclusion of certain groups. They may violate the right to dignity and non-discrimination and the values of equality and justice. Huge business entities already use social scoring – Uber uses scoring to evaluate customer and driver behavior and create black-lists of users. Uber relies on Open AI models, which use a combination of qualitative and quantitative metrics, including customer experience scores.

Based on the potential risks associated with the use of such AI systems, Chapter II of the Act has been dedicated to the list of prohibited AI practices.

The prohibited practices are listed exhaustively in Article 5 and provide a framework for what AI Systems can and cannot do within the EU.

The list of prohibited practices in Article 5 is exhaustive, but not final. The Commission will assess the need for amendment of the list of prohibited practices annually. There are detailed exceptions to many of the prohibitions and each practice should be considered on a case-by-case basis, which requires professional approach within the Company to avoid potential financial and opportunity losses.

With the adoption of the prohibited practices list, Section III of the AI Act regulates High-risk AI systems. These are AI systems that pose a significant risk of harm to the health, safety and fundamental rights of persons in the EU. There are two criterias to defining that AI system may be classified as High-risk:

The AI system is used as a safety component in a product that is regulated by certain EU product safety legislation and is subject to the conformity assessment procedure with a third-party conformity assessment body under such legislation, or constitutes on its own such a product; or

The AI system falls within one of the eight categories set out in Annex III of the AI Act – unless the provider can demonstrate and document that such AI system does not pose a significant risk of harm

We recommend assessing and documenting whether your product falls within the material and territorial scope of the EU AI Act, and determining whether any AI systems or AI models you develop or deploy fall within one or more of the regulated categories. It is critical to verify that the AI system is not classified as a prohibited practice and to monitor regulatory updates on an ongoing basis, as the list of prohibited practices may be revised annually.

Requirements for High-risk AI systems are outlined in the Section II of the Act with the detailed list of obligations for providers and deployers of High-risk AI systems in Sections 2, 3 and 4.

Providers of High-risk AI systems must ensure compliance with the requirements set out in Articles 8–15 throughout the system’s lifecycle, including the establishment of a documented risk management system, robust data governance measures, detailed technical documentation, automatic logging, appropriate human oversight, and safeguards for accuracy, robustness, and cybersecurity. Prior to placing a system on the market or putting it into service, providers must carry out the applicable conformity assessment, draw up an EU declaration of conformity, affix the CE marking, and register the system in the EU database. Ongoing obligations include corrective actions, cooperation with competent authorities, and maintaining a quality management system that enables continuous compliance.

Deployers of High-risk AI systems are required to use such systems strictly in accordance with the provider’s instructions and to implement appropriate technical and organisational measures, including assigning trained and competent human oversight.

Where deployers control input data, they are responsible for ensuring its relevance and representativeness, and they must continuously monitor system performance and report serious incidents or risks without delay. Deployers must also retain system logs, comply with transparency obligations towards affected individuals and workers, and, where applicable, conduct data protection impact assessments or fundamental rights impact assessments prior to deployment, particularly in sensitive use cases such as creditworthiness assessments, insurance pricing, or public-sector decision-making.

If deployers of High-risk AI systems are required to perform a Data Protection Impact Assessment under the Regulation (EU) 2016/679 (GDPR) they should use information provided by the provider under article 13 of the AI Act.

Although the EU AI Act focuses primarily on providers and deployers of High-risk AI systems, it also imposes targeted obligations on other actors in the AI value chain – importers, distributors, and suppliers. Importers and distributors must carry out verification checks before placing a system on or making it available in the EU market, including confirming that the required conformity assessment has been completed, technical documentation and the EU declaration of conformity are in place, the CE marking has been affixed, and storage and transport conditions do not undermine compliance.

Transparency means that AI systems are developed and used in a way that allows appropriate traceability and explainability, while making humans aware that they communicate or interact with an AI system, as well as duly informing deployers of the capabilities and limitations of that AI system and affected persons about their rights.

Chapter IV of the Act contains main requirements for providers and deployers of AI systems. Both providers and deployers are required that users are informed clearly when they interact with AI (Article 50(1)), when content is artificially generated or manipulated (including deepfakes and public-interest text), and when emotion recognition or biometric categorisation systems are used, with information provided in an accessible and timely manner. It mandates labelling of synthetic content in a machine-readable and detectable way where technically feasible (Article 50(2)), while carving out exceptions for law-enforcement uses authorised by law, minor assistive editing, and certain artistic or editorial contexts.

Identify and classify all AI use cases at an early stage, determining the role of the organisation (provider or deployer), the risk category of the system, and whether specific transparency obligations or exemptions apply under the AI Act;

Design and integrate transparency measures by default, ensuring that clear, accessible disclosures (e.g. AI interaction notices or system-use notifications) are provided to individuals at the latest at the point of first interaction or exposure.

Techniques such as watermarks, fingerprints, metadata identifications or other methods should be employed to indicate the content’s artificial nature. It is important that these labels are easily, instantly and constantly visible to the audience.

Implement appropriate technical solutions for marking AI-generated or manipulated content, selecting methods that are effective, interoperable, and proportionate to the type of content and the state of the art.

Establish operational processes for deployer obligations, including informing individuals about emotion recognition, biometric categorisation, or deepfake use, and aligning these practices with applicable data-protection and editorial responsibility requirements.

Document, review, and maintain compliance on an ongoing basis, including records of transparency measures, human oversight, reliance on exemptions, and monitoring of evolving standards, codes of practice, and regulatory guidance.

AI Act emphasizes that Member states should take all necessary measures to ensure that the provisions of this Regulation are implemented, including by laying down effective, proportionate and dissuasive penalties for their infringement. Member states will lay down the upper limits for setting

the administrative fines for certain specific infringements.

For example, The Artificial Intelligence Law in Italy (Law No. 132/2025) which entered into force on October 10, 2025, establishes fine of up to maximum EUR 774,685 and – in the most serious cases – the disqualifying measures provided under Decree 231 for up to one year. These include: disqualification from conducting business; suspension or revocation of authorisations, licences, or concessions instrumental to the offence; a ban on contracting with the public administration; exclusion from grants, financing, contributions, or subsidies (and possible revocation of those already awarded) and a prohibition on advertising goods or services. Law 132/2025 also introduces a new aggravating circumstance for crimes committed with the support of AI. A new criminal offence is established for the unlawful dissemination of AI-generated or altered content (e.g., deepfakes), punishable by imprisonment ranging from one to five years. Market manipulation offences committed through AI are subject to increased penalties under the Financial Consolidated Act (TUF) and Italian Civil Code.

The provisions on penalties under the AI Act exceed even those provided for in the GDPR. The maximum fine was set at EUR 35,000,000 or 7% of annual worldwide turnover. Such fines can be imposed by national authorities, the European Data Protection Supervisor, or the European Commission. The European Data Protection Supervisor can impose fines on Union institutions, agencies and bodies. The European Commission can impose fines on providers of general-purpose AI models. National authorities can impose fines on other operators.

We recommend closely monitoring national legislative proposals and draft laws implementing the AI Act, and adopting a continuously updated approach to AI compliance. This is particularly important as liability for violations of AI-related obligations may extend beyond administrative sanctions and, in certain jurisdictions, may also include criminal liability.

The AI Act establishes a gradual application of its provisions, including transitional arrangements for AI systems placed on the market or put into service before certain dates. Obligations under the Act will apply to all operators of High-risk AI systems in place before 2 August 2026. The key transitional dates are as follows:

from 2 February 2025: Prohibitions and general provisions concerning AI systems that pose unacceptable risks come into effect, allowing for early mitigation of such risks and their impact on civil law procedures.

Prohibitions and general provisions concerning AI systems that pose unacceptable risks come into effect, allowing for early mitigation of such risks and their impact on civil law procedures. from 2 August 2025: Governance infrastructure, including notified bodies and the conformity assessment system, must be operational. Obligations for providers of general-purpose AI models also begin on this date.

Governance infrastructure, including notified bodies and the conformity assessment system, must be operational. Obligations for providers of general-purpose AI models also begin on this date. The AI Act applies in full to all remaining obligations, including High-risk AI systems placed on the market or put into service prior to this date.

This phased approach ensures both early risk mitigation and sufficient time for operators and governance bodies to prepare for full compliance. Please find the relevant transitional dates outlined in the table below.

The EU AI Act introduces a comprehensive regulatory framework for artificial intelligence, creating significant compliance obligations for providers, deployers, importers, and distributors of AI systems. Failure to comply can result not only in administrative fines but also civil and criminal

liability, depending on the jurisdiction and the nature of the violation. Organizations must therefore adopt a proactive and structured approach to risk management and regulatory compliance. Key business risks arising from the AI Act include:

Administrative fines or sanctions as a result of failure to meet the AI Act’s requirements for High-risk AI systems, transparency, documentation, and human oversight.

Large fines (up to EUR 35 million or 7% of worldwide turnover) and negative publicity from non-compliance or misuse of AI can damage brand reputation and investor confidence.

Misclassification of AI systems as High-risk or prohibited practices may lead to mandatory recalls, suspension of deployment, or restrictions on market access.

Civil claims from affected individuals, including claims related to fundamental rights violations, discrimination, or inaccurate AI decisions.

Certain AI-related practices, such as unlawful dissemination of deepfakes or manipulation of AI systems for fraud, may constitute criminal offences under national laws.

Improper handling of personal data by AI systems, especially in biometric or emotion recognition applications, may lead to GDPR-related penalties and enforcement actions (up to EUR 20,000,000 or 4% of annual worldwide turnover).

Until 2 August 2026, we recommend classifying all AI systems, assessing whether they fall under High-risk or prohibited categories, and implementing relevant measures for risk management, human oversight, data governance, and transparency. By 2 August 2026, conformity assessments should be completed, technical documentation finalized, CE marking affixed, and EU database registration for High-risk systems completed. After 2 August 2026, organizations should continuously monitor regulatory updates, respond to consultations, cooperate with authorities, report incidents promptly, and update compliance processes to mitigate administrative, civil, and criminal risks.

We welcome you to book a consultation with one of the Legal Nodes' legal professionals to kickstart your AI compliance journey.

AI Act

The AI Act is the first-ever legal framework on AI, which addresses the risks of AI and positions Europe to play a leading role globally.

The AI Act (Regulation (EU) 2024/1689 laying down harmonised rules on artificial intelligence) is the first-ever comprehensive legal framework on AI worldwide. The aim of the rules is to foster trustworthy AI in Europe. For any questions on the AI Act, check out the AI Act Single Information platform.

The AI Act sets out a risk-based rules for AI developers and deployers regarding specific uses of AI. The AI Act is part of a wider package of policy measures to support the development of trustworthy AI, which also includes the AI Continent Action Plan, the AI Innovation Package and the launch of AI Factories. Together, these measures guarantee safety, fundamental rights and human-centric

AI, and strengthen uptake, investment and innovation in AI across the EU.

To facilitate the transition to the new regulatory framework, the Commission has launched the AI Pact, a voluntary initiative that seeks to support the future implementation, engage with stakeholders and invite AI providers and deployers from Europe and beyond to comply with the key obligations of the AI Act ahead of time. In parallel, the AI Act Service Desk is also providing information and support for a smooth and effective implementation of the AI Act across the EU.

The AI Act ensures that Europeans can trust what AI has to offer. While most AI systems pose limited to no risk and can contribute to solving many societal challenges, certain AI systems create risks that we must address to avoid undesirable outcomes.

For example, it is often not possible to find out why an AI system has made a decision or prediction and taken a particular action. So, it may become difficult to assess whether someone has been unfairly disadvantaged, such as in a hiring decision or in an application for a public benefit scheme.

Although existing legislation provides some protection, it is insufficient to address the specific challenges AI systems may bring.

EU AI Act Update: Navigating the Future

On July 10, 2025, the GPAI Code of Practice and related FAQs were published. The GPAI Code of Practice aims to aid compliance with the AI Act's obligations.

The Commission's decision to forge ahead with AI regulations reflects its proactive stance on ensuring that AI technologies are developed and deployed responsibly. The regulations, which have been in the works for several years, aim to create a comprehensive framework that addresses the ethical, legal, and societal implications of AI. Major technology players and European businesses have recently urged the Commission to delay the AI Act by years.

Adding to the momentum, this week, the European Commission received the final version of the GPAI Code of Practice. This landmark document is poised to set the standard for ethical AI deployment across Europe, ensuring that AI technologies are developed and used in ways that are transparent, accountable, and aligned with fundamental human rights.

The GPAI Code of Practice is voluntary for organizations looking to demonstrate compliance with the EU AI Act. The Commission expects that organizations that agree to follow them will have a "reduced administrative burden" and "more legal certainty" compared to those that choose to show compliance in other ways.

The GPAI Code of Practice has three chapters: "Transparency" and "Copyright," both of which address all providers of general-purpose AI models, and "Safety and Security," relevant only to a limited number of providers of the most advanced models. Within the "Transparency" chapter is a model documentation form for providers to document compliance with the AI Act's transparency requirements.

As the European Commission moves forward with the implementation of the GPAI Code of Practice, the focus will be on fostering a culture of ethical AI development. This will involve ongoing dialogue with stakeholders, continuous monitoring of AI systems, and regular updates to the GPAI

Code of Practice to keep pace with technological advancements. Additionally, Commission guidelines on key concepts related to general-purpose AI models are expected to be published soon.

August 2, 2025: From this date onward, provisions on notifying authorities, general-purpose AI models, governance, confidentiality, and most penalties will take effect.

From this date onward, provisions on notifying authorities, general-purpose AI models, governance, confidentiality, and most penalties will take effect. February 2, 2026: Guidelines should be available specifying how to comply with the provisions on high-risk AI systems, including practical examples of high-risk versus non-high-risk systems.

Guidelines should be available specifying how to comply with the provisions on high-risk AI systems, including practical examples of high-risk versus non-high-risk systems. August 2, 2026: The remainder of the legislation will take effect, except for a minor provision regarding specific types of high-risk AI systems that will go into effect on August 1, 2027.

In the meantime, the final version of the GPAI Code of Practice marks a significant milestone in the implementation of the AI Act.

Ogletree Deakins' Cross-Border Practice Group, Cybersecurity and Privacy Practice Group, and Technology Practice Group will continue to monitor developments and will provide updates on the Cross-Border, Cybersecurity and Privacy, and Technology blogs as additional information becomes available.

US AI Policy

ENSURING A NATIONAL POLICY FRAMEWORK FOR ARTIFICIAL INTELLIGENCE

By the authority vested in me as President by the Constitution and the laws of the United States of America, it is hereby ordered:

Section 1. Purpose. United States leadership in Artificial Intelligence (AI) will promote United States national and economic security and dominance across many domains. Pursuant to Executive Order 14179 of January 23, 2025 (Removing Barriers to American Leadership in Artificial Intelligence), I revoked my predecessor's attempt to paralyze this industry and directed my Administration to remove barriers to United States AI leadership. My Administration has already done tremendous work to advance that objective, including by updating existing Federal regulatory frameworks to remove barriers to and encourage adoption of AI applications across sectors. These efforts have already delivered tremendous benefits to the American people and led to trillions of dollars of investments across the country. But we remain in the earliest days of this technological revolution and are in a race with adversaries for supremacy within it.

To win, United States AI companies must be free to innovate without cumbersome regulation. But excessive State regulation thwarts this imperative. First, State-by-State regulation by definition creates a patchwork of 50 different regulatory regimes that makes compliance more challenging, particularly for start-ups. Second, State laws are increasingly responsible for requiring entities to embed ideological bias within models. For example, a new Colorado law banning "algorithmic discrimination" may even force AI models to produce false results in order to avoid a "differential treatment or impact" on protected groups. Third, State laws sometimes impermissibly regulate beyond State borders, impinging on interstate commerce.

My Administration must act with the Congress to ensure that there is a minimally burdensome national standard — not 50 discordant State ones. The resulting framework must forbid State laws that conflict with the policy set forth in this order. That framework should also ensure that children are protected, censorship is prevented, copyrights are respected, and communities are safeguarded. A carefully crafted national framework can ensure that the United States wins the AI race, as we must.

Until such a national standard exists, however, it is imperative that my Administration takes action to check the most onerous and excessive laws emerging from the States that threaten to stymie innovation.

Sec. 2. Policy. It is the policy of the United States to sustain and enhance the United States' global AI dominance through a minimally burdensome national policy framework for AI.

Sec. 3. AI Litigation Task Force. Within 30 days of the date of this order, the Attorney General shall establish an AI Litigation Task Force (Task Force) whose sole responsibility shall be to challenge State AI laws inconsistent with the policy set forth in section 2 of this order, including on grounds that such laws unconstitutionally regulate interstate commerce, are preempted by existing Federal regulations, or are otherwise unlawful in the Attorney General's judgment, including, if appropriate, those laws identified pursuant to section 4 of this order. The Task Force shall consult from time to time with the Special Advisor for AI and Crypto, the Assistant to the President for Science and Technology, the Assistant to the President for Economic Policy, and the Assistant to the President

and Counsel to the President regarding the emergence of specific State AI laws that warrant challenge.

Sec. 4. Evaluation of State AI Laws. Within 90 days of the date of this order, the Secretary of Commerce, consistent with the Secretary's authorities under 47 U.S.C. 902(b), shall, in consultation with the Special Advisor for AI and Crypto, the Assistant to the President for Economic Policy, the Assistant to the President for Science and Technology, and the Assistant to the President and Counsel to the President, publish an evaluation of existing State AI laws that identifies onerous laws that conflict with the policy set forth in section 2 of this order, as well as laws that should be referred to the Task Force established pursuant to section 3 of this order. That evaluation of State AI laws shall, at a minimum, identify laws that require AI models to alter their truthful outputs, or that may compel AI developers or deployers to disclose or report information in a manner that would violate the First Amendment or any other provision of the Constitution. The evaluation may additionally identify State laws that promote AI innovation consistent with the policy set forth in section 2 of this order.

Sec. 5. Restrictions on State Funding. (a) Within 90 days of the date of this order, the Secretary of Commerce, through the Assistant Secretary of Commerce for Communications and Information, shall issue a Policy Notice specifying the conditions under which States may be eligible for remaining funding under the Broadband Equity Access and Deployment (BEAD) Program that was saved through my Administration's "Benefit of the Bargain" reforms, consistent with 47 U.S.C. 1702(e)-(f). That Policy Notice must provide that States with onerous AI laws identified pursuant to section 4 of this order are ineligible for non-deployment funds, to the maximum extent allowed by Federal law. The Policy Notice must also describe how a fragmented State regulatory landscape for AI threatens to undermine BEAD-funded deployments, the growth of AI applications reliant on high-speed networks, and BEAD's mission of delivering universal, high-speed connectivity.

(b) Executive departments and agencies (agencies) shall assess their discretionary grant programs in consultation with the Special Advisor for AI and Crypto and determine whether agencies may condition such grants on States either not enacting an AI law that conflicts with the policy of this order, including any AI law identified pursuant to section 4 or challenged pursuant to section 3 of this order, or, for those States that have enacted such laws, on those States entering into a binding agreement with the relevant agency not to enforce any such laws during the performance period in which it receives the discretionary funding.

Sec. 6. Federal Reporting and Disclosure Standard. Within 90 days of the publication of the identification specified in section 4 of this order, the Chairman of the Federal Communications Commission shall, in consultation with the Special Advisor for AI and Crypto, initiate a proceeding to determine whether to adopt a Federal reporting and disclosure standard for AI models that preempts conflicting State laws.

Sec. 7. Preemption of State Laws Mandating Deceptive Conduct in AI Models. Within 90 days of the date of this order, the Chairman of the Federal Trade Commission shall, in consultation with the Special Advisor for AI and Crypto, issue a policy statement on the application of the Federal Trade Commission Act's prohibition on unfair and deceptive acts or practices under 15 U.S.C. 45 to AI models. That policy statement must explain the circumstances under which State laws that require alterations to the truthful outputs of AI models are preempted by the Federal Trade Commission Act's prohibition on engaging in deceptive acts or practices affecting commerce.

Sec. 8. Legislation. (a) The Special Advisor for AI and Crypto and the Assistant to the President for Science and Technology shall jointly prepare a legislative recommendation establishing a uniform Federal policy framework for AI that preempts State AI laws that conflict with the policy set forth in this order.

(b) The legislative recommendation called for in subsection (a) of this section shall not propose preempting otherwise lawful State AI laws relating to:

(ii) the functions of the Director of the Office of Management and Budget relating to budgetary, administrative, or legislative proposals.

(c) This order is not intended to, and does not, create any right or benefit, substantive or procedural, enforceable at law or in equity by any party against the United States, its departments, agencies, or entities, its officers, employees, or agents, or any other person.

Executive Order Issued to Restrict State Regulation of AI

On December 11, 2025, the Trump Administration issued Executive Order 14365, titled, Removing Barriers to American Leadership in Artificial Intelligence. Executive Order 14365 was the administration's seventh executive order supporting artificial intelligence (AI).1 This most recent executive order stipulates that a national AI policy is needed to enable the AI industry to innovate without being required to follow the AI laws of 50 different states. The administration is using the executive order as a method of creating a national AI framework by threatening to withhold $21 billion of Broadband Equity and Access Deployment (BEAD) funding, not yet deployed but allocated to the states by the Infrastructure Investment and Jobs Act (IIJA).2 States maintaining AI statutes characterized as burdensome or 'onerous' by the Department of Justice (DOJ) will be disqualified from obtaining undisbursed BEAD funds.

Executive Order 14365 seeks to force states, by withholding the remaining BEAD funds, to repeal or assert that the states will not enforce AI laws that "embed ideological bias within model" that would result in preventing the discrimination of minorities. The executive order will likely face legal challenges on the grounds of federal preemption and the unconstitutional coercion of states through the withholding of authorized funds.

In June 2025, the National Telecommunications & Information Administration (NTIA), an agency of the Commerce Department, revised the prior NTIA bidding of BEAD grants through a new program named the Benefit of the Bargain. The rebidding of the BEAD grants resulted in allowing satellite providers to deliver broadband to areas that were previously approved for fiber. The NTIA thereafter threatened to refuse to deploy the $21 billion of remaining funds. Throughout the fall of 2025, however, the states advocated for the administration to allow states to deploy these remaining funds for various broadband purposes as authorized by Congress. Instead, the administration used the non-deployed BEAD funds as leverage for establishing a national AI framework. The administration issued Executive Order 14365, the Administration's seventh AI-focused executive order,3 requiring Department of Justice (DOJ) to determine which state AI laws were "onerous". If a state's AI law was deemed onerous, the NTIA was required to declare the state ineligible to receive the additional BEAD funds.

The administration asserted in the executive order that a national AI policy was needed to allow the AI industry to innovate without being required to follow a "patchwork of 50 different regulatory regimes."4 However, the most recent Executive Order did not prohibit state AI laws; on the other hand, it allowed states to maintain their AI laws but if the AI laws were deemed onerous, the Administration would hold back congressionally authorized funding.

In the absence of a national AI law, the states pass their own AI laws to protect their citizens. These state AI laws include, for example, protecting children, limiting facial recognition and prohibiting

discrimination against protected groups. The Executive Order seeks to categorize the state AI laws as either onerous or not onerous. For example, state AI laws that prohibit discrimination against protected groups are to be deemed onerous. States with onerous AI laws are prohibited by the Executive Order from continuing to receive federal grants, such as the remaining BEAD funds, and other federal agency funding. The executive order seeks to address the absence of federal AI legislation by leveraging control over infrastructure grants. By conditioning the release of undisbursed federal resources on the rescission of certain local mandates, the administration is attempting to establish a uniform national standard that Congress has yet to codify. The Executive Order explained:

[The]… Administration must act with the Congress to ensure that there is a minimally burdensome national standard — not 50 discordant State ones. The resulting framework must forbid State laws that conflict with the policy set forth in this order. First, State-by-State regulation by definition creates a patchwork of 50 different regulatory regimes that makes compliance more challenging, particularly for start-ups. Second, State laws are increasingly responsible for requiring entities to embed ideological bias within models. For example, a new Colorado law banning “algorithmic discrimination” may even force AI models to produce false results in order to avoid a “differential treatment or impact” on protected groups. Third, State laws sometimes impermissibly regulate beyond State borders, impinging on interstate commerce.

The Executive Order also provides that any state AI laws that the DOJ determines regulate AI beyond the state’s borders may be declared onerous. If true, these laws would violate the Dormant Commerce Clause.5 The DOJ must declare a state AI law as onerous by March 11, 2026.

Under the latest executive order, states must rescind AI statutes classified as burdensome by the DOJ or risk losing undisbursed BEAD grants. The mandate establishes the following primary components:

DOJ Required to “Challenge” AI Laws. Section 3 of the Executive Order requires that, by January 10, 2026, the DOJ establish an “AI Litigation Task Force” to challenge state AI laws deemed inconsistent with Section 2 of the Executive Order. The task force includes the Secretary of Commerce, Special Advisor for AI and Crypto, Assistant to the President for Economic Policy, Assistant to the President for Science and Technology, and Assistant to the President and Counsel to the President.6 Section 2 outlines the primary objective of securing international AI leadership by implementing a streamlined federal regulatory approach. This policy seeks to bolster domestic competitiveness by ensuring national standards remain flexible and do not impose excessive compliance weights on developers. The executive order charges the DOJ with the task of determining which state laws would impede U.S. AI leadership, unconstitutionally regulate interstate commerce, are preempted by existing federal regulations, or are otherwise unlawful, as determined by the DOJ. Task Force to Evaluate State AI Laws. Section 4 of the Executive Order requires that by March 11, 2026, a group of individuals in the Administration publish a report listing “onerous” state AI laws. The state laws to be deemed onerous include those that are a burden to the country’s dominance in the AI race, laws that alter “truthful outputs,” or violate the First Amendment or the Constitution. A National AI Framework Created by Withholding Broadband Funding. Section 5 of the Executive Order requires the NTIA to issue a “Policy Notice” deeming states ineligible for the BEAD funds left over after the second re-bidding of the $42.4 billion of BEAD funds (known as non-deployed funds) if the state AI laws are deemed onerous. The Policy Notice must explain how a “fragmented State regulatory landscape for AI threatens to undermine BEAD-funded deployments … and BEAD’s mission of delivering universal, high-speed connectivity”; however, the Executive Order does not provide an explanation. Withholding All Grant Funds to Enforce the Executive Order. Section 5. (b) of the Executive Order requires federal agencies, in consultation with the Special Advisor for AI and Crypto, to withhold discretionary grant program funding if such AI laws are deemed onerous. As a condition of receiving the grants, states with existing AI laws deemed onerous are required to enter a binding agreement with the agency, agreeing that the state would not enforce the AI law during the grant’s performance period. FCC to Open a Proceeding on an AI

Standard. Section 6 of the Executive Order expands the FCC's regulation of interstate and international communications coverage area (defined under 47 U.S.C. § 151 as radio, television, wire, satellite and cable) to AI by requiring the FCC to open a new proceeding to explore the possibility of adopting a "Federal reporting and disclosure standard for AI models that preempts conflicting State laws." FTC to Determine if State AI Laws are Deceptive. Section 7 of the Executive Order requires the Federal Trade Commission (FTC) to issue a policy statement on how the prohibition of "unfair and deceptive acts or practices" under the FTC Act (15 U.S.C. § 45) would be applied to AI state laws that "require alterations to the truthful outputs of AI models." Notably, the FTC Act applies to businesses and not to state laws. AI and Crypto Advisor to Recommend Legislation. Section 8 of the Executive Order requires the Special Advisor for AI and Crypto, and the Assistant to the President for Science and Technology, to prepare a legislative recommendation establishing a uniform AI federal policy framework. The framework would preempt only the state laws that are deemed onerous, but not other state AI laws. AI laws that could not be preempted by the group's recommended legislation These include those related to child safety protections, AI compute and data center infrastructure, state government AI use and procurement, and other topics to be determined. The agency or advisor that would make future "determinations" was not provided.

The Supremacy Clause requires that only a law passed by Congress can preempt a contrary state law. While an executive order is not a law, the December 11, 2025 AI Executive Order does not seem, by its terms, to directly preempt the state AI laws.7 Rather, it uses a method of withholding congressionally allocated funding as a mechanism to force, or coerce, states to repeal the AI laws the administration disagrees with.

Many groups, such as consumer rights advocates, express concern about the lack of accountability for AI systems that may be created from the process outlined in the Executive Order.8 The groups cited possible harms from the law changing AI software, such as discriminatory AI price-fixing that could inflate prices for protected groups. When an early draft of the Executive Order was released in November 2025, Colorado Attorney General Phil Weiser and California State Senator Scott Weiner had threatened to sue the Administration if the Executive Order was issued.

In case of litigation, states could argue that only Congress, not the President, may preempt state law under Article I. Following the Supreme Court decision National Federation of Independent Business (NFIB) v. Sebelius 567 U.S. 519 (2012), states might also assert that tying compliance of the Executive Order to BEAD funding is coercive. In NFIB v. Sebelius, the conservatives on the Court (Justices Scalia, Kennedy, Thomas and Alito) held that tying federal funding for state Medicaid programs to require states to expand their Medicaid programs was unconstitutionally coercive. Here, the Executive Order requires tying federal funding to the states by either repealing AI laws or agreeing not to enforce the laws, which sets up a nearly identical factual scenario that would deem such a requirement unconstitutionally coercive.9

For more information on this Executive Order, please contact a member of our Telecommunications Team or the Phillips Lytle attorney with whom you have a relationship

Unpacking the December 11, 2025 Executive Order: Ensuring a National Policy Framework for Artificial Intelligence

On December 11, 2025, President Trump issued a new Executive Order (EO) to protect American Artificial Intelligence (AI) innovation from "the most onerous and excessive laws emerging from the States that threaten to stymie innovation." Consistent with the President's July 2025 America's AI Action Plan, the EO further indicates, "[i]t is the policy of the United States to sustain and enhance the United States' global AI dominance through a minimally burdensome national policy framework

for AI."

The EO arrives at a time of extensive state legislative activity concerning AI. Multiple sources, including the White House and the National Conference of State Legislatures, have noted that in 2025 more than 1,000 AI-related bills were introduced across all U.S. states and territories. While other sources, including the Future of Privacy Forum, have reported lower numbers based on different calculation metrics, it is undeniable that AI is an increasing area of focus for both Republican and Democratic lawmakers and that dozens of new state AI laws were enacted this year. Ultimately, the EO seeks to centralize AI policy by mobilizing the DOJ to identify and challenge "onerous" state AI laws, discourages state enactment and enforcement of AI laws that conflict with federal policy—including by conditioning federal broadband funding on policy alignment—and advances federal preemption through litigation and agency action while pressing Congress to enact a uniform national framework. In the near term, this fluid dynamic will underscore the need for adaptive compliance and governance programs and close attention to regulatory and litigation risk.

Prior to the issuance of the EO, the Trump administration supported legislative proposals that would have imposed a sweeping 10-year moratorium on new state AI laws and curtailed enforcement of many existing statutes. The proposed moratorium included limited carveouts for laws that directly promoted innovation, such as AI incubators or infrastructure programs. Those proposals—introduced first in a draft of the One Big Beautiful Bill Act (signed into law on July 4, 2025) and later revived in discussions for inclusion in the National Defense Authorization Act—were omitted after bipartisan opposition.

Rather than imposing a categorical freeze on states, the EO adopts a more "as-applied" approach. It preserves space for state innovation in specific areas—such as child safety (Section 8(b)(i)), infrastructure development (Section 8(b)(ii)), and government procurement (Section 8(b)(iii))—while empowering the federal government to challenge state AI laws that the administration views as unduly burdensome (Section 2), ideologically driven (Section 4), or incompatible with the First Amendment (Section 4).

The administration has framed this shift as necessary to counter concerns of a splintering regulatory regime. In its place, the EO works to advance a national framework intended to "ensure that children are protected, censorship is prevented, copyrights are respected, communities are safeguarded," and ultimately, that the United States "wins the AI race."

The EO creates a new AI Litigation Task Force within the Department of Justice (the "Task Force"). The Task Force is charged with identifying and challenging state AI laws that conflict with federal policy, defined broadly as to "sustain and enhance the United States' global AI dominance through a minimally burdensome national policy framework for AI."

This standard affords substantial discretion to the Attorney General. The EO expressly contemplates litigation grounded in constitutional and statutory theories, including by challenging state laws on the grounds that they unlawfully infringe on national authority to regulate interstate commerce, are preempted by federal law, or violate the First Amendment (Section 3). To operationalize this approach, the Secretary of Commerce, working in consultation with the Special Advisor for AI and Crypto and senior White House policy officials, must publish, within 90 days, an evaluation identifying state AI laws deemed onerous and appropriate for referral to the AI Litigation Task Force.

Notably, the EO singles out state laws that could require AI systems to alter truthful outputs or compel disclosures that may infringe First Amendment protections or other constitutional rights as problematic. This language furthers a clear administration concern with state-level mandates that regulate AI content generation or impose viewpoint-based obligations on AI developers and

deployers, and also continues a theme from the July 23, 2025 Executive Order 14319, “Preventing Woke AI in the Federal Government.”

The EO places existing omnibus state AI laws under scrutiny, potentially including those in California, Colorado, Texas, and Utah. The final text explicitly criticizes Colorado’s algorithmic discrimination statute for potentially compelling AI systems to “produce false results in order to avoid a ‘differential treatment or impact’ on protected groups.” This critique may support the calls to amend the Colorado law, an effort that began last year and has resulted thus far in a delay of the effective date from February 1, 2026 to June 30, 2026.

The EO also incorporates a favored mechanism of the Trump administration: conditional federal funding. Within 90 days, the Secretary of Commerce must issue a policy notice establishing eligibility requirements for the Broadband Equity, Access, and Deployment (BEAD) Program, a federal initiative aiming to increase access to high-speed internet across the country (Section 5). To the maximum extent permitted by law, states identified as having AI laws that conflict with federal policy may be rendered ineligible for certain funds.

Beyond the BEAD Program, the EO directs executive agencies (again in consultation with the Special Advisor for AI and Crypto) to review other discretionary grant programs (Section 5). Agencies are encouraged to consider conditioning grants either on a state’s agreement not to enact conflicting AI laws or on a binding commitment not to enforce such laws during the funding period.

The EO also advances a preemption strategy, and puts the Federal Communications Commission (FCC) and Federal Trade Commission (FTC) in the lead.

First, the FCC Chair is directed, within 90 days of the Commerce Secretary’s report, to initiate a proceeding to determine whether the FCC should adopt a federal reporting and disclosure standard for AI models that would preempt conflicting state laws (Section 6). This directive has prompted questions about the FCC’s institutional role in AI governance, with nearly two dozen state attorneys general filing a letter to the FCC on December 19, 2025, urging it to not issue preemptive AI regulations. Pursuant to the Communications Act of 1934, the FCC regulates “telecommunications services” and, to a more limited degree, “information services,” the latter of which are more likely to apply to AI given recent precedent. In January of this year, for example, the Sixth Circuit found that broadband internet access service (which is arguably closer in kind to telecommunications than AI) was an information service on the basis that “an ‘information service’ manipulates data, while a ‘telecommunications service’ does not.”

Second, the FTC Chair is directed to issue a policy statement clarifying how Section 5 of the FTC Act applies to AI models, particularly with respect to preempting state laws that require alterations to truthful AI outputs (Section 7). This directive aligns with the FTC’s broader AI enforcement posture, including its ongoing “Operation AI Comply,” which focuses on deceptive AI claims, unfair data practices, and algorithmic harms.

The EO continues the Administration’s call for preemptive federal AI legislation. The EO also requires the Special Advisor for AI and Crypto and the Assistant to the President for Science and Technology to jointly prepare legislative recommendations for Congress that would preempt conflicting state laws (Section 8).

Importantly, the EO instructs that these recommendations should not preempt state authority in sensitive areas, including child safety, AI compute and data-center infrastructure (other than generally applicable permitting reforms), and state government procurement and use of AI (Section 8). This carveout appears designed to address concerns raised by certain Republican and Democratic opponents of earlier moratorium efforts.

Despite growing bipartisan attention to AI, federal legislative progress remains limited. To date, only one AI-specific federal statute has been enacted in 2025: the TAKE IT DOWN Act, signed by the President in May, which criminalizes the nonconsensual distribution of intimate images and imposes notice-and-removal obligations on covered platforms. Several Republican lawmakers, including U.S. Senator Ted Cruz (TX), who previously supported a 10-year moratorium of state and local AI laws, and members of House leadership, have praised the EO as a necessary interim step.

The administration has framed this EO as a response to what it views as an urgent crisis: a rapidly fracturing AI regulatory landscape driven by state action that risks undermining economic growth, job creation, national security, and U.S. competitiveness vis-à-vis China. Officials emphasize that the EO builds on earlier AI-related executive actions in 2025 and the President's America's AI Action Plan released in July 2025, all aimed at ensuring that the United States maintains technological superiority in AI. For a more detailed analysis of America's AI Action Plan, please see Sidley's The Trump Administration's 2025 AI Action Plan – Winning the Race: America's AI Action Plan – and Related Executive Orders.

Given the response from certain public-interest groups and state officials, and that the EO itself contemplates the possibility of litigation initiated by the federal government, the stage is set for court battles, although the timing and scope of litigation is highly contingent. Potential areas of litigation controversy include the preemptive sweep of various federal statutes, the extent to which the First Amendment constrains state regulatory authority over AI (and, perhaps, federal authority to intervene on particular subjects), the scope of executive authority to condition federal grants, and whether the Major Questions Doctrine limits agency authority to regulate while waiting on Congress to step in with framework legislation.

Ultimately, many stakeholders agree on one point: durable resolution will require federal legislation. A comprehensive national AI framework could provide legal certainty for businesses, protect consumers, mitigate systemic risks, and preserve constitutional values while sustaining innovation. In the meantime, companies will need to navigate a rapidly evolving AI legal landscape through nimble compliance and governance programs, while closely monitoring regulatory, litigation, and legislative developments.

How different states are approaching AI

The race to implement AI has now shifted to the states and, at present , 34 states are studying AI, including 24 states that have created a group to study AI and another 10 states have delegated the task to a standing committee. In addition, according to tracking by the Brookings Center for Technology Innovation, 1 47 states have introduced 2 AI-related legislation in 2025. New York has introduced the most legislation, but Texas has passed the most (see Table 1). As such, we believe the time has come to analyze the state of the states as it relates to AI, and this is our first of a multi-part series of papers on the topic.

As part of our research, our focus shifted to how the United States is approaching AI in another series of posts using data provided to us by Leadership Connect. Here, we analyzed how the U.S. federal government was spending funds in pursuit of AI and found a widely dispersed purchasing landscape with no real primary focus. It was the classic hallmark of a new, but immature market. In our most recent post on that topic, we analyzed hundreds of federal government contracts and found a sharp rise in AI spending and a maturing market but with massive focus on the U.S. Department of Defense (DoD). As we noted in that paper, "DoD grew their AI investment to such a degree that all other agencies became a rounding error."

It is no surprise that nations are embracing and legislating artificial intelligence. At last count, 54 countries' governments , ranging from China, India, and the United States to Uganda, Armenia, and Latvia, have published their own national AI plans. These plans have run the gamut from plans that are highly focused on defense-related initiatives (e.g., China) to those who are more focused on societal betterment (e.g., Switzerland). Some plans are concerned with broad governance and data privacy guidance while other plans largely ignore such guardrails. Our team's earlier series of posts detailed how countries were approaching national AI plans, as well as how to interpret those plans. In our concluding paper in that series , we identified how the United States could dominate national AI success by applying lessons from the U.S. space race to invigorate people development, adopting a consortium model for AI development, and creating a partnership with another country to share talents and assets. We are gratified to see the United States moving in that direction. In terms of human talent development, the " Computer Science For All Act, " introduced in 2023, would require that the Department of Education award grants to various entities in computer science education. An example of one such consortium is the Partnership for Global Inclusivity on AI (PGIAI) , which brings together the U.S. Department of State, major U.S. tech companies, and international partners for the responsible development of AI, particularly to address quality of life issues in developing countries. Finally, under the former Biden administration, the United States developed a partnership with India (which was our specific partnership recommendation) for sharing AI infrastructure investments.

Our team used data from the Brookings Center for Technology Innovation for this analysis which is current as of June 27, 2025. A total of 260 measures related to AI were introduced in the 2025 legislative session, of which 22 have been passed. The remaining bills are pending or rejected.

When examining the language in summaries of the proposed legislation, it is immediately obvious that the focus of the legislation is protection for citizens as evidenced by two of the most frequently used words in many of the proposed bills: prohibit and disclosure. Given this, it seems clear that, at present, the states see the misuse of AI as something that citizens need to be protected against versus the appropriate use of AI as an opportunity for better services to citizens.

Analysis of the 2025 data show roughly two-thirds (173) of state AI bills were introduced by Democrats, compared to about one-third (84) by Republicans. Only three—in Minnesota, New Jersey, and Tennessee—were bipartisan efforts. Overall, Democrats have been more proactive in pushing AI governance, consistent with their broader tech-regulation tendencies. Republicans have also put forward many AI bills, often on issues like banning certain harmful uses of AI or promoting innovation. Election deepfake bans and child pornography, for example, saw bipartisan interest. On the other hand, sweeping regulatory bills (imposing strict obligations on AI developers or businesses) have mostly come from Democrats in states like California, New York, and New Jersey. Republican-led states often favored lighter-touch approaches, with Texas failing to pass a high percentage of its proposed bills.

Nonconsensual intimate imagery (NCII)/Child sexual abuse material (CSAM): 53 bills introduced and 0 currently signed into law

Some states introduced comprehensive AI governance bills, covering broad obligations for AI developers or government use, while others focused on narrow sector-specific or issue-specific bills. For instance, Colorado's "Consumer Protections for Artificial Intelligence" bill created a wide-ranging regime for "high-risk" AI systems, requiring developers and users of such AI to implement transparency, monitoring, and anti-discrimination measures. In contrast, in 2024, California started to pursue a patchwork of many targeted laws addressing things like election deepfakes, AI-generated content warnings, digital replicas of performers (including one bill for deceased performers), and training-data disclosure, rather than one overarching framework. It is worth noting that California originally pursued a comprehensive approach, but it was vetoed by the governor. As a result, California reverted to the more piecemeal approach. These different approaches are noted by analysts as a major difference between state strategies. Also, some

smaller states have primarily symbolic or exploratory bills (e.g., creating AI task forces or study commissions rather than imposing rules).

Not surprisingly, two major topics considered by legislators in 2025 were NCII and CSAM. A Maryland NCII bill illustrates the intent of take-it-down legislation (the TAKE IT DOWN Act) as it requires certain online platforms to establish a process for individuals to request removal of NCII, including synthetic NCII. A Mississippi NCII bill prohibits knowingly—and with an intent to injure—publishing or disseminating synthetic NCII or deepfakes of another person. A number of other states, including Maryland, have similar bills. The New Mexico bill prohibits non-consensual public dissemination of synthetic content that depicts a person engaging in conduct that they didn't do and the person is identifiable. A total of 16 states proposed bills on this topic and, as to be expected, all the bills proposed significant penalties for those engaging in these acts. Interestingly, most of these bills have already died in committee.

Elections were a hot button issue in bills that were introduced in 2025, and most of the prohibitions were around candidates needing to disclose if AI was used to craft an advertisement. For example, a bill in New York seeks to require any political communications that use "synthetic content" to disclose its use. While these examples cover use cases where candidates seek to improve their standing, there is also legislation focused on opponents using AI to damage it. For example, a Massachusetts bill seeks to prohibit purposefully deceiving people and injuring a political candidate or influencing an election by creating, publishing, or distributing a deepfake video or an "altered image" 90 days prior to an election. Both of these bills are still in committee and are expected to die there.

Generative AI transparency was closely tracked as well in 2025 and the most common fear that the bills addressed was the concern that a person might not know that they were communicating with a chatbot. For example, Hawaii proposed legislation that would require corporations, organizations, and individuals engaged in commercial transactions to clearly and conspicuously notify a consumer if the consumer is interacting with a chatbot or other technology capable of mimicking human behaviors. A similar bill was proposed in Massachusetts. The bill also required the providers to have a red team that would determine if the watermarks could be easily removed and would require them to report those findings to the state's attorney general. The Hawaii and New Mexico bills have died on the floor, but the Massachusetts bill is still in committee.

ADMT and high-risk AI were also major topics covered by legislative bills. The crux of these bills was recognition that safeguards are needed to protect citizens from unintended consequences of AI systems. Most of the bills from this session were styled after Colorado's AI Act, which was signed into law in 2024 and is set to go into effect in 2026. The Colorado legislation focuses on algorithmic discrimination in "consequential decisions," requires identification of the use of AI as a "substantial factor" in such decisions, imposes a duty of care, mandates transparency and accountability for developers and deployers, and affirms consumers' right to an explanation of AI's role in decision making. Legislation in a number of states, including Georgia, Illinois, Iowa, and Maryland, is closely modeled after the Colorado bill. However, the Georgia, Iowa and Maryland bills all died in committee and the Illinois bill remains in committee.

This category covered the waterfront in terms of government use of AI. For example, legislation in Georgia proposed an "AI Accountability Act" that would create a Georgia Board for Artificial Intelligence which would require government entities to develop an AI usage plan to detail specific goals, data privacy measures, and outline the role of human oversight. In contrast, a Montana bill, which was recently signed by the governor, would limit the use of AI by state and local government, would require disclosure of the use of AI systems and would require certain decisions and recommendations to be reviewed by a human in a "responsible position." Finally, a Nevada bill would require the Department of Taxation to inform the taxpayer if they are communicating using AI systems. The majority of the bills were protecting citizens from the state using AI to make any governing decisions. At present, the Georgia and Nevada bills died in committee.

There were a number of AI-related bills surrounding employment and all of them were focused on protection of employees from AI decision making. For example, an Illinois bill (currently in the assignments committee) requires employers to notify applicants if AI is used as part of the interviewing and decision-making process. A similar bill is under consideration in Pennsylvania. A number of states, including California, also put limitations on the use of AI-based workplace surveillance.

Health care was a major focus of legislation, with all bills focused on the potential issues arising from AI systems making treatment and coverage decisions. For example, California's bill was focused on protecting Californians from AI systems and "prohibits the use of specified terms, letters, or phrases to indicate or imply possession of a license or certificate to practice a health care profession." A bill in Illinois, which is awaiting the governor's signature, would outright prohibit licensed health care professionals from using AI to make therapeutic decisions, interacting with clients, or creating generative therapeutic recommendations or treatment plans. Other states, such as Indiana, did not go as far and merely required health care professionals to disclose to a patient if AI was used to make a health care decision or was used in communicating health care decisions. The Indiana bill placed similar requirements on health care insurers.

The majority of bills proposed in 2025 by states were focused on protecting citizens from the overreach of AI. States recognize the need to act on AI, albeit they are taking a variety of approaches. Still, common ground is found in areas where tangible harm is seen or widely considered a problem (i.e. CSAM, NCII, and elections).

There is a storm cloud over this progress at the federal level, and it is unclear how this will affect how states are addressing AI. Republicans in the U.S. Senate recently proposed a 10-year moratorium on states enacting any laws that regulate AI, though it was eventually dropped. Still, several lawmakers supported the measure, and we believe that this battle is just beginning, and it is unclear what direction Congress will eventually go. However, it is clear that the federal government is planning on keeping a close eye on state legislation and the federal AI Action Plan has tasked the FCC with monitoring the impact of state laws.

Until and unless prohibited by federal law, we expect states to continue pressing hard on AI legislation to address their citizens' concerns. The topic is simply too pressing for the states to stand by while the federal government tries to figure it out. Our goal in this series of papers is to better understand how and why state governments are addressing AI and to create a playbook for doing it successfully.

Key AI Regulations in 2025: What Enterprises Need to Know

As artificial intelligence continues to reshape business operations, enterprises face a growing landscape of AI regulations. 2025 marks a pivotal moment for AI regulation as state laws in the U.S. and key provisions of the EU AI Act begin to take effect. This blog post outlines the most important regulations taking effect in 2025 and what enterprises can do today to meet this moment.

On January 1, 2025, California began enforcement of three AI laws relevant to enterprises impacting AI-processed personal information, healthcare services and healthcare facilities.

Applies to: any company employing AI where an AI system could generate personal information subject to the protection of the CCPA.

AB 1008 amends the California Consumer Privacy Act (CCPA) to clarify that personal information can exist in various formats, including artificial intelligence systems that are capable of outputting personal information.

SB 1120 amends Section 1367.01 of the Health and Safety Code and Section 10123.135 of the Insurance Code, introducing specific regulations for health care service plans and disability insurers that employ artificial intelligence, algorithms, or other software tools in their utilization review (UR) and utilization management (UM) processes.

AB 3030 introduces new requirements for health facilities, clinics, and physician practices in California that use genAI to generate written or verbal patient communications pertaining to patient clinical information.

We expect the U.S. States to continue to take a leading role in AI governance regulation in 2025. Initiatives underway to look out for include:

Privacy Protection Agency (CPPA) is conducting rulemaking on automated decision-making technology. The Multistate AI Policymaker Working Group , spanning 45 states, is working toward consistent AI regulation approaches across states in the U.S.

, spanning 45 states, is working toward consistent AI regulation approaches across states in the U.S. Texas ' proposed AI Governance Act (TRAIGA) is expected to pass in 2025 given that the Texas legislature is only in session every two years.

' proposed AI Governance Act (TRAIGA) is expected to pass in 2025 given that the Texas legislature is only in session every two years. New Jersey and Illinois are considering legislation on AI in employment and insurance.

This is only the beginning of AI regulation, with some EU AI Act provisions entering into force and numerous initiatives expected to gain traction in 2025. Key global policy developments include:

In Europe , AI literacy, prohibited practices and obligations for General-Purpose AI (GPAI) models begin to apply: By February 2, 2025 , providers and deployers of AI systems must take measures to ensure a sufficient level of AI literacy of their staff (Article 4). Prohibited AI practices also go into effect on February 2nd (Article 5).

By May 2 , Codes of Practice for GPAI models –to determine if GPAI models pose systemic risk and to help developers, distributors, and deployers comply with their own requirements– must be ready (Article 56).

, Codes of Practice for GPAI models –to determine if GPAI models pose systemic risk and to help developers, distributors, and deployers comply with their own requirements– must be ready (Article 56). On August 2 , GPAI models obligations go into effect (Article 53 and 55). These include: providing technical documentation making publicly available detailed summaries about the content used for training showing compliance with EU copyright law.

South Korea enacted the Framework Act on the Development of AI and Establishment of a Foundation of Trust, a comprehensive AI law sharing common elements with the EU AI Act, on January 10th, 2025.

enacted the Framework Act on the Development of AI and Establishment of a Foundation of Trust, a comprehensive AI law sharing common elements with the EU AI Act, on January 10th, 2025. The UK unveiled the AI Opportunities Action Plan on January 13, 2025 outlining its commitment to shaping the future of AI and including enabling safe and trusted AI development and adoption through regulation, safety and assurance. The UK is expected to continue public consultations on sector specific issues relating to AI.

Common elements that emerge from laws taking effect in 2025 as well as others set to take effect in 2026, such as Colorado AI Act and EU AI Act provisions for high-risk AI. Enterprises beginning their compliance journey should prioritize three governance areas:

Document how personal information is processed by these systems Establish risk management processes Develop AI risk assessment frameworks

Ensure human review is built into risk management processes Promote AI literacy Educate teams about AI use in your organization and policies and processes already in place

As AI regulations continue to evolve, organizations must establish comprehensive yet agile governance frameworks that not only comply with emerging requirements but also enable the responsible and accelerated adoption of AI. Credo AI empowers organizations to navigate this complex landscape by identifying relevant regulations, mapping policy requirements to a streamlined set of actionable controls, and offering implementation guidance to mitigate risks effectively.

Connect with Credo AI today to explore how our AI Governance Advisory Services can help you meet AI literacy requirements and build a governance foundation for your AI initiatives.

China AI Regulation

AI Watch: Global regulatory tracker - China

Laws/Regulations/National Standards directly regulating AI (the "AI Regulations") 1 The Cyberspace Administration of China, the National Development and Reform Commission, the Ministry of Education, the Ministry of Science and Technology, the Ministry of Industry and Information Technology, the Ministry of Public Security, and the National Radio and Television Administration jointly released the Interim Measures for the Management of Generative Artificial Intelligence Services (the "AI Measures"), which is the first administrative regulation on the management of Generative AI services, which came into effect on August 15, 2023.2 On September 1, 2025, new 'Labeling Rules' came into effect, making it mandatory for AI-generated content to be implicitly labeled, and explicitly labeled where applicable. Explicit labels are those that are easily perceived by users and must be added to text, audio, images, videos, and virtual scenes, while implicit labels are embedded within the file's metadata.3 In addition, on April 25, 2025, the State Administration for Market Regulation and the Standardization Administration of China jointly released three national standards aimed at enhancing the security and governance of generative AI. These standards will officially take effect on November 1, 2025. The three standards are: Cybersecurity Technology—Generative Artificial Intelligence Data Annotation Security Specification: This standard outlines security requirements for data labeling processes used in training generative AI models.4 Cybersecurity Technology—Security Specification for Generative Artificial Intelligence Pre-training and Fine-tuning Data: It outlines requirements and evaluation criteria for ensuring the security of datasets used in the pre-training and fine-tuning phases of generative AI development.5 Cybersecurity Technology—Basic Security Requirements for Generative Artificial Intelligence Service: This standard establishes security requirements for generative AI services, encompassing user data security assessments, data protection measures, and the safeguarding of training models and datasets.6

The AI Measures came into effect on August 15, 2023. Additional national standards will officially take effect on November 1, 2025.

There are various other laws and regulations in China that do not directly seek to regulate AI, but that may affect the development or use of AI:

The Administrative Provisions on Deep Synthesis in Internet-based Information Services (the "Deep Synthesis Provisions") which came into effect on January 10, 2023 7

The Administrative Provisions on Recommendation Algorithms in Internet-based Information Services (the "Recommendation Algorithms Provisions") which came into effect on March 1, 2022 8

Data-related laws, such as the Cybersecurity Law, the Personal Information Protection Law (the PIPL), and the Data Security Law may affect several aspects of AI development and use in connection with data privacy and security

Intellectual property laws, including the Copyright Law of China, that may affect several aspects of AI development and use

Both the Civil Code of China and the Criminal Law of China may affect several aspects of AI development and use in connection with the protection of privacy rights, portrait rights and other related rights from being violated by the improper use of AI

The Measures for Review of Scientific and Technological Ethics (for Trial Implementation), which came into effect on December 1, 2023, provide that enterprises engaged in AI and certain other technology activities should undergo scientific and technological ethics reviews 9

The relevant national standards and recommended guidelines set out relevant rules and guidance regarding the use of AI technology, such as the Practical Guidelines for Cybersecurity Standards – Method for Tagging Content in Generative Artificial Intelligence Services10

The current Chinese laws and regulations do not provide a clear definition of "AI." Under the AI Measures, "generative AI technology" refers to models and related technology that have the ability to generate text, images, audios, videos, or other content.11

The AI Measures apply to companies when they provide Generative AI services to the public within China regardless of where they are incorporated.12 For the purpose of this document, "China" refers to the People's Republic of China excluding the Hong Kong Special Administrative Region of the People's Republic of China, the Macau Special Administrative Region of the People's Republic of China and Taiwan. Specifically, the Cybersecurity Administration of China has the authority to take technical and other necessary measures against foreign companies if the provision of Generative AI services (originating from outside of China) to China violates relevant laws and regulations.13

In general, the AI Measures do not adopt a sector-specific approach, while the Deep Synthesis Provisions and Recommendation Algorithms Provisions focus on internet-based information service providers. In addition, some industry sectors, such as the financial, healthcare and automobile sectors, have their own sector-specific regulations relating to the use of AI technology.

Financial sector: Guiding Opinions on Regulating the Asset Management Business of Financial Institutions 14 Evaluation Specification of Artificial Intelligence Algorithm in Financial Applications 15 Guidance on Information Disclosure for Financial Applications Based on Artificial Intelligence Algorithms 16

Guidance for the Classification and Definition of AI-Based Medical Software Products Automobile sector: Opinions on Strengthening the Management of Manufacturing Enterprises and Product Access of Smart Internet-Connected-Vehicles18

"Generative AI service provider" refers to any organization or individual that utilizes generative AI technology to provide generative AI services (including providing such services through the provision of a programmable interface or other means).19

"User of Generative AI services" refers to any organization or individual that uses Generative AI services to generate content.20 The AI Measures do not apply if an organization or institution engages in research, development, or application of generative AI technology but does not offer Generative AI services to the domestic public in China.21

The AI Measures are formulated to "promote a healthy development and regulated application of generative artificial intelligence, safeguard national security and social public interests, and protect the lawful rights and interests of citizens, legal persons and other organizations."22

The AI Measures provide that the government will implement "categorical and hierarchal regulation of Generative AI services."23 Further, the AI Measures provide that the relevant competent departments shall enhance scientific supervision methods compatible with innovative development for Generative AI technology. They shall also create classification and grading-based supervision

rules, or guidelines based on the technology's characteristics and its application in relevant industries and fields.24

However, the AI Measures do not provide for the risk categorizations of AI services, although some provisions suggest that some categories of AI services are subject to stricter scrutiny by the regulators. For example, Generative AI service providers with "public opinion attributes or the capacity for social mobilization" are required to carry out a security assessment and undergo a regulatory filing before launching such services, in addition to complying with generally applicable requirements like content moderation and tagging.25

Lawful use: Processing data lawfully, respecting intellectual property rights and obtaining consent for personal information use26

Clarify and disclose the applicable groups, occasions and uses of their services, guide users to scientifically and rationally understand and use Generative AI technology in accordance with the law and take effective measures to prevent children from overly relying on or being addicted to generative artificial intelligence services 27

Guide users in the lawful and responsible use of Generative AI technology, protect user information and handle user requests for (among other rights) accessing, copying, correcting, supplementing and deleting their personal information 29

Take measures to ensure secure and stable services, promptly address illegal content, and establish a complaints and reporting mechanism for user feedback 30

Cooperate with relevant authorities during supervisory inspections, providing necessary technical, data, and other support and assistance31

Data labeling rules: Establish clear labeling rules, ensure accurate labeling and provide training for labeling staff.32 In addition, the AI Measures specifically require that a Generative AI service provider comply with the following:

"Upholding socialist core values" – including not generating any prohibited content, such as content that incites "subversion of the state power or the overthrow of the socialist system, endangers national security and interests, damages the national image, incites splitting the country, undermines national unity and social stability, advocates terrorism, extremism, ethnic hatred and discrimination, violence, pornography, and false and harmful information" 33

Taking effective measures to "prevent discrimination based on nationality, religion, country, region, gender, occupation and health in the design of the algorithm, training data set, model generation, optimization and service provision" 34

Respecting intellectual property rights and business ethics – including keeping confidential trade secrets and refraining from "carrying out acts of monopoly and unfair competition by taking advantage of proprietary algorithms, data and platforms" 35

Respecting others' legitimate rights and interests – refraining from endangering others' physical and mental health and respecting their "rights of portrait, reputation, honor, privacy and personal information"36

Taking effective measures to "enhance the transparency, accuracy and reliability of the contents generated by Generative AI services"37

Data training: The AI Measures also introduce requirements related to training data. Generative AI service providers must use data and models from legitimate sources, respect intellectual property rights and personal information, and strive to improve the quality, authenticity, accuracy, objectivity and diversity of the training data they utilize.38

Content moderation: Generative AI service providers are required to promptly remove any illegal content, employ measures for model optimization training, and report cases to the relevant authorities.39

Reporting mechanism: Generative AI service providers must establish a complaints and reporting mechanism, where they accept and handle complaints and reports from the public and provide feedback on the outcome of these cases.40

The Cyberspace Administration of China is the lead regulator for Generative AI technology. The cyberspace, development and reform, education, science and technology, industry and information technology, public security, radio and television, press and publication, and other relevant authorities shall strengthen the management of Generative AI services in accordance with the law based on their respective functions and responsibilities.41

To effectively manage Generative AI services, Chinese regulators have been granted regulatory powers. These regulators are responsible for strengthening the management of these services based on their respective functions and responsibilities.42 They have the authority to conduct security assessments, supervisory inspections, and impose penalties for any violations in accordance with relevant laws and regulations.43

If Generative AI service providers violate the AI Measures or any other relevant laws, the relevant authorities will impose penalties in accordance with the Cybersecurity Law, the Data Security Law, the PIPL, the Law on the Progress of Science and Technology, or other applicable laws and regulations. In cases where the laws or administrative regulations are silent, the relevant authorities, based on their functions and responsibilities, may issue a warning or reprimand, order corrections within a specified time limit, or suspend the provision of related services if corrections are refused or the violation is severe.

If a violation constitutes a breach of public security management, public security administrative sanctions will be imposed in accordance with the law. If a violation constitutes a crime, criminal liability will be pursued in accordance with the law.

According to Articles 253 and 291 of the Criminal Law of China, when individuals or organizations illegally obtain personal information or disseminate false information through information networks or media platforms with the intention of disturbing public order or causing severe consequences, they can face penalties of up to seven years in prison or fines.

According to the PIPL, failure to comply with the PIPL can result in penalties such as fines up to 50 million RMB, revenue confiscation (up to 5 percent of annual revenue), and even business cessation.44

1 For the purpose of this document, "China" refers to the People's Republic of China excluding the Hong Kong Special Administrative Region of the People's Republic of China, the Macau Special Administrative Region of the People's Republic of China and Taiwan.

2 See the Interim Measures for the Management of Generative Artificial Intelligence Services, promulgated on July 10, 2023 by the Cyberspace Administration of China, National Development and Reform Commission; Ministry of Education, Ministry of Science and Technology, Ministry of

Industry and Information Technology, Ministry of Public Security, National Radio and Television Administration.

3 See Notice on the issuance of the "Measures for the Identification of Artificial Intelligence-Generated and Synthetic Content", released by the Cyberspace Administration of China, Ministry of Industry and Information Technology on 7 March 2025.

7 See the Administrative Provisions on Deep Synthesis in Internet-based Information Services, promulgated on November 25, 2022, by Cyberspace Administration of China, Ministry of Industry and Information Technology, and Ministry of Public Security.

8 See the Administrative Provisions on Recommendation Algorithms in Internet-based Information Services, promulgated on December 31, 2021 by Cyberspace Administration of China, Ministry of Industry and Information Technology, Ministry of Public Security, and State Administration for Market Regulation.

9 See the Measures for Review of Scientific and Technological Ethics (for Trial Implementation), promulgated on December 1, 2023 by the Ministry of Science and Technology, the Ministry of Education, the Ministry of Industry and Information Technology, the Ministry of Agriculture and Rural Affairs, the National Health Commission, the Chinese Academy of Sciences, the Chinese Academy of Social Sciences, the Chinese Academy of Engineering, the China Association for Science and Technology, and the Science and Technology Commission of the Central Military Commission

10 See e.g., the Practical Guidelines for Cybersecurity Standards – Method for Tagging Content in Generative Artificial Intelligence Services, released by the China's National Information Security Standardization Technical Committee on August 25, 2023.

14 See the Guiding Opinions on Regulating the Asset Management Business of Financial Institutions, promulgated by the People's Bank of China, the China Banking and Insurance Regulatory Commission, the China Securities Regulatory Commission, and the State Administration of Foreign Exchange on April 27, 2018.

15 See the Evaluation Specification of Artificial Intelligence Algorithm in Financial Application, promulgated by the People's Bank of China on March 26, 2021.

16 See the Guidance on Information Disclosure for Financial Applications Based on Artificial Intelligence Algorithms, promulgated by the People's Bank of China on November 8, 2023.

17 See the Guidance for the Classification Defining of AI-Based Medical Software Products, promulgated by the National Medical Products Administration on July 1, 2021.

18 See the Opinions on Strengthening the Management of Intelligent and Connected Vehicles Manufacturing Enterprises and Product Access, promulgated by the Ministry of Industry and Information Technology on July 30,2021.

White & Case means the international legal practice comprising White & Case LLP, a New York State registered limited liability partnership, White & Case LLP, a limited liability partnership incorporated under English law and all other affiliated partnerships, companies and entities.

This article is prepared for the general information of interested persons. It is not, and does not attempt to be, comprehensive in nature. Due to the general nature of its content, it should not be regarded as legal advice.

Preparing for compliance: Key differences between EU, Chinese AI regulations

On 12 July 2024, the Official Journal of the European Union published the EU Artificial Intelligence Act, which aims to comprehensively regulate the design, deployment and use of AI systems in the EU.

Key provisions of the act will become applicable to organizations over the course of the next three years, depending on their roles and the risks and capabilities of their services, with the first provisions and restrictions taking effect 2 Feb.

In parallel, regulatory authorities in China have also turned their attention to AI, issuing regulations and technical standards to bring service providers into compliance. The three most important are the Regulations on the Management of Algorithm Recommendations for Internet Information Services, Regulations on Deep Integration Management of Internet Information Services and Interim Measures for the Management of Generative Artificial Intelligence Services, which have been fully applicable since their adoptions in 2022 and 2023.

The evolution of AI governance frameworks in the EU and China share many similarities. For example, the European Commission's High-level Expert Group on AI published its Ethics Guidelines for Trustworthy AI in 2019, around the same time Chinese government bodies first made commitments to supervise AI.

Policymakers in the EU and China have likewise expressed similar goals of regulation, implicating topics such as ethics, data protection, safety and security. Additionally, AI regulation in Brussels and Beijing have been built off preexisting legal frameworks, including those regulating the processing of personal data.

Despite these similarities, however, there are notable differences between the two approaches. Exploring those differences can help organizations comprehend the scope of requirements in both markets and prepare for compliance.

The EU AI Act aims to lay down a "uniform legal framework … for the development, the placing on the market, the putting into service and the use of artificial intelligence systems" across the EU. Complimenting other data-related laws — including those emanating from the Commission's digital strategy, such as the Data Act and the Digital Services Act — it will operate as a horizontal regulatory framework that brings into scope a variety of actors, organizations and technologies.

Encapsulating this are a series of core definitions, including one for AI systems as machine-based systems "designed to operate with varying levels of autonomy and that may exhibit adaptiveness after deployment" and infer "how to generate outputs such as predictions, content, recommendations, or decisions that can influence physical or virtual environments" based on the input they receive,and one for general-purpose AI, which captures foundation models that can be developed and integrated into other services.

In practice, this means compliance with the EU AI Act will involve determining whether organizations fall within the law's scope, what obligations are triggered and what needs to be implemented as a result. This indicates an intent to promote maximum harmonization. However,

organizations will still need to navigate tricky issues that overlap with other laws, including those that govern personal data processing generally, and other topics like product liability and safety.

By contrast, China has opted for a very different approach. Rather than adopting a single comprehensive law, government bodies have followed a two-pronged strategy: one, draft and implement industry-specific regulations and, two, promulgate technical standards and AI governance pilot projects to build best practice and enforcement experience.

While there is substantial overlap in the approach taken by existing regulations, each turns on the typeof service offered. There is correspondingly no overarching definition of AI, though the term has been defined in certain technical standards.

Moreover, local authorities, such as in Shanghai and Shenzhen have issued their own experimental regulations to test case different regulatory approaches, although these remain relatively light-touch in terms of prescriptive obligations on companies. None of these have so far been adopted at the central level.

Consequently, a key threshold question becomes whether a particular service falls within the scope of the regulation not whether it meets the definition of an AI system. There have been initial attempts to develop a unified AI law in China, which may change this analysis.

AI governance in both the EU and China revolves around assessing levels of risk — albeit defined differently — and mitigating these through legal restrictions. However, the types of risk categories and the process of classification are notably different and carry implications for where organizations should start with compliance.

A key feature of the AI Act is how it assigns differing levels of compliance requirements depending on the risk category of the service. This classification system operates on a tiered basis, with the strictest restrictions corresponding to the highest level of risk. From a compliance perspective, a significant hurdle will be determining which risk category an AI system falls into and whether an exception may apply.

While the classification framework appears sector and technology agnostic, certain technologies are proactively targeted over others, like biometrics and emotion and facial recognition. In practice, organizations will need time and guidance to unravel the interaction between these different risk levels.

Under the EU AI Act, risk categories are classified as prohibited practices, high-risk AI systems and AI with specific transparency obligations.

Article 5 sets forth eight different types of prohibited AI systems, which all share a common feature: prohibited systems substantially undermine EU values by empowering organizations to manipulate or gain predictive control over human social and psychological behavior in ways deemed harmful and unacceptable to their interests. Notably, there are exceptions and conditions that will prevent AI systems from falling into this category. Member states, likewise, have flexibility in authorizing some uses of prohibited AI systems, such as real-time remote biometric identification systems, but do not have the ability to issue licenses for all technologies listed.

AI systems falling into the high-risk category are permitted but are subject to strict compliance obligations. This category intends to capture technologies that pose significant risk of harm to the health, safety or fundamental rights of an individual but are not so harmful as to be outright prohibited. Under Article 6, AI systems become high risk if they meet the conditions for an Annex 1 high-risk system or if they are expressly referred to in Annex 3. The Commission will clarify the

scope of high-risk AI systems through delegated acts and guidelines, which will include a list of practical examples.

Article 50 captures certain AI systems where it is not immediately clear to the user that they are communicating with an AI system, like chatbots, or images, audio, video or other content that have been manipulated by an AI system, like synthetic content or deepfakes. AI systems that generate or manipulate text that is published to inform the public on matters of public interest are also included. Note, this category is not mutually exclusive with high-risk systems.

In contrast, Chinese AI measures build similar compliance requirements off levels of risk but do so in different ways. The AI regulations in force with the strictest obligations target specific services and, within those services, a special category of capabilities carries more restrictions.

There is no overarching framework for AI risk classification that imposes specific legal obligations. When Chinese policy instruments have categorized AI risks generally, they've come in the form of best practices, guidelines and administrative strategy documents, such as the Ministry of Science and Technology's code of practice for AI ethics.

Under the Chinese AI regulations, risk categories are classified as prohibited practices, public opinion attributes and social mobilization capabilities.

China does not have anything approximating a prohibited list of technologies, though in practice some might be due to de facto regulatory expectations. Experts have floated, and continue to float, the idea of creating a negative licensing regime, which would create a list of prohibited AI systems requiring preapproval from authorities. But policymakers have not implemented this in formal regulation. If such an idea were picked up, it is likely Chinese authorities would be able to approve AI systems after an assessment or security filing.

While the Chinese AI regulations in force do not prescribe high-risk categories, they do recognize a special subset of services, those that have "public opinion attributes or social mobilization capabilities," and impose special restrictions on these. Public opinion attributes cover technologies that can provide a platform or channel for the expression of public opinion, while social mobilization capabilities target services that can encourage the public to engage in specific actions.

The concept originates from a 2018 regulation governing news information services, which also provides examples such as forums, blogs, chat rooms, short videos, live webcasts, information sharing services and apps. In the context of AI, these features often overlap and would include technologies like chatbots, virtual assistants and services that use content-recommender algorithms. News information services are subject to a strict licensing regime in China.

The extension of the concept to digital technologies indicates that despite sharing some similarities with the EU AI Act — in terms of a special focus on technologies that manipulate or distort information — administrative context changes how and why risk categorization plays out in practice.

Another key difference between Chinese regulations and the EU AI Act concerns who is governed and when. The EU AI Act defines distinct actors, each with corresponding obligations. These actors play a specific role in the AI supply chain — how an AI product is designed and put into market — and can occupy spaces in numerous industries, as they are sector neutral.

While some of this terminology exists in Chinese regulations, there is no one single document describing all actors. Rather, providers of specific services are covered. These services can involve multiple actors, but in practice they are interpreted to cover very specific entities.

Under the EU AI Act, the scope of actors and their key obligations are as providers, deployers, and importers and distributors of AI systems.

Providers are entities that develop AI systems or general-purpose AI models, or have them developed, and subsequently place them on the market or put them into service under their own name or trademark. Depending on their risk classifications, most obligations of the EU AI Act fall on providers. For high-risk systems these include complying with technical and organizational measures, AI quality management, post-market monitoring, corrective actions, system registration and conformity assessments.

Deployers are entities that use an AI system under their authority, excluding personal, nonprofessional uses. These entities are subject to specific deployment, data quality and monitoring obligations. For high-risk systems, this includes ensuring the real-world application of the AI system adheres to the design and operational regulations of the provider, for example input data control, monitoring and incident reporting. Notably, deployers may assume the responsibilities of providers through certain actions.

Importers are entities that are located in the EU and place AI systems carrying the name or trademark of someone established outside the EU on the market. For high-risk systems, they must verify appropriate conformity assessments have been completed, check the provider's technical documentation, and ensure the necessary CE marking/declaration of conformity and instructions for use are in place.

Distributors, on the other hand, are entities that are neither providers nor importers but otherwise make AI systems available. For high-risk systems, they must verify critical elements before marking the AI system available including the CE marking, declaration of conformity and instructions for use.

In contrast, Chinese AI regulations define main actors as providers of algorithm recommendation services, providers and technical supporters of deep synthesis services, and providers of generative AI services.

Providers of algorithm recommendation services, detailed in Articles 6-22, are entities or individuals that use recommendation algorithms to provide internet-based information services, including content. These providers must ensure content complies with laws, mark algorithm-generated content, establish internal management systems, protect users and avoid harmful influences. They must also protect vulnerable groups, as well as conduct security assessments, and file algorithms with the Cyberspace Administration of China, when applicable.

In this context, providers are entities or individuals that provide deep synthesis service. They must ensure compliance with content moderation, transparency and user protection, as well as conduct security assessments and algorithm filings, when applicable. Technical supporters, on the other hand, are entities or individuals that provide technical support for deep synthesis service. They are responsible for protecting personal information, regularly reviewing and verifying synthetic algorithm mechanisms, cooperating with authorities for remediation, and conducting security assessments and algorithm filings. These provisions are in Articles 6–22

Providers of generative AI services are entities or individuals that utilize generative AI technology to provide generative AI services, including providing such services through the provision of a programmable interface or other means. These providers must protect users, ensure content complies with laws, maintain transparency, use lawful datasets, prevent discrimination, protect privacy and conduct security assessments. They are also required to guide users, monitor AI use, mark AI content and report noncompliance to authorities, ensuring responsible and secure AI services. Generative AI measures are in Articles 5-15.

Notably, the EU AI Act's approach to general-purpose AI models is unique. Article 3(63) defines a general-purpose AI model as "an AI model, including where such an AI model is trained with a large amount of data using self-supervision at scale, that displays significant generality and is capable of competently performing a wide range of distinct tasks regardless of the way the model is placed on the market and that can be integrated into a variety of downstream systems or applications." The inclusion of special rules covering these, including those covering transparency for downstream uses and training data, technical documentation and data security testing, within the structure of the AI Act itself was a major negotiating point during the new law's drafting.

In China, regulators have also focused on general-purpose AI models. However, the level of prescription is far less defined than in the EU. Regulations currently in force may implicate the development of such models, but the bulk of current oversight comes from nonbinding technical standards, which play a different role in enforcement and oversight than in the EU.

A good example concerns the standard TC260-003 Basic Security Requirements for Generative Artificial Intelligence Services, which imposes special requirements around what training data can be used, how it can be used, the level of human oversight expected, and metrics for evaluating and recording the accuracy and quality of output data. These obligations are not explicitly found in AI regulations, although many contain a requirement to comply with standards when applicable.

Major differences are also visible in the types of compliance obligations that apply to organizations. The bulk of the differences arise from the risk classification of the service and the organization's role in providing that service.

Compliance requirements under the EU AI Act are quite dense and will be subject to further clarification through implementing and delegated acts. For instance, high-risk AI systems must meet specific technical and organizational measures throughout the life cycle of the AI system, including requirements around risk management, including implementing identification, analysis and mitigation strategies; data governance, including meeting standards for data training and model testing; technical documentation and recordkeeping, including demonstrating compliance, traceability and post-market monitoring; transparency, including designing AI systems so outputs can be comprehended; providing instructions for use; accurate and robust cybersecurity, including maintaining accuracy, security and resilience; and human oversight, including keeping humans in the loop depending on the AI system's level of risk, autonomy and context of use.

Notably, requirements in China also touch upon these topics and use similar regulatory language. For instance, a key component of Chinese AI governance is implementing comprehensive risk management systems, recordkeeping and providing instructions for downstream uses of AI products. Additionally, many requirements in Article 50 of the EU AI Act mirror those already found in Chinese AI regulations, such as transparency requirements to watermark AI generated content, label content containing deep fakes and disclose to the public that services use AI.

Despite textual similarities, it remains to be seen exactly how much actual overlap arises in enforcement and implementation. However, there are a few key structural differences worth mentioning.

Conformity assessments. The integration of the EU AI Act's requirements into the conformity assessment procedure under relevant EU product safety laws is a unique feature not currently found in China. While conformity assessments are a routine part of testing and accreditation in China, their role in AI governance is relatively small compared to other forms of regulatory review, such as algorithm filing and security assessments. Note, conformity assessments have been floated as an idea in local regulations such as Shanghai's AI regulation.

Algorithm filing. In China, services that meet particular criteria must file their algorithms with authorities. This filing must include basic information about the filing entity and the algorithm, as well as the product/services associated with the algorithm. Regulators have created a dedicated online portal for this. While there are certain registration requirements for AI systems under the EU AI Act, there is currently no central database being maintained for this purpose.

Security review. Another unique feature in China is the use of security reviews. For algorithms with public opinion attributes and social mobilization capabilities, providers must conduct security assessments. This includes multiple dimensions such as risk prevention, user protection, content moderation, model security and data security. Actors must complete a self-assessment before filing.

While regulators in the EU will conduct investigations and ask to see compliance documentation, there is nothing like a detailed security review required before launch of the service. To be sure, conformity assessments under the EU AI Act are ex ante, but their purpose is aligned with product safety and involves certification from third-party accreditation bodies, not outright licensure from a regulatory authority. Additionally, the EU AI Act's requirement to conduct a fundamental risk assessment under certain conditions diverges in scope and substance from China's self-assessment process, although in practice similar topics will be covered.

Under the EU AI Act, enforcement of most provisions — excluding those targeting general-purpose AI — is the responsibility of national competent authorities, which each member state must designate under the law. The national competent authorities will supervise the implementation of the AI Act, including obligations for high-risk AI systems and prohibitions. It is currently unclear which administrative bodies will be designated as competent authorities, but many data protection regulators have already insisted the AI Act be within their competence.

Coordinating enforcement is the AI Office, which will play a supportive role by centralizing oversight across all 27 member states and offering guidance and expertise on key topics. The AI Office, which sits within the European Commission's Directorate-General for Communications Networks, Content and Technology, will assist with joint investigations, streamline communication between national authorities, and create a governance system that includes establishing technical advisory bodies at the EU level and supporting the scientific panel of independent experts.

The office also has oversight to enforce the AI Act's provisions on general-purpose AI; a key signal that enforcement of these models should be uniform across the EU. Finally, the AI Office will assist authorities and other central-level bodies by issuing voluntary codes of conduct and support the drafting of technical standards, guidelines, and implementing and delegated acts.

Similarly, in China, there is a plethora of different administrative bodies involved in AI enforcement. These cut across different ministries and departments with core responsibilities often divided depending on the area of law and industrial sector. These government bodies collaborate on key AI issues, which include drafting and adopting AI regulations and guidance.

For instance, while the CAC remains a main regulator in AI governance, the Ministry of Information and Industry Technology, which supervises the telecommunications, internet of things and mobile apps industries, played a key role in promulgating China's existing AI regulations. Additionally, authorities in China have built off previous expertise to address risks stemming from AI. A good example is how the CAC has taken a lead role in supervising the cybersecurity and data protection aspects of AI, two areas it traditionally oversees for organizations generally.

A key difference between the EU and China is the division of competence between and within administrative units. In practice, this will affect how organizations understand and evaluate their relationships with enforcement bodies. Under the AI Act, member states have responsibilities to

designate enforcement authorities and can do so based on their own considerations. This means enforcement bodies may differ depending on the member state. Organizations will need to get used to different bodies being the competent regulator depending on where they offer services.

In China, ministries have local offices across provinces and city-level units, with these bodies being responsible for local oversight. Consequently, organizations will need to get used to all ministries that have competence over their services, regardless of where they operate.

Relatedly, in China, centralized oversight happens through the central-level offices of each ministry or state body that has political power to supervise their respective industries. These offices dictate rules of procedural administration and policy that local offices must implement. This contrast with the EU ,where the central-level bodies play a supporting role and can issue guidance but cannot outright interpret and enforce the AI Act on behalf of competent authorities, save for the specific provisions the AI Office oversees.

Understanding the key differences between the EU and China's approach to AI regulation can help organizations better prepare for compliance. Beyond these key differences in legislative approach, risk-classification frameworks, scope of covered actors, key compliance requirements and enforcement structure, additional concrete lessons for compliance will emerge as AI regulations are implemented in both jurisdictions.

Hunter Dorwart is an associate, Harry Qu, CIPP/E, is a data associate, and Tobias Bräutigam, CIPP/E, CIPM, FIP, and James Gong, CIPP/CN, CIPP/E, CIPP/US, CIPM, are partners at Bird & Bird.

Artificial Intelligence Law of the People’s Republic of China (Draft for Suggestions from Scholars)

The following document is a preliminary draft of China’s proposed AI Law that has circulated among legal scholars. The draft law specifies various scenarios in which AI developers, providers, or users are liable for misuse of AI tools. It also allows for the use of copyrighted material for model training in most cases, and provides intellectual property protections for content created with the assistance of AI technology.

This Law is enacted in accordance with the Constitution in order to promote technological innovation in artificial intelligence (AI), facilitate the healthy development of the AI industry, regulate AI product and service development, provision, and use activities, as well as their supervision and management, safeguard national security and the public interest, and protect the legitimate rights and interests of individuals and organizations.

This Law applies to AI development, provision, and use activities within the People’s Republic of China (PRC), and to the supervision and management thereof. This Law applies to those AI development, provision, and use activities outside the PRC that may affect the national security or public interests of, or the legitimate rights and interests of individuals or organizations in, the PRC.

The development of AI shall (■■) adhere to a people-centered (■■■■) approach, respect personal freedom and dignity, promote the well-being of the people, and provide assurance for the public interest, so as to guide and regulate the healthy and orderly development of the AI industry. AI developers, providers, and users shall prevent and control the possible ethical risks of AI in accordance with law.

Tracing the Roots of China's AI Regulations

In 2021 and 2022, China became the first country to implement detailed, binding regulations on some of the most common applications of artificial intelligence (AI). These rules formed the foundation of China's emerging AI governance regime, an evolving policy architecture that will affect everything from frontier AI research to the functioning of the world's second-largest economy, from large language models in Africa to autonomous vehicles in Europe.

U.S. political leaders often warn against letting China "write the rules of the road" in AI governance. But if the United States is serious about competing for global leadership in AI governance, then it needs to actually understand what it is competing against. That requires examining the nuts and bolts of both China's AI regulations and the policy process that shaped them. This paper is the second in a series breaking down China's AI regulations and pulling back the curtain on the policymaking process shaping them.

The Chinese Communist Party (CCP) and the Chinese government started that process with the 2021 rules on recommendation algorithms, an omnipresent use of the technology that is often overlooked in international AI governance discourse. Those rules imposed new obligations on companies to intervene in content recommendations, granted new rights to users being recommended content, and offered protections to gig workers subject to algorithmic scheduling. The Chinese party-state quickly followed up with a new regulation on "deep synthesis," the use of AI to generate synthetic media such as deepfakes. Those rules required AI providers to watermark AI-generated content and ensure that content does not violate people's "likeness rights" or harm the "nation's image." Together, these two regulations also created and amended China's algorithm registry, a regulatory tool that would evolve into a cornerstone of the country's AI governance regime.

Contrary to popular conception in the rest of the world, China's AI governance regime has not been created by top-down edicts from CCP leadership. President Xi Jinping and other top CCP leaders will sometimes give high-level guidance on policy priorities, but they have not been the key players when it comes to shaping China's AI regulations. Instead, those regulations have been the product of a dynamic and iterative policymaking process driven by a mix of actors from both inside and outside the Chinese party-state. Those actors include mid-level bureaucrats, academics, technologists, journalists, and policy researchers at platform tech companies. Through a mix of public advocacy, intellectual debate, technical workshopping, and bureaucratic wrangling, these actors laid the foundations for China's present and future AI regulations.

This paper traces the progression of these regulations through the "policy funnel" (see figure 1) of Chinese AI governance. For both recommendation algorithms and deep synthesis rules, the initial spark for the regulation came from long-standing CCP concerns about the creation and dissemination of online content. For the former, the rise of the algorithmically driven news app Toutiao threatened the CCP's ability to set a unified narrative and choose which stories are pushed to readers. In the case of deep synthesis, online face swap videos grabbed the attention of the Chinese public and led government regulators to consider the threat of deepfakes. Over the course of 2017–2020, these concerns made their way through China's bureaucracy. Regulators took a series of stopgap measures in specific applications, while also tasking policy analysts and government-adjacent technical organizations with exploring different regulatory interventions.

Figure 1. Exploring the "Policy Funnel" of Chinese AI Governance Click on the sections above to learn more about them. Source: Updated from Matt Sheehan, "China's AI Regulations and How

They Get Made," Carnegie Endowment for International Peace, July 10, 2023, https://carnegieendowment.org/2023/07/10/china-s-ai-regulations-and-how-they-get-made-pub-90117, figure 3.

At the same time, actors from outside the Chinese party-state were also entering the policy debates, bringing concerns about other social and economic impacts of AI to the table. In 2020, a viral magazine article detailed how food delivery workers were being exploited and put in danger by the algorithms that dictated their routes and delivery times. The piece was built on a foundation of scholarship by Chinese sociologists of labor, and it prompted a massive public outcry against the companies and technologies responsible. Within a year, that outcry translated into provisions on the role of recommendation algorithms in worker scheduling. Chinese technology companies were also actively trying to shape the coming rules, successfully nudging regulators away from using the more ominous-sounding term "deepfakes" and toward the more anodyne version, "deep synthesis."

All throughout this process, Chinese academics, journalists, companies, and even state-run media outlets were actively digesting international AI debates. They would analyze, adopt, and adapt ideas from the United States and elsewhere, covering everything from algorithmic transparency requirements in New York City to deepfake videos created by Buzzfeed News. That willingness to absorb international ideas provides a channel for engagement and even influence by U.S. actors. It also marks a real advantage for Chinese regulators who are willing to learn from and adapt ideas from abroad, regardless of whether they come from a geopolitical friend or foe.

Tracing the genesis of China's AI regulations provides a window into one of the most dynamic and consequential policy areas of today. Understanding domestic Chinese AI governance can clarify the party-state's approach to the technology and to the trade-offs inherent in governing it. Identifying the key actors, institutions, and mechanisms at play also makes it possible to begin constructing a working model of how China makes AI regulations, one that can help predict what might come next.

China is regulating AI, and the rest of the world would be wise to pay attention. Since 2021, the Chinese state has rolled out a series of targeted and binding regulations that constitute some of the first major moves by an AI power to govern one of the most transformative technologies of our time. These regulations target recommendation algorithms, deep synthesis, generative AI, and most recently facial recognition. China is now debating whether to create an overarching national AI law that could be written and rolled out in the years ahead.

This paper is the second in a series of three analyzing China's AI regulations and the forces shaping them. AI is a diffuse technology, and China has begun to introduce policies to address its application in several fields, including autonomous vehicles and medicine. This series of papers analyzes a specific subset of AI-relevant regulations: regulations drafted by the Cyberspace Administration of China (CAC) to deal with public-facing online algorithms and generative AI. The first paper in the series gave an overview of these regulations to date and analyzed their core motivations and structural characteristics. This paper focuses on China's first two significant regulations: the 2021 provisions governing recommendation algorithms and the 2022 regulation on "deep synthesis." The third and final paper in this series will examine the rich debates that shaped China's 2023 regulation on generative AI and what they signal about the road ahead for Chinese AI governance.

Recommendation algorithms are an omnipresent application of AI, powering everything from social media feeds to e-commerce platforms and navigation apps. But they are also an odd target for a country's first AI-focused regulation. Recommendation algorithms do not feature prominently in international AI governance debates because most applications of the technology are seen as relatively anodyne. The regulation on "deep synthesis" provoked similar head-scratching, as the term was unknown outside of China and sparsely used within the country.

Some provisions within the rules themselves also appeared puzzling to outsiders. Alongside some straightforward requirements on recommendation algorithms, such as granting users the right to switch off an algorithm, the recommendation algorithm regulation also contains a grab bag of seemingly unconnected requirements. It bars excessive price discrimination, bans the synthetic generation of fake news, and requires algorithm providers to protect the rights of gig workers. Together the two regulations also established and refined a new algorithm registry, a regulatory tool that has evolved into a cornerstone of China's AI governance framework.

Why did China choose to target recommendation algorithms and deep synthesis? How did these seemingly unrelated requirements get included? And what does that process reveal about the motivations and trajectory of Chinese AI governance? This paper tackles these questions by taking a "reverse engineering" approach (see box 1).

Box 1. Reverse Engineering Chinese AI Governance This paper is one of a series that attempts to better understand Chinese AI governance by reverse engineering the regulations themselves. The analysis begins with the finished product: the text of the regulations on AI and algorithms that China has already adopted. The regulations are then broken down into their component parts—the terminology, concepts, and requirements embedded in them—and those parts are traced backward through China's four-layered "policy funnel" (see figure 1). By tracing the evolution of China's AI regulations through the policy funnel, this approach aims to both deeply understand China's existing AI regulations and to help predict what new provisions may be coming around the bend. Those four layers of the policy funnel are real-world roots, Chinese President Xi Jinping and CCP ideology, the "world of ideas," and the party-state bureaucracies. Real-world roots include China's macro-level economic, political, social, and technological environment. The forces and real-world events captured in this layer create the need for new policy, but they also limit the options for regulators. The elements of the second layer—Xi's worldview and CCP ideology—act as the political and intellectual filters through which policymakers then understand these issues. To date, Xi does not appear to be actively involved in writing specific AI regulations in China; instead, he sets the direction of travel for policy actors and provides the ultimate backstop for decisions. The "world of ideas" layer is the most complex and often the most consequential. It is composed of the policy and academic debates that generate new policy proposals, as well as the corporate lobbying that attempts to steer or water down these proposals. While these public debates do not fully dictate policy, they provide the intellectual grist for the bureaucratic mill. The final layer of the policy funnel, the CCP and government bureaucracies, contains the key ministries and CCP bodies that draft and finalize regulations. While the Cyberspace Administration of China (CAC) has been the key actor driving regulation so far, other institutions may take on leading roles as AI governance evolves and expands, including the Ministry of Science and Technology or China's State Council and National People's Congress. The policy funnel described here and depicted below (see figure 2) is meant to illustrate a conceptual framework for understanding Chinese AI policymaking. In reality, the specific events leading to each new regulation are unique, and the process is often ad hoc. The layers depicted here are porous, and regulations do not proceed through them in a purely linear fashion. Instead, they often pinball forward and backward through these layers, getting shaped and reshaped by different combinations of academics, bureaucrats, public opinion, and CCP ideology. But there is often a pattern to this process, and this paper seeks to provide a grounded and granular look into the underlying forces and key players in that pattern. For a more detailed examination of these four layers and prominent actors within each layer, see the first paper in this series.

The following section highlights the key provisions and terminology in these regulations. The paper then traces the history and evolution of these ideas over the past seven years, revealing a dynamic policymaking process. The initial impetus for these regulations came from longstanding concerns of the CCP and the Chinese government over online content. Those concerns about how online content is created and disseminated represent the through line in all of China's early regulations on algorithms and AI.

But as these regulatory proposals gathered momentum, a range of other actors from across Chinese society—academics, journalists, tech companies, and Chinese citizens—reshaped the regulations. Public outcry and intellectual ferment from these groups led to changes in the scope and requirements of the regulations, with new provisions added to address their concerns. Many of these policy actors were absorbing and engaging with ideas about AI governance that were developing in the United States and other countries, constituting an informal channel for both influence and debate.

Tracing the development of these regulations reveals a dynamic policymaking process that weaves together CCP concerns, bureaucratic politics, public discourse, corporate lobbying, academic analysis, and international AI governance debates. Whether one wants to compete against, engage with, or simply understand China's AI governance initiatives, this wider policymaking process must be examined as well.

China's rules on recommendation algorithms are as much about internet regulation as AI. They were initially conceived to target not AI technology in itself but a specific technology application—online recommendation algorithms—that happened to be powered by AI and machine learning. In doing so, they incorporated many components from earlier Chinese regulations on online content, some of which are described below.

The central goal of the rules is to regulate the way in which algorithms affect the promotion and dissemination of content online. Recommendation algorithms presented a new challenge to the Chinese state's long-standing internet controls, and so the CCP singled them out for regulation. In fact, the term "artificial intelligence" (■■■■) does not even appear in the regulation itself.

And yet, the regulation is inextricably linked to AI. The rapid adoption of machine learning algorithms by leading Chinese internet platforms is what created the need for the regulation. And it created precedents and regulatory tools that have become key pieces of China's subsequent AI regulations, laying a foundation for the country's wider governance regime.

Most of the requirements in the algorithmic recommendation regulation can be grouped into three categories: content requirements for service providers, user rights, and business practices (see box 2.). The provisions pertaining to online content reflect the original motivation for the regulation, while those around user rights and business practices appear to have been added on during the policy incubation process.

Box 2. Provisions on the Management of Algorithmic Recommendations in Internet Information Services Draft: August 27, 2021. Signed: December 31, 2021. Implemented: March 1, 2022. Signatories: Cyberspace Administration of China (CAC), Ministry of Industry and Information Technology, Ministry of Public Security, State Administration for Market Regulation. Draft regulation issued solely by the CAC. Available in Original Chinese, English Translation. Key Requirements Recommendation algorithm service providers must: "actively transmit positive energy" and not "disrupt economic and social order"

create mechanisms for manual intervention in "top searches" and "hot topics" so that promoted content conforms to "mainstream value orientations" Users have new rights to: turn off algorithmic recommendations for an app or website

obtain an explanation if an algorithm has a major impact on a users' rights Businesses must: not use algorithms for monopolistic or unfair business practices

protect the rights to fair compensation and adequate rest for workers whose schedules are set by algorithms Finally, the regulations creates a new regulatory tool, the algorithm registry (■■■■■■). Algorithms that can affect public opinion or "have social mobilization capabilities" must provide

basic information on the algorithm and conduct algorithm security self-assessments.

China's second major regulation targets "deep synthesis" (■■■■), a relatively new term that refers to the use of deep learning to synthesize or generate content—what is today called generative AI. The regulation casts a wide net, including the use of algorithms to generate or alter text, images, video, audio, and other virtual renderings, such as the metaverse (see box 3). The term deep synthesis and the wide regulatory scope are an expansion from narrower CCP concerns about deepfakes that sparked the development of the regulation. This paper traces that evolution, revealing both the roots of, and corporate influences on, policymaking.

The deep synthesis regulation aims to guard against the creation of misleading or politically sensitive content, while also protecting the data and privacy rights of Chinese citizens from those using or developing AI. It tries to address these problems at three levels: broad ideological guidance, specific prohibitions and requirements, and the watermarking of synthetically generated content.

Box 3. Provisions on the Administration of Deep Synthesis Internet Information Services Draft: January 28, 2022. Signed: November 25, 2022. Implemented: January 10, 2023. Signatories: Cyberspace Administration of China (CAC), Ministry of Industry and Information Technology, Ministry of Public Security. Draft regulation issued solely by the CAC. Available in Original Chinese, English Translation. Key Requirements Broad ideological guidance: "respect social mores and ethics"

"adhere to the correct political direction, public opinion orientation and values trends" Specific prohibitions and requirements for deep synthesis service providers: Do not produce, publish, or transmit "fake news"

Conduct in-house or third-party security assessments if either editing biometric features or producing content that "might involve national security" or "the nation's image" Requires organizations providing "technical support" for deep synthesis service providers to conduct security assessments and make filings in the algorithm registry Watermarking requirements for service providers: All generated or edited content must contain a digital watermark for identification

If generated content "might cause confusion or mislead the public," providers must include a "conspicuous label in a reasonable position"

The signing of the deep synthesis regulation on November 25, 2022, marked the end of the first phase of China's regulations on AI and algorithms. The regulations on algorithmic recommendations and deep synthesis had developed in parallel, moving through China's policy funnel and party-state bureaucracy together. These two regulations would lay the foundations of China's emerging AI governance regime, creating new precedents and regulatory tools that would be embedded in subsequent regulations on generative AI and facial recognition. (The term "party-state" refers to the combination of institutions comprising both the Chinese government and the CCP.)

The motivations driving and forces shaping these two regulations are best revealed by looking backward. Tracing the regulations' terminology and provisions back to their origins reveals how CCP concerns, social forces, and intellectual inputs all shaped these regulations.

The regulations on recommendation algorithms and deep synthesis emerged separately, but they developed in parallel and were soon bureaucratically woven together, moving through the policy funnel in near unison.

In both cases, the initial spark for the regulation came from party-state concerns about how a new technology application was affecting the creation and dissemination of online content. This core concern provided the impetus to begin the policy process, and it would go on to become the backbone of the final regulation. Understanding how that content-focused backbone came to be offers insight into the CCP's core motivations and process for AI regulation. The second half of this paper explores how other nongovernmental forces expanded the scope of the regulations to include other social and economic concerns related to the technology.

The first clue to the origins of the algorithmic recommendation regulation comes in the title itself. "Algorithmic recommendation" (■■■■) is a sufficiently unique term that it is possible to identify when and in what context it first surfaced in Chinese state media. Those early media discussions of recommendation algorithms reveal how a backlash against algorithm-driven news apps planted the seeds that would grow into one of the world's first major binding regulations on the technology.

The CCP uses China's state- and party-run media outlets as a venue to communicate its priorities and explain its policies. When prominent outlets highlight a problem or promote an industry, that can act as a signal to actors both inside and outside the party. That signal can lead lower-level regulators to turn their attention to that problem, and it can lead businesses to double down on (or withdraw from) that industry. How often a particular term is mentioned in state-run newspapers can serve as a rough proxy for CCP interest in a topic, and the arguments found in those articles are a window into how the party is thinking about that issue. In some cases this media coverage is forward-looking, laying the political and intellectual groundwork for future action. In other cases it is backward-looking, explaining to the public and party members why a certain action was taken. In the case of state media discussion of recommendation algorithms, these early discussions were forward-looking, signaling and sharpening the ideas that would go on to drive policy.

Between 2014 and 2016, the terms "recommendation algorithm" and "algorithmic recommendation" began appearing in some state-run media articles about business and technology. These articles were relatively rare, one every couple of months, and the use of the terms was marginal to the central focus of the article. In late 2016 and early 2017, that began to change. The number of articles referencing algorithmic recommendation increased, and many of them discussed the risks of this new technology. Several of them specifically singled out one popular news app for attack: Jinri Toutiao (■■■■, Today's Headlines).

Toutiao (pronounced roughly: "toe-teeow"), as the app is often called, is an algorithm-driven news aggregator and content platform that delivers a personalized stream of news and other content to users. It was created in 2012 by ByteDance (the company that would go on to create TikTok), and by 2016 it had become China's most popular news app. But just as Toutiao was surging in popularity, ByteDance founder Zhang Yiming stumbled into trouble with Chinese authorities.

In December 2016, Zhang gave a wide-ranging interview to leading business magazine Caijing, discussing everything from his personal life philosophy to how Toutiao dealt with vulgar content. In the conversation, Zhang portrayed Toutiao as a neutral platform whose algorithm's sole goal was to match users with content that they were interested in, regardless of whether that content was high-minded or "vulgar." He explicitly disavowed the idea that his company should try to guide users or inculcate any sort of "values" in them. When asked by the interviewer, he rejected the idea that Toutiao was a media company or the suggestion that it needed an "editor-in-chief."

That laissez-faire approach to content moderation reflected Zhang's own beliefs and showcased how empowered China's leading technology entrepreneurs felt at the time. They were rich, publicly idolized, and often supported by a government that was going all out for "innovation." But the approach also ran counter to a growing push for ideological conformity in Chinese media. Earlier in 2016, Xi had demanded that official media outlets serve the party, saying that the surname of these outlets must be "the party" (■■■■■■). In China, nonstate media outlets are restricted from publishing content on a wide variety of public events and government actions.

Zhang likely believed his algorithm-driven platform to fall outside of these structures, but the People's Daily—the flagship newspaper of the CCP—quickly set him straight. On December 23, 2016, the People's Daily ran an editorial that never mentioned Toutiao by name but directly rebutted Zhang's statements. It emphasized the need for proactive "guidance" of public opinion to diffuse social conflicts. It said, "In an era dominated by algorithms, there is an even greater need for 'editors-in-chief' who can guard, guide, and lead."

Over the spring and summer of 2017, criticism of Toutiao in state-run media outlets ramped up, with articles covering everything from the app's alleged copyright infringement to the proliferation of vulgar content. That criticism came to a head in September 2017. From September 18 to 20, the website of the People's Daily ran a series of three articles on consecutive days attacking recommendation algorithms, Toutiao in particular. Each of the three articles criticized recommendation algorithms from a slightly different angle.

The first article, titled "Algorithms Cannot Be Allowed to Decide Content," claimed that vulgar content was spreading on these platforms because algorithms had replaced human editors in deciding what content was featured. It once again invoked the need for algorithm-driven platforms to have an "editor-in-chief." The second article criticized recommendation algorithms for trapping readers in an "information cocoon" where they are not exposed to other viewpoints. The third article attacked Toutiao on copyright issues and argued that platforms driven by recommendation algorithms could harm creators and end up producing "the opposite of innovation." It concluded by calling on companies to exercise self-discipline and on the government to improve relevant regulations. With these three articles, the CCP signaled its intense displeasure with the way recommendation algorithms were being deployed in China's media landscape (see figure 3).

In the following months, both official media and regulators began turning up the heat on Toutiao and other algorithm-driven news apps. In December 2017, the Beijing branch of the CAC called managers of Toutiao and another media app to its office to lecture them for spreading vulgar content and operating an online news platform without a permit. Two days later, Toutiao shut down the app's "Society Channel" (a feed for social news) and replaced it with the "New Era Channel" as the default feed for users. The "new era" (■■■) is a political catchphrase popularized under Xi, and the channel promoted news in line with the CCP's narratives. In a public announcement, Toutiao management said the move was made to better "promote the spirit of the [CCP's recent] 19th Party Congress and report on the construction of the new era."

ByteDance tried to fend off further criticism by holding an event in Beijing in January 2018 on "making algorithms open and transparent." At the event, Toutiao's senior algorithm architect explained the fundamentals of how the app worked. But the blows kept coming. In late March, a prominent business show on the state-run broadcaster CCTV ran an exposé accusing Toutiao of violating the law by presenting users with shady medical advertisements. In April, both the CAC and the National Radio and Television Administration (NRTA) separately called in representatives from Toutiao and a short video app, Kuaishou, for more criticism. A few days later, the CAC ordered app stores to block all downloads of Toutiao for three weeks and of a few other apps for shorter periods of time. Not to be outdone, the NRTA then permanently banned another ByteDance app, Neihan Duanzi, that was used for sharing jokes and memes. After the block on downloads, Zhang publicly apologized and pledged to do more to ensure recommendation algorithms would be better aligned with national values.

China's regulatory apparatus had both identified recommendation algorithms as a target and taken some blunt measures to discipline the companies. ByteDance survived the first regulatory onslaught, but just barely. The struggles of Toutiao following Zhang's 2016 interview also provided a smaller-scale preview of the punishment that Alibaba-affiliated Ant Group would face—cancellation of its initial public offering—after a defiant speech by founder Jack Ma in late 2020.

What drove the CCP to single out Toutiao and recommendation algorithms? For starters, there were genuine issues with the way the company and other algorithmically driven content platforms were operating. Aside from copyright infringement, deceptive medical advertisements were a real problem. In 2016 a college student died after undergoing shady medical treatments he had found through an ad on Baidu's search engine, leading to a popular uproar and a government crackdown. There had also been persistent issues with vulgar and borderline pornographic content on Toutiao and other news apps. At the time, an employee of a different algorithm-driven app relayed to this author their team's strategy for user retention: whenever the app's monthly users would begin to dip, they would simply inject more sexual content into users' feeds and the numbers would go back up.

But for all the potential criticisms of Toutiao and recommendation algorithms, one would prove to be core driver of CCP regulation going forward: the role of algorithms in deciding what content gets pushed to Chinese people. Since it took power in 1949, the CCP has largely been able to decide which news stories are featured prominently in the media. This power faced challenges with the rise of nonstate media and then websites, but in both cases the party could always hold a person responsible for the decision to feature a particular story. Toutiao changed that by handing over decisionmaking power to an algorithm.

The party could always hold a person responsible for the decision to feature a particular story. Toutiao changed that by handing over decisionmaking power to an algorithm.

The power of the app came from the fact that the content it pushed to each user reflected the user's interests, not the CCP's. The app further fractured and segmented China's information environment, undermining the party's goal of a unified public narrative of events. Zhang's declaration that Toutiao did not need an editor-in-chief and would not push any particular "values" to users crystallized the conflict. This was unacceptable and required a response, one that would develop over the next few years.

In the year following the initial crackdown on recommendation algorithms, Chinese regulators developed a new set of concerns that would eventually evolve into the 2022 regulation on deep synthesis. But to trace that regulation back to its origins, the first step is changing the terminology. "Deep synthesis" (■■■■) is a term that only entered the Chinese AI lexicon in 2020. To follow this regulatory cycle back before 2020, a different term is needed for the specific technology application that first attracted the authorities' interest: deepfakes (■■■■).

"Deepfake" is a portmanteau of "deep learning" and "fake" that refers to highly realistic images or sounds created using AI. In late 2017, the term entered the English lexicon from the username of a poster on Reddit, who used AI to swap the faces of female celebrities onto the bodies of performers in adult films. The action sparked a public outcry and greater attention in U.S. technology and policy circles to the threats deepfakes posed for individual privacy and mass disinformation.

In China, deepfakes entered mainstream discourse in a similar way: an internet user swapped the faces of female celebrities in videos. In February 2019 an internet user who went by the online moniker "Face Change Bro" (■■■) uploaded a short clip from the 1994 television show Legend of the Condor Heroes in which he swapped the face of actress Athena Chu out for another actress, Yang Mi. The video immediately went viral on the Chinese internet, sparking debates among users about the implications for copyright protections, actors' incomes, and AI ethics. It also led to a swift backlash. Within two days Face Change Bro had removed the video and apologized, calling the outcry "a warning to all of us." He added, "AI face-swapping should be used in the correct way. Everyone should respect copyright and likeness rights, and devote attention to the technology itself."

In the following months, Chinese public discussion of deepfakes ramped up in both the press and online. Notably, almost every early article in Chinese state media that discussed deepfakes cited research emerging from the United States. The very first article in Chinese state media to mention deepfakes, an August 2018 piece in Science and Technology Daily, reported on a project by the U.S. Defense Advanced Research Projects Agency (DARPA) for identifying fake videos. A February 2019 piece in Beijing Youth Daily discussing the Face Change Bro scandal cited a video explainer from the Carnegie Endowment for International Peace titled “How Should Countries Tackle Deepfakes?” Other articles cited technical research from New York University, reports by other U.S. think tanks, and even a deepfake video of former president Barack Obama created by Buzzfeed and actor-director Jordan Peele. Despite the two countries’ different political systems and growing tensions between China and the United States at the time, Chinese journalists, scholars, and analysts frequently looked to their peers in the United States for analysis of technology trends and regulatory responses.

This early attention to deepfakes was quickly translated into law. In the spring of 2019, a subcommittee of China’s legislature was debating amendments to the civil code. In April, two months after the viral face-swap video, it announced amendments to the section on “likeness rights.” These amendments explicitly banned the infringement of a person’s likeness rights by “using information technology fabrication.” The subcommittee said that “some ministries” had raised the issue of deepfakes and asked that the amendments address problems such as “face-swapping.” The amendments were formally adopted in June.

The Face Change Bro incident illustrates the Chinese party-state’s attention to online discourse and its rapid regulatory response time. But the amendments to the civil code were just a stopgap solution to a more complex problem. Over the following two years, the Chinese state marshaled its own bureaucracy and government-adjacent technical organizations to build a more robust regulatory response to the threat of deepfakes.

By early 2019, the CCP had identified both recommendation algorithms and deepfakes as issues in need of attention. While recommendation algorithms posed a threat to China’s online content environment, deepfakes first arose as a threat to citizen privacy and “likeness rights.” The party-state had taken some initial measures—corporate punishments and tweaks to civil code—but more substantial regulation was needed to achieve its objectives.

Over the course of 2019 and 2020, the Chinese party-state began laying the groundwork for those regulations by galvanizing an array of government-adjacent technical organizations. These organizations form an ecosystem that is a key part of how the Chinese state acquires knowledge and generates touchpoints with industry that help it to understand emerging technologies. Exploring how that government-adjacent technical ecosystem operates gives new windows into how Chinese AI governance is incubated and eventually enforced. Identifying key organizations and committees can also help with forecasting future regulations.

The initial nudge that got these organizations working on recommendation algorithms and deepfakes came from the CAC. Founded in 2014, it quickly rose to a prominent role in Chinese technology policy under its charismatic (and later purged) first leader. The CAC’s central mission has always been regulating online content, but through a series of bureaucratic reorganizations and technocratic land grabs it has expanded its remit to include cybersecurity, data policy, and a wide array of issues touched by the internet. The CAC is “one organization with two nameplates,” meaning it sits within the CCP bureaucracy but also has an identity within the government. Its exact institutional structure and legal identity are very complex and are the subject of close examination by leading scholars. For the time being, it is enough to know that the CAC is a powerful, content-focused internet regulator, and it has largely driven the first wave of China’s AI regulations.

To begin that process, the CAC engaged the Artificial Intelligence Industry Alliance (AIIA). The AIIA was created in 2017 at the behest of several ministries and is de-facto led by the China Academy of

Information and Communications Technology (CAICT), an influential technical think tank underneath the Ministry of Industry and Information Technology (MIIT). The AIIA acts as a service provider and coordinator between the alliance's many members: private companies, state-owned enterprises, academic institutions, and the government. In early 2019, the CAC charged the AIIA with creating a new "Cyberspace Technical Committee" (■■■■■■■) that would investigate issues related to AI's role in online content. The new committee was led by the CAICT, in partnership with Zhejiang University and the Institute of Information of the Chinese Academy of Sciences.

Though the operations of the Cyberspace Technical Committee were not highly publicized, a slide deck uploaded to the website of the AIIA in 2020 recapped the group's work over its first year. That slide deck gives a snapshot of how the CAC first began trying to understand and tackle recommendation algorithms and deepfakes.

The committee chose three technical areas to focus on in its first year: recommendation algorithms, deepfakes, and "content security" (■■■■), the latter largely focused on using AI to detect illegal content. For these AI applications, the committee began developing standards, certifications, technical tools for governance, and "industry self-discipline pledges." To do that, it convened technical seminars and workshops that included academia, think tanks, and leading Chinese platform companies.

For recommendation algorithms, the committee focused on industry self-discipline and technical requirements. It drafted an "industry self-discipline pledge" in which companies promised to protect user data, to respect users' right to choose, and to "spread a positive and healthy online culture." The committee also crafted a set of technical requirements for recommendation algorithms in both the news and finance. By 2020, those draft requirements had been turned into an official evaluation and certification offered through the AIIA for "news information domain trustworthy intelligent recommendation systems." The AIIA and CAICT offer evaluation and certification services for many AI applications, ranging from deep learning software frameworks to turnstiles that use facial recognition. Obtaining these types of evaluations and certifications does not offer full legal protection to companies, but it gives them a sense of where regulation is heading, as well as a selling point when marketing their AI-driven products and services to government procurers. Together, these pledges, standards, and certifications were a rough sketch of where the rules for recommendation algorithms were headed.

For deepfakes, the committee focused on gathering information and developing technical tools for detection. In January 2020, it began collecting examples of deepfake technologies. It requested that companies submit descriptions of software for either creating or detecting deepfakes, as well as any relevant recommendations from the companies. In parallel, the committee began drafting its own evaluation methods for deepfake detection technology, and it built a "prototype system for fabricated video management," though details on this are unclear.

The third area, "content security," focused on AI's role in detecting illegal or harmful content. This work was oriented toward developing AI tools and technical standards to detect undesirable content, not toward developing rules or regulations governing it. As such, this paper does not examine it further.

The committee's work is an example of what the Chinese government calls "collaborative governance" (■■■■). This multistakeholder approach, deployed over the past decade in governing the sharing economy, would have the government avoid simply dictating policy and instead "collaborate with industry associations, platform companies, resource providers, and consumers to address issues that arise" from new technologies and business models. This approach continues today, but the balance between the party-state regulators and other actors is always in flux. From 2015 to 2019, the CAC left much in the hands of companies to self-regulate, but from 2020 to 2022 the administration took a more hands-on—and in some cases, uncompromising—approach to technology regulation.

Exactly what role did the Cyberspace Technical Committee play in the development of China's future regulations on recommendation algorithms and deep synthesis? It offered a chance for the CAC to familiarize itself with the technologies and workshop potential technical and regulatory interventions (see figure 4). For outside observers, the committee's work gave a preview of where regulation was headed and a window into the development of government-adjacent technical certifications. Understanding these committees and technical certifications will offer insight into how the Chinese party-state is likely to approach issues of AI security and safety in the future.

Particularly noteworthy is a new group founded by the AIIA in September 2023, the Safety and Security Governance Committee. (In Chinese, the word ■■ can mean either "safety" or "security" depending on context. In this case, it appears to reference both concepts.) The new committee has been tasked with supporting policy development, building basic infrastructure, and collaboratively exploring issues related to AI safety and security. The new committee is headed by a director at the CAICT, Shi Lin, who also led the earlier Cyberspace Technical Committee covering recommendation algorithms and deepfakes. In December 2023, the committee previewed a new and updated AI risk management framework, which it aims to build on in 2024. Following the work of these types of groups and committees, as well as the development of new standards and certifications for "trustworthy large models," can offer a grounded sense of where regulation in China might be headed.

China's technology regulations tend to be highly iterative. Rules are often rushed out in response to a new issue, and when these prove inadequate the regulators will simply roll out another regulation that clarifies, tweaks, or adds to the requirements in the initial rules. This can appear haphazard, but it is often an intentional strategy that allows regulators to evaluate the effectiveness of different interventions and then double down on the ones that work best. Between 2019 and 2023 this pattern played out in China's AI regulations.

In late 2019, when it was overseeing the AIIA's work on deepfakes and recommendation algorithms, the CAC also released two regulations that addressed aspects of these problems. In November 2019, the CAC released the Provisions on the Management of Online Audiovisual Services, which targeted deepfakes. China's Ministry of Culture and Tourism and the NRTA co-signed the regulation. A month later came the Provisions on the Governance of the Online Information Content Ecosystem, a sweeping, content-focused regulation that also addressed online recommendation systems. These two sets of rules were the CAC's first stab at formally regulating these technologies.

The audiovisual regulation banned the use of deep learning to create or spread "fake news," and it required conspicuous labels to be applied when deep learning is used to create "false audiovisual information." The content ecosystem regulation was much broader and largely focused on sorting content into broad categories: "positive" (encouraged), "negative" (discouraged), and "illegal" (banned). It then applied these categories to the use of personalization algorithms, requiring that providers use those algorithms to "actively present" positive information, to "prevent and resist" negative content, and to not display illegal content. The regulation also required that providers "optimize" their recommendation systems and use manual interventions so that lists of "hot topics" and "top searches" display politically desirable content. These two 2019 regulations gave a relatively crude preview of many requirements that would eventually make their way into the more robust regulations of 2021 and 2022.

In the case of the audiovisual regulations, the provisions were likely a slapdash response to the public outcry over a popular new face-swapping app, Zao. The app debuted in August 2019, and within days it was the most-downloaded app in Apple's Chinese app store. In the weeks that followed, online discussion and news articles criticized Zao and discussed the risks to user privacy and misinformation resulting from deepfakes. The CAC's deputy director fielded questions from journalists about regulating the app. The audiovisual regulation was released two months later. One

analysis reposted on its website praised the CAC for “conforming to the urgent demands of the broad masses of people” in its adherence to the principle of “urgent needs come first” in regulation. But what the regulation gained in speed, it lost in specificity. The scope of applications covered was vague, perhaps intentionally so, as were many of its requirements, such as those for labeling AI-generated content.

In the content ecosystem regulations, the rules addressing recommendation algorithms were not the central focus. The provisions targeted one related concern—what content gets promoted on homepages and on lists of “hot topics”—but they largely focused on enumerating types of undesirable content. The regulation requires platforms to respect users “right to make selections,” but it didn’t clarify what that right meant in practice. Getting clarity would require the creation of a regulation that was wholly focused on recommendation algorithms.

One month after the content ecosystem regulation was signed, COVID-19 swept through the city of Wuhan and then China and the world. The massive upheaval caused by the virus dramatically slowed down regulatory action on these topics during the first half of 2020. But by the end of the year, the highest levels of the party turned their attention back to these topics.

In December 2020, the powerful Central Committee of the CCP issued a document intended to guide much of the legislation and regulation over the next five years: the Implementation Outline for the Establishment of a Rule of Law-Based Society (2020–2025). The document contained some overarching guidance and a massive laundry list of specific issues. Buried within a section on internet law issues that needed to be addressed were two AI applications: recommendation algorithms and deepfakes.

The inclusion of these two marked a success for the CAC’s policy entrepreneurship. Between 2017 and 2020 the CAC had successfully established itself as the go-to regulator for both topics, and it had managed to get them elevated to the level of the Central Committee. It also meant that more work was needed to hammer down these topics. Over the course of 2021 and 2022, the CAC would lean on Chinese academics, think tanks, and companies to get that work done.

Between August 2021 and January 2022, the CAC went on a spree of releasing both drafts and finalized regulations on algorithms and AI. In August and December of 2021, the CAC released a draft and then final version of its regulation on “algorithmic recommendation.” The final version was co-signed by the MIIT, Ministry of Public Security (MPS), and State Administration for Market Regulation (SAMR). Between these dates, the CAC and eight other CCP and government departments also jointly released guidance on “strengthening overall governance” of online algorithms. (This document provides high-level principles for regulation of algorithms and will not be covered in this paper.) In January 2022, the CAC released a draft of the new regulation on “deep synthesis,” which it finalized ten months later in November, in conjunction with the MIIT and MPS.

How did the policy ideas in these regulations move from the amorphous intellectual ferment of Chinese policy discourse—the “world of ideas” layer of the policy funnel—into concrete requirements in regulations?

Most of this process happened behind closed doors, with limited windows for outsiders to peek inside. But some characteristics of the processes at play are apparent. The CAC has both power and experience with controlling online content but more limited in-house expertise when it comes to the technical and legal specifics of regulating AI. So it relies heavily on extensive external consultations with legal and technical experts from across academia, think tanks, and industry.

These consultations occur over various stages of the drafting process and range from informal exchanges to formal workshops and specially commissioned reports. A single regulation may see over one hundred different experts consulted. Some of those outside advisers will simply offer

general ideas or feedback, while others may be asked to take the pen and draft pieces of the regulation. In some cases, the regulating organization will create a formal “expert drafting small group” (“■■■■■■” or “■■”), while in other cases the process is conducted more informally. The membership of these advisory groups is not always publicized, but some scholars will publicly tout their role as advisers. The CAC also occasionally announces the appointment of groups of legal advisers to the organization as a whole.

Which scholars and institutions are consulted will depend partly on which bureau leads the drafting of a regulation. The CAC is home to over a dozen bureaus, each with a different specialization. The head of a given bureau may have a preexisting network of experts they rely on or a tendency to lean more heavily on technical or legal experts. While the final text of regulations will not list which bureau led the drafting process, the leadership can sometimes be inferred based on which bureau is designated to receive feedback on the draft regulation. Both the recommendation algorithm and deep synthesis regulations appear to have been drafted by the Bureau of Cyber Rule of Law (■■■■■) of the CAC. The draft generative AI regulation appears to have been drafted by the Cybersecurity Coordination Bureau (■■■■■■■).

Once a regulation emerges from a bureau at the CAC, it goes through several layers of internal bureaucratic consultation. This process is not identical for all regulations, but it appears to have been followed for the two regulations covered in this paper. First, a regulation is circulated to the other bureaus within the CAC for feedback, then to other relevant ministries and agencies across the government. Once it has gone through these two rounds of internal consultation, a draft is published to solicit public feedback, including from outside experts, companies, and regular citizens. A revised version is then circulated for feedback once more among relevant ministries and agencies before a final draft is published. While the first drafts are often issued solely by the CAC, the final versions are sometimes co-issued by other ministries or agencies who oversee aspects of the relevant topics (see figure 5).

Previous sections of this paper have laid out the roots of these regulations in CCP and state concerns over online content and shown how those concerns made their way through party and state apparatuses. This section looks beyond those content-centric concerns and focuses on key additions to these regulations during the process of drafting and revising them. Tracing the origins of these changes shows how the regulations were shaped by nongovernmental actors from throughout China’s broader social, economic, and intellectual ecosystem.

The regulation on recommendation algorithms saw some of the most significant additions between the draft and the final version. A provision barring the use of algorithms for unfair competition or monopolistic practices was added. This is likely directly attributable to the addition of the SAMR as a co-signatory of the final version. The SAMR, China’s top antitrust authority, took a series of anti-monopoly and unfair competition actions against China’s leading platform companies in 2020–2022, culminating in changes to the Anti-Monopoly Law and Anti-Unfair Competition Law. This campaign followed Xi’s pledge to build “common prosperity,” a broad push to narrow the wealth gap that helped fuel a range of crackdowns, several led by the SAMR, on large corporations. The provision on algorithms’ role in reinforcing monopolies was most likely added at its behest, a manifestation of the kind of bureaucratic bargaining that takes place as the CAC seeks buy-in and co-signatories among other agencies.

More mysterious was the addition of Article 20 of the regulation. That article requires recommendation algorithm providers that offer “work dispatch services” (■■■■■■) to protect the rights and interests of workers, including their legal rights to compensation and rest. It also demands that providers improve any algorithms used to assign orders or to dictate workers’ salaries and scheduling. This provision stands apart from the rest of the regulation in its focus on worker rights. It has no clear connection to the core concerns of the CAC or its fellow regulators. So where did this issue come from, and how did it make its way into the regulation? The following section will trace its origins.

The most notable changes in the deep synthesis regulation are its title and scope. The roots of the regulation are found in the party-state's concern over deepfakes (■■■■). Prior documents, including the December 2020 directive from the CCP Central Committee, explicitly used the term "deepfakes" in its call for regulation on the topic. But somewhere between that December 2020 directive and the January 2022 release of the draft regulation, the CAC changed both the terminology and the scope of the regulation to "deep synthesis" (■■■■). In the regulation, deep synthesis is defined as the use of "generative synthesis algorithms" such as deep learning to generate or edit text, images, audio, video, or other "virtual scenes." The term "deepfake" is not in the regulation itself, only appearing in an accompanying article explaining that the regulation responds to the Central Committee's call for addressing deepfakes. Why did the CAC change both the key term and the overall scope of the regulation? This will be answered in a subsequent section.

A final, notable addition to both the recommendation algorithm and deep synthesis regulations is the creation of the algorithm registry (■■■■■■). Both regulations require that providers whose algorithms have "public opinion properties or capacity for social mobilization" submit a filing on their algorithm within ten business days of operation. What that filing should contain is left vague in the text of the original regulation, but the algorithm registry has gone on to become a key component of China's AI governance regime. The origins and mechanics of the algorithm registry are explored further below.

To understand where these additional regulatory requirements came from, this analysis will turn back the clock to 2020 and widen the field of vision beyond the machinations of China's official bureaucracies.

There was a clear inciting incident behind the algorithmic recommendation regulation's protections for workers: the public outcry in 2020 over the plight of food-delivery workers whose schedules and routes are set by algorithms. That outcry followed the publication of a magazine exposé on the struggles of these workers that went viral, an article that was in turn built on the work of Chinese and American sociologists and anthropologists. Tracing the progression of these ideas from academic papers to a magazine article to public outcry to the party-state's response showcases alternative routes to policy influence in China.

Over the past decade, the restaurant industry in Chinese cities has been transformed by food delivery apps. By 2020, China's food-delivery employed 6 million delivery drivers who were zooming across cities on electric bikes and scooters. Two companies—Meituan and Ele.me—dominate that industry, combining to control 98 percent of the market.

Both companies rely on machine learning algorithms to assign deliveries, set expected wait times, decide delivery routes, and determine driver compensation for each trip. Amid the fierce competition between them, those algorithms were often tuned to reduce both delivery times and driver compensation to an absolute minimum. Workers found themselves scrambling to meet these new demands, while struggling to make the deliveries needed to earn a livable income. Many of these workers are migrants from the Chinese countryside who lack both proper contracts and any social safety net.

Outrage about the treatment of these workers erupted in September 2020 following the publication of a long-form magazine exposé titled "Delivery Workers, Trapped in the System" (original Chinese). The piece was published by Renwu ("People") magazine, an outlet under a state-owned publishing house. But Renwu is known for some edgy publication decisions, including an interview with an early coronavirus pandemic whistleblower that was later censored. The article about delivery workers detailed how the unreasonable delivery times and low compensation set by Ele.me and Meituan's algorithms forced delivery drivers to run red lights and ride against traffic in their

desperation to make a living. It described in gory detail the way these pressures led to the death of drivers:

One afternoon this spring, Wei Lai and another rider wearing the same color clothing as him were stopped at an intersection waiting for the light to change. There were only a few seconds before the light turned green, but the other rider darted into the intersection. At the same time, a fast-moving bus sped through and the rider and his scooter went flying. He died on the spot. Wei Lai said he saw the badly mangled body in the middle of the road, but he did not stop at all. His own order was late. At that time, another order came in and the familiar female voice of the app's delivery assistant chimed in: "Order! From 'Point A' to 'Point B,' please respond after the beep to accept.

The article also explored the technical and sociological underpinnings of the industry, drawing extensively from the work of Chinese sociologists of labor, including Sun Ping from the Chinese Academy of Sciences and Zheng Guanghai from Central China Normal University. These researchers had spent years surveying delivery workers and developing sociological concepts to describe the interplay of algorithms, platforms, workers, and consumers. Those academics in turn frequently cited and built on the work of American academics like Nick Seaver, who developed the concept of "algorithms as culture." Weaving all of these strands together, the Renwu article was a remarkably in-depth yet accessible investigation into the role algorithms played in the exploitation of gig workers.

Despite its length and heavy subject matter, the article quickly went viral, dominating social media feeds and sparking intense public backlash against Ele.me and Meituan. The companies attempted to placate these concerns with minor changes to their apps. Ele.me gave users the option to "wait five more minutes" for their delivery, while Meituan added eight minutes of "flexible delivery time." But many commentators dismissed these as cosmetic changes that just shifted the burden onto consumers. Chinese state media piled on the criticism. One leading television anchor demanded the platforms treat their employees as "people, not machines" and called for increased government oversight of the industry. Newspapers associated with different parts of the party-state bureaucracy proposed different solutions, ranging from technical interventions in the algorithms to resetting the contracting relationship between firms and workers.

The government's regulatory response came the following summer. On July 7, 2021, the State Council issued a statement on protecting the rights of workers in the platform economy, which included a call for delivery platforms to improve the algorithms used to assign orders and to determine worker compensation. Later that month, the SAMR and the CAC led a group of seven agencies in co-issuing new guiding opinions to "safeguard the rights and interests of food delivery workers." The first provision in the document required platforms not to use "the strictest algorithm" but instead to adopt a "moderate algorithm" in determining the distribution and time expectations for orders. Other provisions in the regulation target issues highlighted in the Renwu article, such as setting floors under worker income and increasing enrollment in insurance programs.

In August 2021, the CAC followed up with its draft regulation on algorithmic recommendations, which included the provisions calling for the protection of workers whose schedules and compensation are determined by algorithms (see figure 6). During author discussions with advisers involved in drafting the regulation, those advisers explicitly cited the Renwu article and the outcry as the spark that led to the inclusion of these provisions. Among regulatory bodies, the CAC is particularly attuned and responsive to public outcry, a sensitivity that stems from its role monitoring and censoring online discussion.

Though the requirements in both of the regulations are relatively vague—"protect workers' rights and interests"—the companies responded. In July 2022, an engineering manager at one of the delivery platforms told the author that augmenting the delivery algorithm became a top engineering priority during that time. Some of that work can be seen in the public-facing versions of the companies' algorithm registry filings. Both Meituan and Ele.me took pain to highlight how they had

loosened up time requirements, with Meituan claiming that its algorithm generates four estimates for delivery time and it chooses the longest one. Whether those changes materially improved the experience of delivery drivers is beyond the scope of this paper, but anecdotal evidence suggests that these workers still face grueling and precarious work conditions.

Despite the Chinese state's eagerness to protect workers from the excesses of algorithms, it was not receptive to direct agitation by those workers. Strikes by delivery workers had been ramping up in the years prior to the Renwu article, and the leaders of these strikes were often detained by police. Shortly after one Ele.me delivery driver self-immolated over unpaid wages in January 2021, a prominent labor organizer was arrested. This combination of police crackdowns and new regulations is part of the party-state's playbook for responding to social issues: arrest those who organize unsanctioned collective action and simultaneously enact policies intended to address the underlying cause of public anger.

While worker protections stemmed from investigative reporting and public outcry, the use of the term "deep synthesis" had a different origin: corporate thought leadership.

By early 2020, the party-state had clearly singled out deepfakes as a target for regulation. For a major Chinese technology corporation with big ambitions for AI-generated content, there was reason to worry that regulatory overreach would preemptively smother a whole new AI-driven industry.

The Chinese tech giant Tencent is best known as the creator of WeChat, the Chinese super-app that allows users to chat, make payments, book a doctor's appointment, and hundreds of other functions. But Tencent is also a digital media empire consisting of streaming platforms, a movie studio, and some of the most successful online games ever created. The opportunities to leverage AI across this portfolio are massive. Looking at both the commercial opportunities and the potential regulatory headwinds, Tencent's in-house policy think tank conducted some AI policy thought leadership.

In May 2020, Tencent Research Institute (■■■■■) launched a report titled, "AI-Generated Content Development Report 2020: Deep Synthesis's First Year of Commercialization." The report deftly reframed discourse about AI-generated media, switching both the terminology and the focus of the discussions. It argued that using the term "deepfake" to describe all AI-generated content was "amateur and unscientific" and that it unduly stigmatized the technology. Instead it proposed "deep synthesis" as a more neutral and inclusive term. The report then laid out the potential applications of deep synthesis across industries ranging from entertainment to education to healthcare. It predicted that 2020 could be the year that commercial applications of the technology flourished.

The report did not deny that deep synthesis presented new problems nor dispute that some regulation may prove necessary. It described the sordid origin of the term "deepfake" on Reddit and described threats posed from fraud, misinformation, and even election interference. A substantial section of the report outlined U.S. and European legislative proposals targeting deepfakes, as well as China's own regulations. It discussed technical countermeasures like detection techniques and industry attempts at self-regulation.

But the report warned against regulatory overreaction. It claimed that deep synthesis "won't erode truth in society, and it isn't a threat to world order." It cautioned that governments should not rush to regulate deep synthesis but instead adopt a tolerant and prudent approach that does not obstruct useful applications and technological innovation.

These conclusions were convenient for the interests of Tencent, but the arguments in the report were not without basis. At the time, much of China's policy discourse was using the loaded term "deepfake" to refer to a wider batch of AI applications, and "deep synthesis" was a more neutral and

appropriate term for AI-generated content. (The term “generative AI” had not yet entered mainstream policy discourse.) As with much of the work from Tencent Research Institute, the report was well-researched, technically sophisticated, and forward-looking. It was also an attempt to steer the policy conversation away from the threat of deepfakes and toward things the Chinese party-state wants, like innovation and increased productivity.

On one level, this appears to have worked. Prior to Tencent’s report, the term “deep synthesis” had never appeared in Chinese state media. In the eighteen months following the report’s publication, the term gained some traction in policy discourse, and in January 2022 it jumped to the center of that discourse when the CAC published its draft regulation. How exactly the term entered and held sway within the drafting process remains a mystery. It is possible that Tencent’s government relations team directly pushed for its adoption or that it simply entered the Chinese policy lexicon and then was taken up by regulators.

But on another level, the report may have partially backfired for Tencent. Chinese regulators had previously been focused explicitly on deepfakes, most often in the form of face swapping. In an attempt to showcase the many productive uses of the technology, Tencent’s report widened the aperture to include many other kinds of AI-generated media that would not properly be called deepfakes. Regulators ended up using not only the term but also the much more expansive definition in the regulation, including things like AI-generated text and image enhancement. It is possible that regulators were already planning to cast a wide net on the promised deepfake regulation and that the report did not actively expand that scope. But it is also possible that this bit of corporate thought leadership contributed to a more far-reaching—if somewhat less severe—regulation.

The most consequential feature of the two regulations was the creation of the algorithm registry (■■■■■■, sometimes translated as “algorithm filing system”). This system went largely unremarked upon at the time, but it has gone on to become a key piece of China’s AI governance toolkit—a way for the party-state both to gather information and to impose requirements on Chinese AI developers. So, what is the registry, and where did it come from?

China has a long tradition of requiring companies and individuals to register projects, products, and much more with the government. These registration systems (■■■■) keep track of everything from construction projects to published books to the color of each car. In most cases, they only require basic information: the individual’s name, the company or organization, and some information specific to the object or activity. They are intended simply to create a record for the government and are explicitly not licensing systems—though, in the case of the algorithm registry, this line is becoming blurred.

After the dawn of the internet, the Chinese government began creating registries for cyberspace, beginning with noncommercial websites (as commercial sites require a license). These early online registries were created and managed by the bureau that would go on to become the MIIT. But the CAC entered the fray in 2019 by creating a registration system for blockchains. Given this background, it is not entirely surprising that the CAC would eventually create a registry for algorithms, but it also was not a given. These registries are usually created only after a new regulation demands them. For example, China did not have a formal registry for mobile apps until August 2023, when a new law covering online fraud mandated it.

China’s recommendation algorithm regulation created that mandate. It required that all providers whose algorithms have “public opinion properties or the capacity for social mobilization” submit a registration within ten days of operating. The information explicitly mandated by the regulation is fairly basic: name and type of algorithm, form of service, application setting, and an “algorithm self-assessment report.” When the first batch of registrations were made public in August 2022, they included some additional details, such as a very high-level and nontechnical description of the functioning of the algorithm. These details, and the fact that they were released publicly, were new

additions compared to previous registries. If this were the full extent of information given to Chinese regulators, it would provide them with few meaningful insights into the algorithms.

But the CAC was publicizing only a fraction of the information it gathered through the registrations, with much more being kept private due to concerns over trade secrets. Some insight into that additional information is available in a downloadable user manual that the CAC posted on the algorithm registry website. The document is intended to guide algorithm providers through the registration process, and it contains screenshots of various stages of the process. These screenshots show that the CAC is asking for significantly more information on input data, such as whether it contains biometric features or personal identity information. Providers are also asked to list all the open-source and "self-built" datasets that the algorithm was trained on, as well as to upload an "algorithm security self-assessment report." The user manual does not contain any further details on the contents of that report, but these reports are likely modeled on a 2018 document from the CAC on conducting security assessments for internet services with "public opinion properties or social mobilization capabilities."

Many Chinese companies were initially confused as to the scope of applications covered and how to satisfy certain requirements. The first batch of registrations included everything from search-engine filtering algorithms to apps that recommend parenting tips and child-rearing content. And there remained a huge knowledge gap between companies and the CAC regulators when it came to the functioning of the algorithms. The Wall Street Journalreported on a meeting between representatives from ByteDance and the CAC, in which officials from the agency "displayed little understanding of the technical details" of algorithms, forcing company representatives to "rely on a mix of metaphors and simplified language."

But both the algorithm registry and the regulators involved have evolved. When the deep synthesis regulation was finalized in December 2022, the registry was updated to include a separate registration process for technical support teams, which face new requirements under the regulation. The biggest changes came with the explosion of generative AI applications in 2023. The new generative AI regulation from July 2023 did not explicitly change the registration process and simply mandated that providers fulfill their existing obligations under the algorithmic recommendation and deep synthesis regulations. But in practice regulators began treating the registration process more like a licensing regime than a simple registration process. They did this by withholding their official acceptance of registrations until they felt satisfied with the safety and security of the models.

This process of ensuring that models are "safe enough" for release began as an informal back-and-forth between companies and regulators. But in October 2023 a key Chinese standards body released a draft standard that laid out specific tests a generative AI model must pass before registering. If implemented, that draft standard would mandate that each dataset used for training contain less than 4 percent objectionable content. For sensitive prompts—covering ideology, politics, race, and gender, among other topics—over 90 percent of the model's outputs must be acceptable. The draft standard includes detailed requirements for several other tasks, such as what percentage of user prompts on sensitive topics a model can refuse to respond to.

Over the past two years, Chinese regulators have repeatedly iterated on and expanded the function of the algorithm registry. It has gone from being a straightforward tool for limited algorithmic transparency to a flexible, multipurpose tool of AI governance.

So where did the idea for creating a new registry for algorithms come from? The earliest public call for the creation of an algorithm registry appears be in 2016. At the time, the Chinese search engine Baidu was embroiled in controversy. As mentioned above, a Chinese student had died after undergoing bogus medical treatments, which he had found through an ad on Baidu, leading to a major public outcry. At a forum with internet policy experts debating the issue, legal academic Liu Deliang advocated for the creation of an algorithm registry, one that would include Baidu's search engine algorithm. Liu argued that without a registry, the public and regulators had no way of

knowing whether Baidu had tampered with the search results for financial gain. In the following years, Liu continued to argue for greater regulation on algorithms, including in a 2019 interview with the influential newspaper Legal Daily calling for mandatory public disclosure of corporate algorithms to facilitate societal oversight.

But it is also possible that inspiration for China's algorithm registry came from a far different source: New York City. Chinese academics—and to a certain extent, regulators—frequently engage with their counterparts abroad, discussing how different countries have dealt with issues China is now facing. These comparative policy studies go back many decades and span dozens of countries, with economists from Eastern Europe and industrial planners from Japan deeply engaging with and influencing leading Chinese scholars as they confront these policy areas.

During the CAC's consultations with policy experts for the algorithmic recommendation regulation, one expert pointed to New York City's newly created directory of algorithms used by city agencies. The directory was the product of an executive order by the mayor requiring that agencies disclose their use of algorithms in decisionmaking. It includes basic information on the purpose and functioning of the algorithms, including the types of data they were trained on. For example, it describes a "growth model" used by the New York State Education Department to assess teacher performance. The model compares student achievement to projections made by the model, which are based on economic disadvantage indicators, disability, prior achievement, and other factors.

The New York City directory bears many similarities to the public-facing versions of China's algorithm registrations. Both include a basic description of the purpose and functioning of the algorithm and sometimes of the data it is trained on. There are also clear differences between the systems, which reflect different priorities of the two governments. New York's registry exclusively targets algorithms used by public agencies, facilitating greater citizen oversight of their government. The Chinese algorithm registry targets algorithms used by the private sector, giving citizens some information—and the state even more information—about them.

This priority reflects a dichotomy that underlies China's governance of data and AI. The regulations provide Chinese citizens with meaningful protection from Chinese companies, but they do not provide that same protection from the actions of the party-state. To many non-Chinese observers, this difference appears to be a fundamental contradiction, one that invalidates a core purpose of technology regulation. In China, this dichotomy is simply the reality.

The regulations provide Chinese citizens with meaningful protection from Chinese companies, but they do not provide that same protection from the actions of the party-state.

This paper has traced the policymaking process for two of China's foundational regulations on algorithms and AI, illustrating the diverse forces that spark and shape the country's technology policymaking. The initial spark for these regulations derived from a real-world change: a new technology application or business model that affected China's economic, social, and political order. In the case of the recommendation algorithm regulation, that change was the rise of algorithm-driven news apps like Toutiao and the new challenges they posed to content moderation and ideological control. Other real-world shifts that eventually found their way into the regulation included the rise of algorithm-driven food delivery apps and the adoption of algorithmic price discrimination by e-commerce platforms. In the case of deep synthesis, it was a technological innovation that enabled face-swapping applications that sparked a public outcry over their implications.

Those shifts were then filtered through the twin lenses of CCP ideology and, in some cases, Xi's own preferences. Algorithm-driven news apps undermined CCP conceptions of the role of media, and the party weighed in through a series of editorials. The impact of algorithms on Chinese migrant laborers similarly ran counter both to professed party doctrine and to Xi's push for "common

prosperity," the latter being a factor contributing to the crackdown on China's leading technology platforms. While the rise of deepfakes did not go through a period of rigorous ideological analysis, the threats to information environments posed by fabricated images were clear to policymakers in China and around the world.

These issues then entered China's "world of ideas," where they were further defined and shaped by academic research, technical think tanks, investigative journalism, public opinion, and corporate lobbying. Here, the broad party-state imperative to do something about these technology applications was steadily molded into concrete interventions, such as new technical standards and certifications. Chinese sociologists and magazine journalists sharpened the focus on algorithmic exploitation, and a widespread public outcry brought this issue to the forefront. Meanwhile, corporate thought leadership reframed the debate on deepfakes, swapping the terminology and broadening the scope. All throughout this process, international discussions of these same topics permeated China's world of ideas, with China's state media and policy community consistently engaging with and sometimes adapting those ideas for the Chinese context.

All of these influences fed into bureaucratic attempts to formally regulate the technologies. These attempts were highly iterative, beginning with blunt corporate punishments before evolving into progressively more robust regulations. The CAC led the charge in both cases, relying on extensive consultations with Chinese legal academics and technologists to shape the text of the regulations. It also looped in related ministries and agencies, some of which added new provisions to address their own concerns, such as algorithm-fueled monopolistic behavior. What emerged from the end of this process were two of the world's earliest major regulations on algorithms and AI.

While the recommendation algorithm and deep synthesis regulations laid the foundations for China's AI governance regime, they were just the beginning. In November 2022, just five days after the deep synthesis regulation was signed, OpenAI shook up the world with the release of ChatGPT. The chatbot's ability to produce credible writing on innumerable topics wowed users around the world, including many in the CCP and China's government. The deep synthesis regulation technically already covered tools like ChatGPT—under the phrasing "technologies for generating or editing text content"—but it did not anticipate the power or popularity of this new generation of large language models.

Regulators quickly decided new rules were needed. Over the next eight months, the CAC and other regulators quickly produced a draft and then an "interim" generative AI regulation. Technological and economic shifts during this time were altering the state's calculus on trade-offs between political security and AI development, and the regulation was the subject of intense intellectual debate and bureaucratic wrangling. The next paper in this series will examine these debates and the continued evolution of China's AI governance regime, both at home and on the international stage.

Those developments followed a similar, though much accelerated, path through the policymaking process detailed in this paper. The process featured many of the same players—ministries, policy advisers, and technical organizations—as well as some new ones. And it further elucidated the key tensions and likely trajectories of China's AI regulations.

What comes next will hold profound implications for the development and governance of AI, both within China and around the globe. Over the next few years, Chinese companies will continue pushing AI products onto global markets, and Chinese diplomats will engage with and reshape international governance regimes. The country's domestic regulations will both constrain and enable those efforts, and understanding the individuals, ideas, and institutions guiding China is all the more important.

The author would like to thank Zhang Peng, Jamie Horsley, and Jeremy Daum for their helpful feedback and corrections to an early draft of this paper. The responsibility for the content and any remaining errors remain exclusively with the author.

China is leading the world on AI governance: other countries must engage

The nation's regulators are attempting to match the rapid pace at which AI is evolving. That needs to become a global initiative.

China has been forging its own path on the regulation of technologies based on artificial intelligence.Credit: Hector Retamal/AFP/Getty

The past few years have seen no shortage of international dialogues, white papers and recommendations from advisory groups on the development and use of artificial intelligence. Yet when it has come to turning these into globally agreed rules to maximize the benefits and minimize the harm of AI, there has been a leadership vacuum.

As Nature reported last week (Nature https://doi.org/qhbv; 2025), one country is pushing forwards with plans to change that. China is proposing to set up a global body to coordinate the regulation of AI, to be known as the World Artificial Intelligence Cooperation Organization (WAICO). Establishing such a body is in all countries' interests, and governments around the world should get on board.

AI models have astounding power, and abilities that could supercharge science and boost economic growth. But they do not fully understand the world, and can fail in unpredictable ways. There are many ways in which they could cause harm, including exacerbating inequality, aiding criminality and assisting the spread of mis- and disinformation. Some prominent researchers even argue that superintelligent AI could one day destroy humanity.

So far, such risks have not been given due attention in the breakneck race to develop AI — a race that many fear has created an economic bubble that is on the brink of bursting. The United States, which is home to many of the companies making the most powerful and widely used models, has no national AI regulations, just a patchwork of state-level laws. On the whole, companies in the United States are expected to police themselves and establish their own internal guardrails — while also being in relentless competition.

The latest assessment of large technology companies' safety and risk policies — the AI Safety Index — by the Future of Life Institute, based in Campbell, California, was published on 3 December. On a scale from A to F, no US firm scores higher than a C+ (see go.nature.com/48ikyhv). Yet last month, US President Donald Trump launched an initiative dubbed the Genesis Mission, which will give companies and researchers developing AI models unprecedented access to government data sets. The administration has compared it to the Apollo programme to reach the Moon.

In the European Union, the AI Act, introduced last year, requires the makers of the most powerful advanced AI systems to strengthen their analyses of the threats their models pose. The act is being implemented in stages, and it is not yet clear what effect the threat of substantial fines for non-compliance will have. Media reports suggest that companies are pressuring the EU to water down its laws.

China, meanwhile, has charted its own course. Its government is pushing hard to embed AI across society, from local government chatbots to factory robots that can accelerate production. But its regulators have also been pushing for AI outputs to be traceable and companies to be accountable (prominent China-based AI companies score even lower on the AI Safety Index than do their Western counterparts, partly, according to the Future of Life Institute, owing to failure to address catastrophic misuse).

Since 2022, the country has introduced an array of laws and technical standards that require developers to submit generative AI models to regulators for pre-deployment safety assessments and to include visible, indelible watermarks on AI-generated content, in a bid to prevent fraud and misinformation. That push is accelerating: China issued as many national requirements on AI in the first half of 2025 as it did in the previous three years, according to the Beijing-based consultancy Concordia AI.

Chinese officials want to prevent AI from rocking political stability, but still allow it to be used to boost economic growth, says Angela Zhang, a law researcher and technology-regulation specialist at the University of Southern California in Los Angeles. Any attempt to control the outputs of AI models in the name of safety, whether for political purposes or to align outputs with social values, remains fraught. But, as neuroscientist and AI entrepreneur Gary Marcus told reporters at an event in London in October, in many respects China's motivations are "pretty similar" to those of other countries. "They don't want to be annihilated by robots."

Understanding China's standards on AI is important for everyone. The country's policy of producing free or cheap-to-access 'open weight' models means that companies around the world are increasingly building their services on Chinese AI. At the same time, it's important that China's researchers engage multilaterally to help assess what kind of global governance would be both effective and realistic.

Existing efforts to oversee AI at the global level — for example, the Organisation for Economic Co-operation and Development's AI Principles and the Council of Europe's Framework Convention on AI — are either non-binding or not enforced. To effectively regulate AI, unique solutions will be required, but one model that WAICO could learn from is that used by the International Atomic Energy Agency in Vienna to regulate nuclear energy. This involves countries agreeing to limits on development and to opening nuclear facilities for inspection, which allows nations to verify that others are adhering to the system.

Much of the rhetoric surrounding AI frames it as a geopolitical contest, whereby a country's safety is guaranteed by technological dominance. But 'winning' is not an option for most nations, and an AI arms race won't guarantee that people will be safer or wealthier. A better plan would be for everyone to come together to agree on what is safe and how the world wants to harness AI. China's initiative should be welcomed, and researchers and authorities around the world should engage.

India AI Governance

India’s New AI Governance Plan is Much Ado About Nothing

Indian Prime Minister Narendra Modi addresses armed forces personnel during a Diwali celebration onboard the INS Vikrant on October 20. (Prime Minister's Office)

India’s Ministry of Electronics and Information Technology earlier this month released governance guidelines for artificial intelligence, marking the first major regulatory document that has emerged in the country — if one can call it that.

Following an earlier draft in January, this 66-page final report confirms that the Indian government plans to create a regulatory framework for AI that focuses largely on industry self-regulation.

In fact, while introducing these guidelines, Ministry of Electronics and Information Technology Secretary S. Krishnan said that there had been a “conscious and deliberate approach of not leading with regulation” instead emphasizing the overall hands-off approach that the Indian government has taken when it comes to regulating AI.

The timing of the release was significant, as the guidelines would enable the Indian state to present a governance framework at the India AI Impact Summit scheduled in New Delhi in late February, where many of the world’s top decision-makers on the issue will descend. It would be more useful to see these guidelines unveiled in conjunction with a broader set of policy and industry interventions in the run up to the summit, including the India AI Mission, the Indian government’s primary AI development project.

Much of the guidelines are focused on Modi’s whole of government approach. They seek to ensure clear channels for communication and coordination between different arms of the government, to build capacity for the bureaucracy and direct regulators to make sure AI is used responsibly in the public sector.

While the AI guidelines are light on details, the document signals the Indian government’s overall vision on AI policymaking.

This vision views all emerging technologies from the lens of digital public infrastructure (DPI). The issue of infrastructure lends itself naturally to AI policymaking — the issues of developing GPUs, increasing access to datasets and expanding AI infrastructure feature prominently.

However, unlike the discussion on AI infrastructure in most other countries, the Indian policy document also views AI as an extension of its existing digital infrastructure. Amid concerns over India’s ability to develop globally competitive AI solutions, this may help India leverage its thought-leadership on DPIs and position itself as relevant in the AI race.

The guidelines clearly recommend the integration of AI with Aadhaar, India’s national digital identity system; UPI, the state endorsed mobile payment infrastructure; and DigiLocker, a personal online repository the government made available to Indian residents that is used primarily for authentication purposes. India’s existing digital language repository, Bhasini, also finds a key mention. Similarly, India’s regulatory technology privacy initiatives, such as DEPA (Data Empowerment and Protection Architecture), a techno-legal system for permission-based data

sharing, is touted as a potential solution to the privacy-innovation tradeoff questions that have dominated policymaking around AI globally.

While acknowledging the risks posed by the use of AI, it largely passes the buck on taking greater action, concluding that “many of the risks associated with AI can be addressed under existing laws.” In such cases, it diagnoses the primary problem as a lack of predictable and timely enforcement. However, even on this issue, aside from some homilies about regulators needing to be vigilant, it offers no substantive contribution.

On the question of liability, the guidelines appear to recommend a graded system of penalties, while in the same breath warning against accountability mechanisms stifling innovation. The guidelines immediately move towards self-regulations and have prescribed the setting up of an AI Governance Group (AIGG).

Any specificity in the guidelines is reserved for the issue of regulating deepfakes. It suggests that the AIGG could review tools for watermarking and labeling AI-created material to pinpoint its source, potentially tracing it back to the specific databases or large language models (LLMs) used to generate it. This is a clear reference to the latest amendment to the IT Rules (2021), which mandates labeling by platforms.

Together, these facets clearly align with India’s hands-off approach to AI regulation, choosing to give even some of the more controversial use-cases such as facial recognition in public services a wide berth. Only the politically sensitive issue of deepfakes has so far been singled out for regulation.

While the guidelines envision self-regulation, it remains to be seen what form ensuing industry codes of conduct or model regulations may take.

Aside from allowing the India government to tick off the need for a governance framework at February’s AI Summit, they do not achieve much.

Global AI Governance Law and Policy: India

From single computer in Kolkata’s Indian Statistical Institute in the 1950s to powering the world’s digital infrastructure, India’s technical journey has been tremendously transformative. In the 1990s, the government focused on the ‘information technology’ sector by encouraging the export of software services. This resulted in India eventually emerging as a global IT hub first as an outsourcing destination and eventually as the birthplace of digital public infrastructure that offers a new form of governance using population-scale digital architecture. Today, India has the world’s largest digital identity system, the biggest digital payments system by volume and a population that is, for the most part, digital by default.

This dramatic digital transformation has naturally fuelled the adoption of AI. According to the Stanford Artificial Intelligence Index Report 2025, India ranks second in the list of countries with the highest AI skill penetration from 2015 to 2024 and is among the top 10 countries in the world that received the most private investment in AI from 2013 to 2024.

India's AI policy has two different facets. One set of policy initiatives focuses on promoting AI adoption amid rapid advances in generative AI. The government has pledged to invest USD1.25 billion in AI development, and, to that end, launched the IndiaAI Mission. Additional initiatives

promote the integration of AI across a range of use cases and sectors. The other set of policy initiatives relates to the governance of AI and the risks that it could pose. While India currently does not have standalone AI legislation, existing intellectual property, data protection, cybersecurity and content regulations are being adapted to apply to AI.

Historically, India's digital space was exclusively regulated by the Information Technology Act, 2000, an omnibus law that addresses online contracts, data protection, cybercrimes and digital harms, such as phishing and identity theft. The IT Act has undergone numerous amendments to respond to new threats in the digital space. Originally intended to regulate computers and electronic records, it has been used to regulate a range of digital products and services due to the expansive definition of computer resources. AI systems and models will likely fall within the purview of its legislative scope.

Various subordinate legislations were enacted under the IT Act, including the IT (Reasonable Security Practices and Procedures and Sensitive Personal Data or Information) Rules, 2011, which governed how entities should process personal data. Until they were replaced by the Digital Personal Data Protection Act, 2023, these rules served as the primary data protection regulation. Other laws, such as the Bharatiya Nyaya (Second) Sanhita, 2023, and the Indian Copyright Act, 1957, broadly extend to the digital space and also apply to AI.

In 2018, India's National Institution for Transforming India, the government's apex policy think tank, released the National Strategy on Artificial Intelligence. The policy took an "AI for all" approach and aimed to address challenges of accessibility, affordability and skilled expertise.

The NSAI set out four key areas. One part concentrated on research to boost core and applied research in the AI field. Another focused on reskilling the current workforce to facilitate large-scale employment generation through AI; this initiative called for realigning the education sector to harness the potential AI. A third priority promoted investments in AI and product development that would enhance AI adoption. Finally, there was a focus to manage concerns around ethics, privacy and security, such as through the use of privacy preserving technologies.

In 2021, the IndiaAI Mission was launched to develop a comprehensive ecosystem to foster AI innovation by democratising access to compute, enhancing data quality, developing indigenous AI capabilities, attracting top AI talent and promoting ethical AI. The comprehensive national-level program houses initiatives like the IndiaAI Compute Capacity, which intends to scale AI computing infrastructure by deploying over 18,000 graphics processing units through strategic public-private partnerships.

The Ministry of Electronics and Information Technology indicated it was looking to replace the IT Act with new legislation in 2023, tentatively titled the Digital India Act. Since the Indian government is unlikely to introduce a standalone legislation on AI, it is likely the Digital India Act will also apply to potential AI risks.

The MeitY issued an AI advisory addressed to all intermediaries and platforms on 15 March 2024. It required them to ensure that their use of AI models, large language models and generative AI do not make it possible for users to share unlawful content on the platform. Additionally, they were required to ensure the use of AI technology does not result in any bias or discrimination or threaten the integrity of the electoral process.

All platforms and intermediaries were obliged to test their AI models; if they were unreliable in any way, intermediaries and platforms had to appropriately label them as such. Users had to be informed through terms of service or user agreements that their user account could be terminated if they were dealing with unlawful information.

The MeitY has recently constituted a subcommittee to examine whether the IT Act, the Bharatiya Nagarik Suraksha Sanhita, 2023, various content laws, the Information Technology (Intermediary Guidelines and Digital Media Ethics Code) Rules, 2021, cyber security laws, DPDPA and India's intellectual property regulations are sufficient to regulate AI. Since AI spans multiple sectors, the view was that a fragmented regulatory approach may result in inefficiencies. To address this concern, the committee recommended an integrated government approach, calling on the MeitY and principal scientific adviser, who advises the prime minister on matters of science and technology, to establish an empowered mechanism to coordinate on AI governance.

While there is no official guidance on how existing intellectual property laws apply to AI, the NSAI notes that the IP regime must enable AI developers to innovate. The policy, therefore, has recommended the creation of a task force comprised of the Ministry of Corporate Affairs and the Department for Promotion of Industry and Internal Trade to consider the issue.

The Consumer Protection Act, 2019, offers broad protections to consumers that could extend to AI harms. The government recently issued the "Guidelines for Prevention and Regulation of Dark Patterns," prohibiting the use of dark patterns on the grounds that they amount to deceptive advertising.

Various sectoral regulators are looking to regulate AI within the specific domains they supervise. Some key initiatives are listed below:

India's Advance on AI Regulation

This publication was produced under Carnegie India's Technology and Society Program. For details on the program's funding, please visit the Carnegie India website. The views expressed in this piece are solely those of the authors.

Artificial intelligence (AI) is a general-purpose technology that has existed since the early 1950s. Its trajectory is marked by cycles of hype and innovation, followed by periods of stagnation and disillusionment. In 2022 alone, more than thirty laws related to AI were passed in over a hundred countries. What explains this sudden rush to regulate AI?

Some say the launch of ChatGPT in late 2022 was a defining moment. It brought generative AI to the forefront, and along with it, concerns about bias, misinformation, copyright violations, and the impact on labor markets. One might also point to a confluence of factors—the massive breakthroughs in machine learning, new and powerful capabilities of large language models, and the global reach of social media—which has stoked the fears of policymakers and prompted new regulations in some countries.

And yet, nobody seems to have a clear collective vision for how AI should be regulated. This has resulted in divergent approaches around the world—from comprehensive legislation in the European Union (EU) to technology-specific rules in China, and voluntary commitments in the United States. Despite these differences, global policymakers seem to agree on one thing—we must leverage the power of AI while mitigating its risks.

Where does this leave India on AI regulation? The existing body of literature concerning India's approach to regulating AI is disconnected, narrow, or superficial. It includes news coverage of regulatory proposals; brief commentaries on national and global policy developments; editorials on what India's approach should be; summaries of the legal landscape; readouts of roundtable

discussions; and analyses of specific regulatory issues involving AI. What is missing is a clear and comprehensive analysis of India's overall advance on AI regulation. Who is driving the debate in India? What are the views of different stakeholders? Does India need new AI regulations? What should they look like?

To that end, this paper will capture the views of government, industry, and civil society in India and suggest a policy roadmap to inform India's advance on AI regulation.

Part I provides an overview of the current sentiment in government, industry, and civil society in India on the topic of AI regulation.

Part II explores the scope and objectives of AI regulation, the nature of AI risks, and areas where additional regulations may be required.

Our analysis is based on multiple discussions with key stakeholders over several months, both in closed-door and public settings.

We also conducted seventeen interviews with senior government officials, industry executives, lawyers, technologists, and scholars working specifically on AI policy in India. All interviewees have been granted anonymity to protect their privacy given their ongoing engagements on these issues and the sensitivity of these discussions.

Since 2022, the Indian government has oscillated between a hands-off approach to AI regulation and one that is more direct and interventionist, which has led to some confusion.

Broadly, India supports a "pro-innovation" approach to AI regulation. It wants to unlock the full potential of AI while taking into account the anticipated risks. This is reflected in the G20 Ministerial Declaration made during India's presidency, as well as a statement in Parliament in April 2023 that "[the Indian government] is not considering bringing a law or regulating the growth of AI in the country."

However, around the same time, the Ministry of Electronics and Information Technology (MeitY) published a blueprint for a new Digital India Act, which includes a specific reference to the "regulation of high-risk AI systems." Then, after a lull of close to a year, the government issued an advisory in March 2024 that jolted the industry. The advisory, which mandated compliance with immediate effect, directed companies to obtain the government's permission before deploying certain AI models in India, and to take steps to prevent algorithmic discrimination and the distribution of deepfakes. Following sharp criticism, the advisory was withdrawn and replaced with a fresh one that continues to remain in force.

The reason for the government's fragmented approach is that there are multiple, differing views within the establishment. MeitY, which is the nodal ministry for technology regulation in the country, favors a "light touch approach." An official suggested that the much-criticized AI advisory from March 2024 was the product of another agency's influence and not the brainchild of MeitY. Some factions within the government want more regulation. For example, a member of the Prime Minister's Economic Advisory Council has published a report that characterizes AI as a "complex adaptive system" that requires "proactive regulatory intervention." Another key official involved in framing India's AI policy said that MeitY was "not doing enough" to address the risks of AI.

For now, it appears the Indian government is still building consensus while adopting a cautious approach. It has tasked the Office of the Principal Scientific Advisor (PSA), set up to advise the Prime Minister and the cabinet on matters of science and technology, to consult with different ministries and provide "strategic guidance" on AI regulation. A sub-committee, convened by MeitY

and reporting to the PSA, has prepared a draft report on “AI Regulation,” though it has not yet been published.

MeitY, for its part, is considering various regulatory options, including amending the Information Technology Act, 2000 (henceforth IT Act) which would be less time-consuming than adopting new legislation such as the proposed Digital India Act.

Meanwhile, sectoral regulators such as the Reserve Bank of India (RBI) and the Telecom Regulatory Authority of India (TRAI) have begun to articulate the risks of AI. Going forward, they are likely to play an important role in shaping policy and regulation.

Lastly, the Prime Minister’s Office (PMO) and National Security Council Secretariat (NSCS) will be highly influential in AI policymaking, given their cross-agency mandate and strong leadership.

The technology industry, as a whole, does not have a single view on what India’s approach to AI regulation should be. As one lawyer put it, “This is a fragmented ecosystem consisting of big tech companies, startups, industry bodies, and VCs [venture capital firms] … there is no one position.”

That said, the predominant view is that any strict regulation would stifle innovation and make it difficult for India to achieve the ambitious goals of the India AI mission, launched in March 2024 with an initial budget of Rs. 10,300 crores ($1.3 billion) over five years, spanning strategic initiatives across compute, foundational models, datasets, skilling, and safe and trustworthy AI. For example, none of the industry representatives we spoke to expressed support for a new standalone AI law. One cautioned that it was more important for India to “get it right rather than to act swiftly.” Another said that “urgency will create problematic regulation.” Instead, one startup founder suggested that India adopt an “iterative, light-touch, and collaborative” approach to AI regulation.

There are also some extreme positions in the industry. One tech policy executive said that there was no need for new regulations at all, arguing that AI presents no novel risks, a view not shared by many others. On the other end, some have explicitly called for some form of regulation. Google president Kent Walker has previously stated, “AI is too important not to regulate, and too important not to regulate well.” Similarly, Microsoft president Brad Smith has also observed, “there has never been an industry that has successfully regulated itself entirely… we need more laws, more regulation.”

Some companies have advanced specific regulatory proposals for AI regulation in India. Microsoft has advocated for new laws targeted at “highly capable AI foundation models” IBM has called on governments to “recognize co-regulatory mechanisms” and Google has called for a “a risk-based and proportionate approach to AI regulation [in India] focused on use cases,” a model supported by several companies.

Level 1 : Self-regulation that enables firms to proactively address the risks of AI through voluntary commitments, self-certification, and similar models.

: Self-regulation that enables firms to proactively address the risks of AI through voluntary commitments, self-certification, and similar models. Level 2 : Additional regulations to fill the legal vacuum and deal with high-risk AI use cases through bespoke rules, guidelines, and advisories.

Here, civil society refers to the third sector of society, distinct from government and business. This includes activists, scholars, academics, and lawyers.

Some civil society representatives have called for greater representation of women, gig workers and other marginalized groups in the debate on AI regulation because they are most likely to be

impacted by the negative effects of AI deployments.

Some activists also expressed distrust for industry lobbying efforts on self-regulation. One scholar called the argument that regulation would stifle innovation a “convenient oversimplification” that benefits incumbent commercial actors. One scholar argued that voluntary commitments are inadequate because they “merely outlined a set of principles,” and that it encourages “experimentation” which could cause harm to individuals and communities.

Most representatives did, however, agree that India should not adopt a comprehensive AI law, at least for the time being. One scholar argued that an omnibus law might lack the nuance and context required to regulate AI. Another academic put it succinctly: “India needs more guidelines, less hard-coded legislation.” Academics are also wary of overbearing regulation and worry that new rules could restrict their access to AI systems required for public interest research in areas such as disaster management and cybersecurity.

At the same time, the academics we spoke to believe that AI presents new risks which, according to them, existing laws are ill-equipped to handle. They want regulators to intervene in areas where AI could cause irreversible harm and violate fundamental rights.

According to one activist, “the government’s use of AI requires immediate intervention since there is a greater likelihood of impact on legal rights.” They called for a review of public procurement guidelines and the use of facial recognition technologies in public services.

In the legal community, there are broadly two camps. One group believes that there is an unnecessary rush to “create and circulate legislative drafts,” when in fact, only a narrow range of issues are a matter for rule-making. Lawyers in this camp believe that the focus should be on applying existing laws to AI to mitigate risks. The other group believes that a separate law for AI is required to deal with the unintended, downstream risks of AI being deployed across India in potentially harmful ways.

self-regulation should be encouraged at this stage. Additional rules are required to protect consumers, especially for specific high-risk use cases.

This raises some important follow-up questions—What are the risks of AI? Are they novel? What are the gaps in existing laws? What aspect of AI should we regulate?

To answer the question of how AI should be regulated in India, it may be useful to reason by analogy. In a recent paper for RAND, Michael J. D. Vermeer compares AI with four other general-purpose technologies: nuclear technology, the internet, encryption, and genetic engineering. He lists out various factors that would inform their governance, such as the risks posed, consensus on these risks, and the role of public-private partnerships in its development. Based on a similar analog, we suggest that for a dual-use, general-purpose technology such as AI, three fundamental aspects of regulation need to be clear upfront:

Objectives : Policymakers should clearly state the purpose of any coercive state action in the form of regulation. Generally, regulation is deemed necessary to address market failures that may be present in the form of negative externalities, information asymmetry, consumer harm, and abuse of market power. Another objective may be to enforce principles such as transparency, fairness, privacy, security, and accountability through a system of rights and obligations. The Indian government has outlined certain objectives for AI regulation—to mitigate user harm, reduce misinformation, increase accountability, and balance commercial interests—but has not explained why regulation is required in each case. For example, it is not clear why the government should intervene in contracts between creators and publishers, as proposed by the IT Minister under a “new AI law.” Scope : Regulating AI presents multiple challenges, including, for example, the

“pacing problem,” definitional issues, and what is called the “black box” problem. As one expert put it, “regulating AI is like regulating operating systems.” Therefore, it is important to clarify upfront the aspects of AI that should be regulated. Experts suggest that AI regulation should cover three aspects—inputs, outputs, and processes. Inputs may include training data and copyright material; outputs include automated decisions and AI-generated content; and processes may include models and algorithms.” Another expert suggested that India look at regulation of “data, models, and applications separately.” This appears to be a common view, with a senior MeitY official publicly stating that “Instead of trying to regulate [artificial intelligence] technology, [the Indian government] is looking at regulating its applications.” Liability : On the issue of whom the regulation should apply to, the de facto approach is to impose different obligations on the developer and deployer, as has been done in the EU AI Act. The question of liability is also an important strategic issue. As one startup founder put it, “India’s competitive advantage lies in the application layer, so any regulation [on deployers] should account for this.” Multiple experts conceded that India’s legal regime requires a significant update to clarify issues of liability. Under the current IT Act, digital service providers are classified as an “intermediary” or “publisher” (or their sub-categories). However, as one lawyer commented, these existing definitions under the IT Act when applied to AI systems are “being stretched too thin.” This has caused confusion about which market actors need to comply with the rules since “generative AI systems may not fall neatly within the purview of either publisher or intermediary.”

The government should clarify the objectives of regulation, wherein any proposed intervention should be mapped to specific market failures. Instead of regulating the underlying technology itself, policymakers should focus on standalone issues (for example, deepfakes) and individual aspects of the AI value chain (for example, data inputs). Laws must be clear about to whom these regulations apply and when they can be held liable. For example, the current platform classification and intermediary liability framework under the IT Act must be updated to reflect the AI value chain.

A risk-based approach to AI regulation is the most popular. Though implementation varies across jurisdictions, its core objective is to mitigate harm to individuals and society. A risk-based approach is also supported by senior officials in the Indian government, including the IT minister, the IT secretary, and several big tech companies.

In the current discourse, there is a general tendency to conflate the notions of risk and harm. A key difference exists between the two. As explained in a report prepared by a group of scientists led by professor Yoshua Bengio, considered one of the “godfathers” of AI, risk is derived from the “probability of an occurrence of harm and the severity of that harm.” In other words, risk has a “future-orientation” and “looks at the aggregate impacts of the system on groups of people and tries to (often controversially) quantify these harms.” For that reason, as one legal expert put it, “AI regulation should be risk-based, not harm-based because the harm has already occurred.” Therefore, an important exercise in a risk-based approach to regulation is to gather evidence of harm in order to measure and anticipate the associated level of risk.

Next is the question of how to classify these risks. A report commissioned by the UK government classifies AI risks into three types: malicious use risks, risks from malfunctions, and systemic risks. Risks can also be categorized as safety risks and fundamental rights risks, with overlaps between the two. Risks may also vary based on stages (for example, during design, development, or deployment); scope (systemic risks); time-scale (short, medium, or long-term); and the source of risk (inputs vs. outputs). Even as our understanding of AI risk continues to evolve, scholars at the Massachusetts Institute of Technology (MIT) have developed a repository of more than seven hundred AI risks as of 2024. Some governments have also incorporated these risk taxonomies into their legislative frameworks. For example, the EU’s Artificial Intelligence Act (AI Act) classifies risk into four levels: “unacceptable, high, limited, and minimal.”

In the table below, we synthesize the available literature and present our own risk classification framework to inform future AI policy debates.

The table below summarizes key AI risks with examples of harm to inform future policy debates. For the purpose of this analysis, “AI” includes general-purpose AI, generative AI, and artificial general intelligence.

Potential risk Examples of harm 1. Harmful content Unauthorized impersonation of an individual using AI-generated deepfakes, leading to violation of personality rights and financial loss.i 2. Privacy violations Use of personal data without the consent of the individual for the purpose of training AI models, resulting in a violation of privacy rights.ii 3. Cybersecurity threats Sophisticated attacks using AI systems, leading to breach of critical infrastructure through AI-powered cyber attacks.iii 4. Discrimination AI systems discriminating against certain groups in hiring decisions leading to violation of the fundamental right to equality.iv 5. Loss of control AI systems making autonomous decisions without human oversight, leading to unintended consequences or failure to stop harmful actions which could compromise individual safety.v 6. National security Public safety is compromised due to AI-powered cyber attacks, misuse of AI in CBRN (chemical, biological, radiological, and nuclear) weapons, leading to national security threats.vi 7. Product safety Misleading advertisements about the performance of AI systems that exploit consumers and cause physical or financial harm.vii 8. Intellectual property rights violations AI models being trained on copyrighted data without proper consent or compensation, undermining the rights of creators and owners.viii 9. Market concentration Increased ownership and control of AI technologies by a small number of powerful corporations, limiting access to compute resources.ix 10. Global inequality Widening digital divide between the Global North and South, with wealthier nations gaining access to AI technologies faster, deepening economic and technological disparities.x 11. Job displacement Automation of jobs through AI systems, leading to large-scale unemployment in industries such as manufacturing, customer service, and transportation.xi 12. Environmental degradation Increased energy consumption from training and deploying large AI models, contributing to environmental harm such as carbon emissions and resource depletion.xii 13. Superintelligence Risk of AI systems surpassing human intelligence, leading to existential risks if AI’s goals diverge from humanity’s interests.xiii i Nayan Chandra Mishra, “Urgently Need: A Law to Protect Consumers From Deep Fake Ads,” The Indian Express, October 24, 2023), https://indianexpress.com/article/opinion/columns/law-deep-fake-ads-anil-kapoor-personality-rights-8997267/) ii “LinkedIn Users Say Their Data is Being Collected for Generative AI Training Without Permission,” The Hindu, September 19, 2024, https://www.thehindu.com/sci-tech/technology/linkedin-users-say-their-data-is-being-collected-for-generative-ai-training-without-permission/article68658540.ece. iii Kevin Williams, “Cyber Physical Attacks Fueled By AI are a Growing Threat, Experts Say,” CNBC, March 3, 2024, https://www.cnbc.com/2024/03/03/cyber-physical-attacks-fueled-by-ai-are-a-growing-threat-experts-say.html#. iv Jeanita Lyman, “Workday Facing Discrimination Lawsuit Over AI Hiring Software,” Pleasanton Weekly, July 18, 2024, https://www.pleasantonweekly.com/courts/2024/07/18/workday-facing-discrimination-lawsuit-over-ai-hiring-software/ . v Adam Satariano and Roser Toll Pifarré, “An Algorithm Told Police She Was Safe. Then Her Husband Killed Her,” New York Times, July 18, 2024, https://www.nytimes.com/interactive/2024/07/18/technology/spain-domestic-violence-viogen-algorithm.html. vi Raina Talwar Bhatia, Evi Fuelle, and Lucia Gamboa, “AI CBRN Risks: Governance Lessons from the Most Dangerous Misuses of AI,” Credo AI, August 30, 2024, https://www.credo.ai/blog/ai-cbrn-risks-governance-lessons-from-the-most-dangerous-misuses-of-ai#. vii Adam Lashinsky, “Watch out. Companies Are Using “AI Washing” to Mislead Consumers,” The Washington Post, January 29, 2024, https://www.washingtonpost.com/opinions/2024/01/29/sec-ftc-ai-c3-internet-fraud/. viii Annapurna Roy, “Indian Publishers Seek Rules for Copyright Protection Against Generative AI Models,” The Economic Times, January 26, 2024, https://economictimes.indiatimes.com/tech/technology/indian-publishers-seek-rules-for-copyright-protection-against-generative-ai-models/articleshow/107154425.cms?from=mdr ix Jai Vipra and Sarah Myers West, “Computational Power and AI,” AI Now Institute, September 25, 2023, https://ainowinstitute.org/publication/policy/compute-and-ai. x Philip Schellekens and David Skilling, “Three Reasons Why AI May Widen Global Inequality,” Centre for Global Development, October 17, 2024, https://www.cgdev.org/blog/three-reasons-why-ai-may-widen-global-inequality#. xi Mark Talmage-Rostron, “How Will Artificial Intelligence Affect Jobs 2024-2030,” Nextford University,

January 10, 2024, https://www.nexford.edu/insights/how-will-ai-affect-jobs xii Renée Cho, “AI’s Growing Carbon Footprint,” State of the Planet, Columbia Climate School, June 9, 2023, https://news.climate.columbia.edu/2023/06/09/ais-growing-carbon-footprint/ xiii Nick Bostrom, Superintelligence: Paths, Dangers, Strategies, (Oxford University Press, July 2014).

Currently, there is no credible AI risk classification framework for India that is based on empirical evidence of harm. There are some examples of standalone risk assessments, but they are incomplete and unsubstantiated. For example, a 2021 NITI Aayog report outlines certain “risks and considerations” to operationalize responsible AI practices but doesn’t spell out these risks. Another report by TRAI briefly mentions certain risks—“low quality data, data biases, data security, data privacy, inaccurate or biased algorithm, and unethical use of AI”—but jumps straight into regulatory principles without analyzing the risks. Similarly, the Telecom Engineering Center has published a draft AI risk assessment framework but focuses entirely on fairness outcomes.

Therefore, to support future policymaking, we have identified five categories of risks that are most relevant from a regulatory perspective:

Malicious uses , including the creation and distribution of harmful AI-generated content that could result in the violation of legal rights or threat to public order and safety. Algorithmic discrimination , which could arise due to various factors, such as the use of unreliable or unrepresentative datasets in AI systems, and which may result in financial losses, loss of opportunity, and pose a threat to other fundamental rights and freedoms. Transparency failures resulting from the lack of adequate and relevant disclosures about how AI systems are being used to make decisions, what data is being used to train AI systems, and other forms of information asymmetry that pose a threat to privacy and other rights. Systemic risks resulting from the dependence on a small number of AI systems in critical sectors like finance, healthcare, national security, and other forms of market concentration that could cause disruptions on a broad scale with potentially catastrophic effects. Loss of control due to lack of human oversight in the development and use of autonomous AI systems, which could result in unintended consequences and threats to national security and public safety.

We recommend that these specific categories of risks be studied in detail in an effort to develop an appropriate AI risk classification framework for India. Moreover, since the general consensus is that India should regulate “high-risk use cases,” we suggest analyzing proposed AI deployments based on these risk vectors. Although there is no standard definition, “high-risk” applications generally include those in critical infrastructure, lending, credit scoring, insurance, product safety, consumer rights, law enforcement, and justice delivery. A regulatory approach focused on “high-risk use cases” is likely to find favor in India, though it will have to be an iterative exercise grounded in scientific research.

Risk classification should also be grounded in tangible evidence of harm and account for local factors, such as AI adoption, consumer awareness, and digital literacy levels in India. Market studies, like the one being conducted by the Competition Commission of India, are useful references to identify specific market failures in India’s AI ecosystem.

From a regulatory standpoint, a well-considered AI risk assessment would also help identify risks for which regulation is not warranted. In some cases, risks can be addressed through industrial policy (for example, enabling access to compute through subsidies and reskilling individuals to mitigate job displacement). Long-term risks may not warrant immediate regulation (for example, existential threats to humanity). Many other risks can be addressed through existing regulations—an issue that we analyze in more detail in the next section.

India has a complex legal system of laws and regulations, comprising the Constitution, statutory laws, rules, regulations, and guidelines. From an AI perspective, the relevant areas of law are

privacy law, intellectual property law, competition law, media law, employment law, consumer law, criminal law, contract law, and tort law.

The common view across government, industry, and civil society groups is that this existing legal framework can be applied to address many of the AI risks outlined in this paper. The table below illustrates how existing statutory laws can be applied to deal with some AI risks.

Nature of harm Applicable statutory laws Depiction of a child in a sexually explicit video that is AI-generated Information Technology Act, 2000

Information Technology Act, 2000 Discrimination in hiring decisions using AI recruitment tools Rights of Persons with Disabilities Act, 2016

Scheduled Castes and the Scheduled Tribes (Prevention of Atrocities) Act, 1989 Use of an individual's personal data without consent to train AI models Digital Personal Data Protection Act, 2023

Information Technology Act, 2000 Misleading ads about the reliability or performance of an AI service Consumer Protection Act, 2019 Use of copyright-protected material in AI-generated content without permission of the author or owner The Copyright Act, 1957

One legal expert suggested that policymakers should spend the next six to twelve months applying existing laws to AI use cases to understand potential gaps in the current framework, and that the Department of Legal Affairs, Ministry of Law and Justice, Government of India, should be entrusted with this task.

Another important factor will be the role of courts in interpreting, adapting and enforcing existing provisions of law. So far, Indian courts have dealt with only a handful of cases involving modern AI systems, focused primarily on AI-generated content involving the likeness of public figures. A cursory review of cases indicates that Indian courts are content to apply existing legal provisions to provide redress for now.

The table below illustrates which existing statutory laws would apply in the case of a deepfake image or video being circulated without the permission of the individual.

Section 319 of the Bharatiya Nyaya Sanhita, 2023 Transmitting obscene material Section 67A and 67B of the Information Technology Act, 2000 Causing harm to reputation Section 356 of the Bharatiya Nyaya Sanhita, 2023 Failure to observe due diligence guidelines for intermediaries Section 79 of the Information Technology Act, 2000, read with Rule 3 of the Intermediary Guidelines.

Experts have, however, warned that additional regulations are required because "there are fundamentally new risks and harms emerging from AI that existing laws are not equipped to deal with." One expert gave the example of a hospital collecting sensitive health data for a medical diagnosis, which was then repurposed to train AI models without the person's knowledge —an example of how "consent-based regimes break down completely in the AI context." Multiple experts also said that new legal rights are required to protect individuals and society given the pervasiveness of AI systems, the lack of transparency, and the disproportionate impact of these systems on vulnerable communities.

Therefore, on the question of whether new regulations are required to deal with the risks of AI, we suggest the following approach to guide future policymaking:

the risks can be adequately addressed through industrial policy (for example, reskilling initiatives to mitigate job displacement or subsidies to enable access to compute);

there is no pressing need to adopt regulation at this stage (for example, existential threat to humanity from artificial general intelligence);

risks can be addressed through existing regulations (for example, consumer harm arising from misleading advertisements about the performance of AI systems can be adequately addressed under existing laws).

2. Clarifications or targeted amendments are required in cases where advisories, guidelines or targeted legal amendments can sufficiently address risks, instead of adopting new and comprehensive rules. For example, the risks relating to the circulation of deepfakes can be broadly addressed through existing laws. However, regulators should clarify to whom these existing rules apply and in which cases they can be held liable.

3. New regulations should be considered to address market failures and protect consumers. For example, introducing a right to compensation in the case of misuse of AI and the right to object to automated decision-making to protect fundamental rights. New transparency obligations for certain AI systems should also be considered to address the risks relating to information asymmetry.

It is beyond the scope of this paper to conduct a comprehensive analysis of all the relevant laws and regulations and where they fall short in the context of AI systems. Therefore, we recommend that a comprehensive regulatory gap analysis be conducted to help inform future policymaking in India.

In Part II, we identified certain AI risks for which new regulations may be warranted, subject to evidence of harm. In Part III, we examine the different regulatory approaches that can be adopted to address these risks.

Self-regulation : In this approach, a “group of firms in a particular industry, or entire industry sectors agree to act in prescribed ways, according to a set of rules or principles.” In the AI context, self-regulation is generally organized around principles of privacy, security, transparency, fairness, accountability, and trust. Self-regulation can be implemented in many ways, such as voluntary commitments, self-certifications, and impact assessments. Co-regulation : In this approach, the government or regulators play a more proactive role by developing, recognizing, endorsing, or implementing standards. This is a more stringent approach and “represents a midpoint in the continuum between self-regulation and full government regulation.” Examples of co-regulation include codes of practice and risk management frameworks. Binding regulations : In this approach, policymakers enact a law or some other binding framework that makes them legally enforceable. The most prominent example is the EU’s Artificial Intelligence Act, which contains rights and obligations in relation to AI systems.

Each of these approaches have distinct advantages and disadvantages. According to a report prepared by the Stanford Cyber Policy Center, self-regulation helps tap into industry expertise, provides flexibility, and encourages rapid innovation, but lacks sufficient accountability and enforceability. Co-regulation enables collaboration between companies and regulators, allowing for an iterative approach, but often lacks the necessary enforcement mechanisms required to address market failures. Binding regulation provides clear accountability mechanisms and government oversight but could stifle innovation due to bureaucratic delays and the lack of expertise and adaptability.

Therefore, policymakers should carefully evaluate the relative cost and benefit of each approach before adopting them into domestic frameworks.

Regulatory approaches in different countries reflect their own socio-economic priorities, legal traditions, and governance models. For example, the EU's rights-based approach, in the form of a comprehensive statutory law, seeks to "protect health, safety, and fundamental rights." On the other hand, China, with its strong desire for state control, prioritizes social order and the protection of "Socialist Core Values" in its rules on AI-generated content. Japan, on its part, has professed a "human-centric" approach to AI that aligns with its broader societal goals. Singapore and the United Kingdom have adopted a principle-based approach that reflects a pragmatic style, tailoring rules to specific industries. The table below provides an overview of these different global approaches.

The table below provides a summary of the different approaches to AI regulation in various jurisdictions as of the date of publication of this paper.

Jurisdiction Summary of approach Type of regulation Australia The government has released a discussion paper proposing mandatory guardrails to regulate AI in high-risk settings and general-purpose AI models.i Binding government regulations are currently under discussion. Brazil Reviewing proposals for a new AI law that protects fundamental rights and ensures secure, reliable AI systems while categorizing them by risk and imposing various compliance requirements.ii Binding government regulations are currently under discussion. Canada Published the draft Artificial Intelligence and Data Act (AIDA) that focuses on responsible AI use, consumer protection, and fair competition.iii Binding government regulations are currently under discussion. China Techno-centric approach with specific regulations aimed at algorithm recommendations and generative AI.iv Binding government regulations have been adopted and are in force. European Union Statutory framework in the form of the AI Act that categorizes systems by risk levels, imposes stringent requirements on high-risk applications, and aims for transparency and accountability.v Binding government regulations have been adopted and are in force, with provisions for co-regulation. Japan Through the G7's Hiroshima process, Japan has promoted a light touch approach and a voluntary code of conduct. Japan has since established a set of domestic guidelines for businesses and is considering a statutory framework.vi Self-regulatory approach with ongoing discussion on binding regulations. Singapore Voluntary, use-case based approach that emphasizes a sectoral approach based on governance frameworks.vii Self-regulatory approach. United Kingdom Context-based and cross sectoral framework that focuses on core principles that will be implemented by sectoral regulators.viii Self-regulatory approach, with the option for sectoral regulators to frame binding regulations. United States of America Voluntary commitments and executive orders that emphasize a principle-based, cross-sectoral approach to promote industry best practices, and risk mitigation tools with input from various federal agencies.ix Self-regulatory approach, with limited downstream impact on advanced AI model providers from executive orders. iSafe and Responsible AI in Australia, (Department of Industry, Science and Resources, Australian Government, September 2024),https://storage.googleapis.com/converlens-au-industry/industry/p/prj2f6f02ebfe6a8190c7bdc/page/proposals_paper_for_introducing_mandatory_guardrails_for_ai_in_high_risk_settings.pdf.) iiBill no. 2338, of 2023,https://www25.senado.leg.br/web/atividade/materias/-/materia/157233.. iiiBill C-27, (First reading June 16, 2022). ivCyberspace Administration of China, National Development and Reform Commission, Ministry of Education Ministry of Science and Technology, Ministry of Industry and Information Technology, Ministry of Public Security, State Administration of Radio, Film and Television, Interim Measures for the Management of Generative Artificial Intelligence Services, (Enforced on August 15, 2023). vThe EU AI Act, Regulation (EU) 2024/1689. vi "AI Guidelines for Business Ver1.0," Ministry of Economy, Trade and Industry, Government of Japan, April 19, 2024, https://www.meti.go.jp/shingikai/mono_info_service/ai_shakai_jisso/pdf/20240419_9.pdf; Kohei Sakai, Takai Tsuji, and Ryohei Yasoshima, "Japan Weighs Regulating AI Developers, Following U.S. and EU," Nikkei Asia, May 2, 2024, https://asia.nikkei.com/Business/Technology/Japan-weighs-regulating-AI-developers-following-U.S.-and-EU. vii "Model Artificial Intelligence Governance Framework Second Edition," InfoComm Media Development Authority and Personal Data Protection Commission Singapore, January 21, 2020, https://www.pdpc.gov.sg/-/media/Files/PDPC/PDF-Files/Resource-for-Organisation/AI/SGModelAIGovFramework2.pdf. viii "A Pro-Innovation Approach to AI Regulation," Office for Artificial Intelligence, Department of Science Innovation and

Technology, Government of the United Kingdom, (Presented to Parliament on March 29, 2023), https://www.gov.uk/government/publications/ai-regulation-a-pro-innovation-approach/white-paper.
ix “Ensuring Safe, Secure, and Trustworthy AI,” The White House, July 21, 2023, https://www.whitehouse.gov/wp-content/uploads/2023/07/Ensuring-Safe-Secure-and-Trustworthy-AI.pdf; “Executive Order on the Safe, Secure, and Trustworthy Development and Use of Artificial Intelligence,” The White House, October 30, 2023, https://www.whitehouse.gov/briefing-room/presidential-actions/2023/10/30/executive-order-on-the-safe-secure-and-trustworthy-development-and-use-of-artificial-intelligence/ .

Based on our research and discussions with experts, there are four key reasons for supporting self-regulation. Firstly, industry stakeholders support self-regulation for AI because they have already been proposed by NITI Aayog and the Indian Council of Medical Research. Secondly, according to one executive, there are strong market incentives for companies to comply with self-regulation, for example, to secure valuable procurement contracts. This applies even to startups, for whom self-regulatory frameworks are “low-effort promises” to signal seriousness in the market. Third, voluntary commitments can co-exist with other regulatory models, such as co-regulation, so it is not a binary decision. Lastly, India has successfully implemented self-regulation in other parts of the digital sector, such as in digital advertising.

On the other hand, there are some strong skeptics of self-regulation. “The future trajectory of these technologies and their social impacts cannot be decided by the very individuals who stand to profit from them,” asserted one expert. Some industry executives also point to the failure of recent industry efforts to adopt generative AI guidelines, because they were “too broad in scope” and prescriptive. Another executive claimed that “the Indian government does not trust big tech companies enough to make self-regulation work. Moreover, according to one expert, self-regulation “doesn’t do much to separate good actors from bad actors.” This skepticism is also shared by some government officials who said that binding rules are necessary to contain some AI risks.

Many experts said that co-regulatory models were urgently needed in India to reduce the burden on public institutions while providing a higher degree of accountability, especially when compared to self-regulation. They also agreed that co-regulation may also be the most effective way to deal with “high risk use cases” for which government oversight is necessary.

Lack of conceptual clarity: The difference between self-regulation and co-regulation is not clearly understood. Many “self-regulatory bodies” (SROs) are subject to strict oversight, though they are presented as voluntary, industry-led approaches. Structural issues: The failure of co-regulatory models in India indicates deep-seated issues, including a lack of political will to cede authority, poor governance structures, lack of financial independence, and conflicts of interest. Delays in implementation: The government has rejected multiple proposals for co-regulation from the online gaming industry, leading to a regulatory vacuum.

For these reasons, we recommend that co-regulatory models should not be adopted for AI governance in India until they have been proven to work in other domains.

Many experts believe it would be premature for India to adopt binding rules for AI for several reasons. First, there is no comprehensive risk assessment on the basis of which new rights and obligations can be developed. Second, there is no empirical evidence of market failure to justify the increased compliance cost of new regulations. Third, existing laws can address many of the anticipated risks of AI, and a gap analysis is required to identify areas where new rules are required. And fourth, less expensive methods such as self-regulation may be sufficient to address the anticipated risks.

In the next section, we examine India’s broader policy agenda on AI and analyze various socio-economic factors that influence the choice of regulatory architecture.

In this section, we analyze three factors that should inform India's approach to AI regulation—economic opportunity, cost of regulation, and state capacity.

There is significant economic upside to the successful implementation of India's national AI strategy. According to the Indian government, AI is expected to add nearly $500 billion to India's gross domestic product (GDP) by 2025. The government has also identified potential applications of AI in agriculture, healthcare, disaster management, transportation, law, and finance that could have a transformative social impact.

The private sector is also bullish on the AI opportunity for India. Accenture anticipates that AI will increase the annual growth rate of India's economy by 1.3 percent by 2035, while a Google-commissioned report estimates that at least Rs 33.8 lakh crore of economic value will be generated by 2030 through AI adoption.

For these reasons, some researchers from the Global South have identified a new type of AI risk for developing countries—"the risks or opportunity costs of not implementing AI, [and] missing out on potential benefits." Indeed, this view is shared by many in India, as demonstrated in a survey conducted by Ipsos, which finds that "Indian respondents are more optimistic about AI than their global counterparts."

Therefore, we recommend that India adopt a light touch and pro-innovation approach that is aligned with its broader AI strategy to help fully realize the socio-economic benefits of AI.

First is the obvious industry cost of compliance. There are no estimates of what it would cost an Indian business to comply with a model AI law. However, for comparison, the European Commission estimates the cost of complying with the AI Act to be between €1.6 billion and €3.3 billion (Rs 143.5 billion to Rs 296.1 billion or $1.7 billion to $3.5 billion). Although this is a small fraction of the EU's overall GDP, which stood at $16.6 trillion in 2022, the new regulation has been criticized for going "too far" and setting the regulatory barrier "too high." Moreover, as India prepares to implement a new data protection law, new AI regulations will likely increase the cumulative cost for businesses.

Besides compliance costs, some industry executives fear the "psychological costs" associated with new regulations. They point to existing rules under the Apprentices Act, 1961, which grants officials wide discretionary powers to impose fines. In fact, one startup founder said that some new rules, such as the appointment of a local Data Protection Officer, only serve to increase compliance costs and create fear in the minds of entrepreneurs.

Other costs include the administrative expenses associated with implementing a new law. New regulations for AI could also increase litigation costs, especially if they are premature, unclear, or duplicative.

On the other hand, timely regulations could provide business clarity and promote innovation. Some say that the introduction of the Information Technology Act, 2000, helped propel the growth of India's e-commerce industry. Policymakers should also consider the costs of delayed regulation, though it is difficult to quantify the loss of legal rights and freedoms.

Multiple interviewees highlighted capacity constraints that, according to them, would prevent the Indian government from effectively implementing new AI regulations given the status quo. These limitations fall into five categories: (1) lack of technical expertise; (2) failure to issue clear and timely regulatory guidance; (3) lack of investigative powers; (4) ineffective or inconsistent enforcement; and (5) lack of grievance redressal mechanisms.

Addressing each of these issues will require a multi-fold approach. To promote effective AI governance, we recommend increasing state capacity in at least two respects:

Independent expertise : India should consider setting up a national “AI Safety Institute” (AISI) to develop state capacity in foundational research, safety and testing, training and awareness, and cross-border collaboration on AI governance. AISIs that have been set up in the UK, the EU, Singapore, Japan, and the United States guide industry compliance and facilitate information exchange on AI safety issues with other like-minded countries. While the Indian government has commenced consultations on a proposed AISI, experts caution that the scope and structure of the AISI should be carefully thought out to ensure that it functions effectively. Independent enforcement : Poor enforcement stems from a variety of factors—lack of resources and technical expertise, no separation of powers between the government and regulator, lack of consumer awareness, and ineffective grievance redressal mechanisms. There are two possible pathways to increase enforcement capacity—boost MeitY’s resources so that it can act as both policymaker and de-facto regulator, or set up a new regulator to enforce the IT Act and oversee AI governance.

Developing state capacity in these ways, we anticipate, would help separate policymaking from enforcement, promote industry compliance, and protect consumers.

Policymakers in India must better understand the current capabilities and unique risks posed by AI, the dynamic and evolving nature of the AI ecosystem, and the gaps in the existing legal framework. They will also need to adopt a balanced regulatory approach and be prepared to address market failures as and when the need arises.

Understand the risks and benefits of AI: Having a clear understanding of the current capabilities of AI is crucial for cost-benefit analyses, regulatory impact assessments, and risk assessments. We recommend adopting measures similar to the U.S. executive order on AI, which directs government agencies and regulators to issue reports on how AI is being used in their respective sectors and to identify potential risks. Market studies on India’s AI ecosystem, such as the one being conducted by the Competition Commission of India (CCI), are also required to identify market failures. In India, the Office of the Principal Scientific Adviser has been entrusted with inter-agency coordination on issues of AI governance. These discussions should be used to inform future policy and facilitate awareness, training, and capacity-building efforts. Classify AI risks based on evidence of harm: The government should identify specific AI risks for which new regulations are required. The five categories of AI risks identified in Part II of this paper provide a useful framework to gather evidence of harm in India and inform future regulation. Further, we recommend identifying “high-risk applications” based on these risk vectors for which additional regulations are required. A comprehensive risk assessment will require input from multiple agencies. We suggest that this exercise be coordinated by an inter-ministerial committee consisting of MeitY, the Department of Science and Technology, the Ministry of Consumer Affairs, the Ministry of External Affairs, the Ministry of Communications, the Ministry of Agriculture, and the Ministry of Health and Family Welfare, amongst others. Alternatively, the Parliamentary Standing Committee on Communications and Information Technology can coordinate this process. The proposed AISI for India should also be involved in conducting these risk assessments. Identify gaps in existing laws: A comprehensive gap analysis is essential to identifying areas where new regulations are required to address AI risks. We suggest an approach that identifies areas where: (1) no additional regulations are required (for example, in relation to job losses or superintelligence); (2) clarifications or targeted amendments would suffice (for example, in relation to deepfakes); and (3) new regulations are required to address market failures (for example, obligations with respect to transparency or rights in relation to algorithmic discrimination). This comprehensive analysis should involve participation from the Ministry of Consumer Affairs and the Department of Legal Affairs, and may be supplemented by research from academic institutes, law firms, think tanks, and the proposed AISI. Encourage self-regulation: Based on India’s larger strategic goals, we recommend that India adopt a light-touch, voluntary, and principle-based approach to AI regulation, at least for now.

Self-regulation is a good starting point because it allows regulators to adapt to the pace of innovation, develop technical expertise by collaborating with the private sector, reduce administrative costs, and develop baseline industry norms before introducing other compliance measures. To start with, we suggest that companies operating in India adopt voluntary commitments in relation to the safety and testing of AI systems, misinformation, privacy, and security. We also see an important role for specialized industry bodies in developing and adopting voluntary codes for different market actors (advanced models, enterprise companies, startups, etc.) and sectors (healthcare, finance, education, etc.). Empower the government to address AI risks: We recommend that legal provisions be introduced to enable the government to swiftly address market failures, now or in the future. This would include measures to:

(a) enable the government to notify key entities involved in the AI value chain based on their evolving roles and responsibilities, and to clarify to whom specific regulations apply;

(b) empower the government to create new rights for citizens or impose legal obligations on any class of entities if there is evidence of market failure (for example, in relation to testing, transparency, audits, human oversight); and

(c) develop robust enforcement and consumer protection mechanisms that are designed to address the specific AI risks for which there is evidence of harm. Adopt a “whole of government” approach: AI governance requires a multi-faceted approach. Broadly, we recommend a three-part framework:

(a) The relevant sectoral ministries, departments, and regulators (such as the Reserve Bank of India (RBI), the Telecom Regulatory Authority of India (TRAI), the Central Consumer Protection Authority (CCPA), the Ministry of Health, the Ministry of Education, and so on) should take the lead in developing, monitoring, and enforcing regulations.

(b) MeitY, as the nodal ministry responsible for technology regulation, should set baseline requirements across sectors and support the adoption of voluntary codes of conduct.

(c) Inter-agency coordination on AI governance should be handled by the Prime Minister’s Office (PMO) or the National Security Council Secretariat (NSCS), given their cross-agency mandate and strong leadership. Separately, a national AI Safety Institute (AISI) would help provide independent expertise and coordinate with the international network of AISIs. Consult with experts on AI regulation: While the government has initiated a series of public consultations on the proposed Digital India Act, there is a need for a more focused debate on AI regulation with a variety of experts. Specific issues such as bias, privacy, security, and copyright require an interdisciplinary approach and more substantive policy discussions with multi-stakeholder groups to shape future policy.

As India’s policymakers carefully mull the next steps on AI regulation, the brief pause in this continuing advance offers an opportunity to reflect and readjust, lest policymakers get trapped in path dependency and a mindless rush to regulate.

As the sentiment analysis in Part I illustrates, there is broad agreement that India should not adopt a comprehensive AI law. Some new regulations are warranted to address the risks of AI, but the scope of new rules and the ideal regulatory approach remains contentious.

Part II explains the intricate relationship between risk and harm, and the need for empirical evidence grounded in the local context, to inform India’s regulatory approach. We suggest focusing on five categories of AI risk and identifying high-risk use cases based on these risk vectors. We also offer some examples of areas where new regulations are required.

As the overview of different global approaches in Part III demonstrates, there is no one-size-fits-all approach to AI regulation. For a developing country like India, which is committed to reaping the full range of benefits from AI, we suggest self-regulation at least for the next six to twelve months (because binding regulation would entail significant costs, co-regulation is broken, and self-regulation is relatively efficient). However, as AI systems continue to evolve, the government should be empowered to prevent harm with clear legal mandates.

This paper suggests a few ways in which such provisions can be introduced. Regulation encompasses more than just laws. It also includes norms, standards, ethical practices, policy frameworks, institutional oversight, and soft laws. Therefore, we suggest a “whole of government” approach in which sectoral agencies, MeitY, and an inter-ministerial body collaborate in a dynamic fashion. An AISI, designed from the ground up keeping India’s unique needs in mind, can also supplement state capacity.

Finally, and this is important, the process by which AI regulations are framed must be both participative and inclusive. Not only should the data, models, and applications that power India’s AI ecosystem be representative of its culture, but so too should the policy frameworks that shape its future trajectory. To that end, it behooves the government to initiate a series of consultations on this topic before continuing its advance on AI regulation.

The authors would like to thank Rudra Chaudhuri and Anirudh Burman for their feedback on a draft of this paper, and the interviewees and reviewers for their thoughtful comments.

AI Laws and Regulations in India as of 2026

As over two decades navigating the intersections of technology, cybersecurity, and the law, I've witnessed India's digital journey from the early days of the IT Act to today's AI-driven ambitions. In February 2026, artificial intelligence is no longer a futuristic buzzword it's embedded in governance, finance, healthcare, and daily consumer experiences. Yet, the regulatory landscape remains a deliberate mix of existing statutes, soft guidelines, and forward-looking proposals. India has chosen a pro-innovation, techno-legal approach rather than rushing into a heavy-handed standalone AI law like the EU's AI Act.

This blog breaks down the current framework as of early 2026, key statutes that govern AI (directly or indirectly), emerging risks like deepfakes and synthetic media, and what lies ahead with initiatives like the Digital India Act and the AI (Ethics and Accountability) Bill.

In November 2025, MeitY released the India AI Governance Guidelines under the IndiaAI Mission. These are not enforceable law but serve as a foundational reference for responsible AI adoption. They emphasize a "light-touch" model that prioritizes innovation while mitigating harms through existing laws and voluntary measures.

They propose institutional mechanisms such as an AI Governance Group (inter-ministerial), a Technology & Policy Expert Committee, and the IndiaAI Safety Institute for monitoring and capacity building. A January 2026 white paper from the Office of the Principal Scientific Adviser further advances a "techno-legal" framework, advocating embedding compliance (legal safeguards + technical controls like watermarking and bias detection) directly into AI system design.

Practical takeaway: Organizations should align internal Responsible AI policies with these principles and map them to existing legal obligations. Self-certification, industry codes, and

standards (e.g., ISO/IEC 42001 adopted by BIS) are encouraged.

on 20th Feb 2026, MeitY notified The IT (Intermediary Guidelines and Digital Media Ethics Code) Amedment Rules, 2026 specifically targeting "synthetically generated information" (SGI) - content created or altered by AI/algorithms that appears authentic. The rules define SGI broadly and mandates:

- Significant social media intermediaries must require user declarations + deploy tech tools to verify & label synthetic media

- Platforms must use automated systems to block harmful synthetic content (CSAM, non-consensual deepfakes, fake events, explosive-related material, etc.)

This is one of the most proactive regulatory moves globally to fight misinformation, digital fraud, and AI-enabled harm while encouraging responsible innovation. India is leading here and the World may follow us ..

The DPDPA, with its Rules notified in 2025 (phased rollout until 2027), is highly relevant for AI systems processing personal data. Key implications:

Consent-centric regime → Explicit, informed consent is required for most processing, posing challenges for large-scale AI training on personal data (though exemptions exist for publicly available data and certain research).

Data Fiduciary obligations → Purpose limitation, data minimization, security safeguards, and breach notification apply to AI developers/deployers.

Rights of Data Principals → Access, correction, erasure, and grievance redressal can intersect with AI decision-making (e.g., explainability in automated decisions).

AI training on scraped or inferred data must be carefully evaluated for compliance. The Act's risk-based approach complements the Governance Guidelines.

Unfair trade practices and misleading advertisements - E.g., overhyping AI capabilities or opaque recommendation algorithms in e-commerce/finance.

Product liability provisions (Chapter VI) - Manufacturers, service providers, and sellers can be held liable for harm caused by defective AI products (including software/AI systems viewed as products).

The Central Consumer Protection Authority (CCPA) can investigate and penalize misleading AI claims. Courts are increasingly applying CPA to algorithmic harms.

Deepfakes are prosecuted under these technology-neutral provisions, often combined with IT Act charges. Evidentiary challenges persist under the Bharatiya Sakshya Adhiniyam (new Evidence Act), as hash values and digital signatures may not fully counter sophisticated manipulations.

Negligence - Developers/deployers owe a duty of care; breach (e.g., inadequate testing for bias or safety) leading to foreseeable harm can result in damages. Algorithmic opacity complicates proving causation and breach.

Product liability/strict liability - Limited application (primarily statutory via CPA); courts may impose where AI is deemed inherently dangerous.

Vicarious liability - Employers liable for employee misuse of AI tools; platform liability for user-generated synthetic content (subject to safe harbor).

Key challenge: Attribution in autonomous AI systems. No specific AI tort regime yet, so cases turn on facts, foreseeability, and control over the system.

Copyright Act → Ongoing consultations on AI training data; potential mandatory licensing or royalties for using copyrighted works (DPIIT proposals).

Sectoral regulations → RBI guidelines for AI in finance (responsible AI, bias mitigation); SEBI for markets; ICMR for health AI; etc.

Digital India Act (DIA): Expected to be a comprehensive overhaul/replacement of the IT Act. It is likely to introduce risk-based classifications for digital platforms and services, enhanced intermediary obligations, content accountability, and specific provisions for AI and emerging tech (including deepfake regulations and algorithmic transparency). Draft discussions have been ongoing; watch for public consultation in 2026.

Artificial Intelligence (Ethics and Accountability) Bill, 2025: A Private Member's Bill introduced in December 2025 in the Lok Sabha. It proposes a statutory Ethics Committee for AI, mandatory ethical reviews (especially for surveillance and high-risk systems), bias audits, developer obligations, restrictions on certain AI uses in law enforcement/employment, grievance mechanisms, and penalties up to ■5 crore. While not yet enacted, it signals growing parliamentary interest in binding accountability and could influence future policy or be incorporated into broader legislation.

Additional developments may include amendments for data portability, techno-legal tools (e.g., mandatory watermarking), and the operationalization of the IndiaAI Safety Institute.

India's approach is pragmatic: leverage existing tools, embed governance by design, and scale responsibly. As we move toward Viksit Bharat, clear, predictable rules will build trust and attract investment.

If you're dealing with AI compliance, deepfake incidents, or regulatory queries, feel free to reach out. Stay vigilant, innovate ethically, and let's shape a trusted AI ecosystem together.

(Disclaimer: This is an overview based on publicly available information as of February 2026. Laws evolve rapidly consult a qualified legal professional for specific advice.)

AI Regulation in India: Current State and Future Perspectives

Artificial intelligence (AI) presents big opportunities and potential risks for countries around the globe, and India is no exception. India has a vast, burgeoning high-tech labor force. The country also attracts millions of dollars in foreign direct investments, putting it on pace to become a major player in the global technology supply chain. With this growth, AI technologies are, and will, make their way into numerous Indian industries, such as healthcare, technology, the workforce, and education, forcing the Indian government to take steps toward regulating AI.

In recent publications, we discussed US President Joseph Biden's sweeping executive order on AI and the compromise reached on an AI act among representatives of the EU Council, European Parliament, and European Commission. Similar to these global counterparts, the Indian government recognizes AI's potential to have a sweeping impact on society, including negative impacts such as bias and privacy violations. In the last several years, India has introduced initiatives and guidelines for the responsible development and deployment of AI technologies, but there are currently no specific laws regulating AI in India.

The Indian government tasked the NITI Aayog, its apex public policy think tank, with establishing guidelines and policies for the development and use of AI. In 2018, the NITI Aayog released the National Strategy for Artificial Intelligence #AIForAll strategy, which featured AI research and development guidelines focused on healthcare, agriculture, education, "smart" cities and infrastructure, and smart mobility and transformation.

In February 2021, the NITI Aayog released Part 1 - Principles for Responsible AI, an approach paper that explores the various ethical considerations of deploying AI solutions in India, divided into system considerations and societal considerations. While the system considerations mostly deal with the overall principles behind decision-making, rightful inclusion of beneficiaries, and accountability of AI decisions, societal considerations focus on the impact of automation on job creation and employment. In August 2021, the NITI Aayog released Part 2 - Operationalizing Principles for Responsible AI, which focuses on operationalizing principles for responsible AI. The report breaks down the actions that need to be taken by both the government and the private sector, in partnership with research institutes, to cover regulatory and policy interventions, capacity building, incentivizing ethics by design, and creating frameworks for compliance with relevant AI standards.

The government of India also recently enacted a new privacy law, the Digital Personal Data Protection Act in 2023, which it can leverage to address some of the privacy concerns concerning AI platforms.

Additionally, India is a member of the Global Partnership on Artificial Intelligence (GPAI). The 2023 GPAI Summit was recently held in New Delhi, where GPAI experts presented their work on responsible AI, data governance, and the future of work, innovation, and commercialization. The GPAI website provides that "as a vital branch of the initiative, GPAI's Experts produce deliverables that can be integrated into Members' national strategies to ensure the inclusive and sustainable development of AI. Under the 2023 themes of climate change, global health and societal resilience, Experts worked to ensure that AI is used responsibly to address current challenges around the world. GPAI's Members, on the other hand, adopted the 2023 Ministerial Declaration, reaffirming their commitment to the trustworthy stewardship of AI in line with the OECD AI Principles, as well as their dedication to implementing those principles through the development of regulations, policies, standards and other initiatives. In doing so, they highlighted efforts to bridge the gap between theory and practice, and advance AI that is responsible, sustainable, and inclusive for all."

Other Indian agencies are also working on AI policies for the country, including the Ministry of Electronics and Information Technology, which has created committees on AI that have submitted reports on the development, safety, and ethical issues related to AI. The Bureau of Indian Standards, which is the national standards body of India, has also established a committee on AI that is proposing draft Indian standards for AI.

While the government of India has taken steps to regulate AI, its approach has mainly been one of pro-innovation with the development of policies and guidelines that acknowledge the ethical concerns and risks around the use of AI that may require the adoption of best practices. Given India's advantage of having a robust software development industry, this approach makes sense until the government formally enacts AI regulations.

The AI landscape in India is continuously evolving. Uncertainty, however, has not stopped both local and international interest and growth in this space. General Atomics Aeronautical Systems, Inc. announced a partnership with an India-based AI company, 114ai, to develop advanced technology for complex military systems. In September 2023, the US-based chip firm NVIDIA Corporation announced partnerships with Indian conglomerates Reliance Industries Ltd. and Tata group to develop cloud infrastructure and language models, wherein NVIDIA will provide the computing power required for building a cloud AI infrastructure platform.

An entity looking to enter the AI space in India should carefully consider the best legal route for such entry, whether through a joint venture, a strategic alliance, or a wholly owned subsidiary. Each route can be leveraged, and structures can be put in place depending on the level of investment and control required by the investing foreign entity. This is particularly important in a dynamic space like AI, where regulation is continuously evolving. Issues such as liability for harm caused, rights to intellectual property for AI systems, and privacy and data protection have not been fully fleshed out in regulations. Therefore, entities looking to enter this space should carefully consider the best legal and contractual protections.

Morgan Lewis's India initiative comprises more than 30 lawyers, several of whom are India qualified, and is led by lawyers from our Silicon Valley, Singapore, London, and Dubai offices. Collaborating across teams and offices, we advise our clients on a variety of business, regulatory, and governmental matters, from complex cross-border transactions and disputes to international tax planning and investment management. We also assist global companies with matters related to investments, expansion, and outsourcing as they enter into and further develop their business ventures in India. Contact the authors or any member of our India initiative for more information.

Corporate AI Governance

9 Key AI Governance Frameworks in 2025

With significant investment in AI development and more enterprises exploring private AI, it’s crucial to consider the importance of AI governance.

Technology as powerful as AI should be developed correctly. But at present, AI governance is primarily left to individual enterprises to adopt. This leaves enterprises and even countries at risk.

While legislation may lag behind, it’s prudent to build AI systems now that avoid compliance risks, reinforce fairness, in order to help organizations avoid future legal, ethical, and reputational challenges.

AI governance frameworks are structured systems of principles and practices that guide organizations in developing and deploying artificial intelligence (AI) in a responsible and compliant manner.

These frameworks aim to ensure that AI systems are ethically aligned, secure, transparent, and compliant with applicable regulations. Most AI governance frameworks emphasize fairness, accountability, and explainability, while remaining adaptable to differing ethical norms across domains and regions.

AI governance frameworks support organizational oversight of AI systems and provide a foundation for responsible AI adoption in regulated and high-impact environments.

Beyond addressing long-term concerns, AI governance frameworks also support immediate risk mitigation, including the prevention of biased outputs, data misuse, and privacy breaches.

A shared framework fosters consistent understanding of AI goals, limitations, and oversight responsibilities across teams.

AI governance frameworks demonstrate to internal and external stakeholders that the organization maintains formalized practices for AI-related security and compliance.

Across the various AI principles, orders, and standards—some of which are outlined in later sections—several common themes emerge in responsible AI governance.

: AI systems must remain under meaningful human control. Transparency : Users and regulators should be able to understand how an AI system generates outputs or decisions.

: Users and regulators should be able to understand how an AI system generates outputs or decisions. Accountability: There must be clearly defined responsibility for AI outcomes.

There must be clearly defined responsibility for AI outcomes. Safety : Systems should be secure, reliable, and resilient to failures or adversarial attacks.

: Systems should be secure, reliable, and resilient to failures or adversarial attacks. Fairness and non-discrimination : AI should be developed and applied to mitigate bias and support equitable treatment.

: AI should be developed and applied to mitigate bias and support equitable treatment. Privacy and data protection : AI must uphold individuals' data rights and comply with applicable data protection laws.

: AI must uphold individuals' data rights and comply with applicable data protection laws. Proportionality : Oversight and intervention should correspond to the potential impact of the AI system.

: Oversight and intervention should correspond to the potential impact of the AI system. Human-centric design: AI should support human well-being and align with fundamental rights.

Different frameworks emphasize different principles. For example, the UNESCO Recommendation on the Ethics of Artificial Intelligence places greater focus on environmental sustainability and gender equality.

The European Union's AI Act introduces a tiered, risk-based classification system, categorizing AI systems as unacceptable, high, limited, or minimal risk. This structure includes explicit prohibitions on certain high-risk uses and is not replicated in most other governance frameworks.

Legal and regulatory frameworks can help organizations develop AI in ways that are secure, ethical, and compliant with local and international laws. The following frameworks offer guidance on building secure, compliant, and ethical AI systems.

Because it bans certain uses (for example, social scoring) and imposes strict controls on high-risk applications, such as healthcare and financial services, AI deployment within the EU must comply with this governance framework or risk significant fines.

The UK pro-innovation AI framework is a non-statutory whitepaper that sets out five core principles: fairness, transparency, accountability, safety, and contestability.

It emphasizes a flexible, context-driven approach, making it useful for enterprises seeking regulatory alignment without heavy compliance burdens.

An initial Executive Order on AI was issued during the Biden administration, Executive Order 14110. In 2025, a new order was introduced, Executive Order 14179, 'Removing Barriers to American Leadership in Artificial Intelligence.'

This national initiative guides how federal agencies oversee the use of AI in civil rights, national security, and public services.

The updated order emphasizes that AI development must maintain U.S. leadership in AI while remaining free from ideological bias or engineered social agendas.

The NIST AI Risk Management Framework (AI RMF) offers structured, risk-based guidance for building and deploying trustworthy AI. It is widely adopted across industries and favored for its practical, adaptable advice across four principles: govern, map, measure, and manage.

Published by the White House Office of Science and Technology Policy in 2022, the AI Bill of Rights is not a legally binding document. However, it served as an early framework introducing key principles, including data privacy, fairness, and human fallback, to guide the responsible design and use of AI.

It laid the groundwork for subsequent regulatory actions, including the Executive Order on AI issued by the Biden administration in 2023.

State-level regulations are evolving rapidly and are enforceable locally. These laws regulate the use of AI in areas such as hiring, surveillance, and consumer protection.

For example, Colorado's "Concerning Consumer Protections in Interactions with Artificial Intelligence Systems Act" (Colorado AI Act or CAIA) prohibits algorithmic discrimination in high-risk AI systems, including those used in healthcare, recruitment, and education.

In California, proposed legislation such as the Automated Decision Systems Accountability Act seeks to increase transparency and accountability in the use of AI for consequential decisions, particularly in employment and housing.

The OECD AI Principles are non-binding international guidelines that promote human-centric, transparent, and accountable AI development.

First established in 2019 and updated in 2024, the principles encourage governments to regularly review and adapt their AI-related policies and regulatory frameworks. Broad adoption has been seen globally, especially in OECD member countries.

Originally known as the Recommendation on the Ethics of Artificial Intelligence, the UNESCO framework represents the first global standard on AI ethics. It is voluntarily adopted by United Nations member states.

In line with UNESCO's broader goals of promoting peace and human rights, it encourages the development of inclusive, sustainable, and ethical AI.

Established by the G7 nations, the G7 Code of Conduct is a voluntary commitment that outlines best practices for the safe and responsible development of foundation models and generative AI.

The code works with the overarching G7 Action Plan, which promotes the human-centered adoption of secure and trustworthy AI across member nations.

Selecting the right governance model ensures AI initiatives align with business objectives while maintaining security, ethics, and compliance.

Establish policies that govern data quality, privacy, model development, deployment, and ongoing monitoring. These operational rules ensure that high-level AI principles are consistently translated into everyday practices across teams.

AI governance requires designated ownership. This may include assigning data stewards to oversee data quality, AI leads to manage implementation, and compliance officers to manage regulatory risks, as well as incorporating mechanisms for end-user feedback to support continual improvement.

Governance structures should support the broader goals of your organization. For example, in sectors such as healthcare or finance, accuracy and transparency must be embedded from the outset, using tools such as execution graphs, explainability layers, and audit logs to trace and validate AI decisions.

Security and regulatory alignment should be foundational. In sensitive or regulated environments, opt for private deployments using virtual private clouds (VPCs) or on-premises infrastructure to maintain control over data access, model behavior, and system integrity.

AI governance must span the entire lifecycle of an enterprise's systems — from initial design through deployment, monitoring, and continuous improvement. The following guidance outlines how organizations can operationalize an AI governance framework in practice.

During the design phase, teams should map the AI system's intended functions and enterprise use cases before development begins. It is essential to establish mechanisms for traceability, regulatory compliance, and model performance standards early in the process.

Simulating real-world scenarios using representative training data can help identify edge cases and verify that the system performs reliably under varied conditions prior to deployment.

In the deployment phase, AI systems should be implemented in secure environments, such as VPCs or on-premises infrastructure. Integrating access controls, output validation mechanisms, and audit logging enables oversight, accountability, and traceability from the outset. Organizations should promote internal reporting of risks and may consider structured incentives to support the identification of issues that conflict with responsible AI practices.

During the monitoring phase, real-time observability should be maintained using monitoring dashboards and structured user feedback collection. Automated retraining pipelines can support model adaptation as data distributions or use cases evolve. Continuous monitoring for model drift, emerging bias, or operational risks is critical to ensuring sustained compliance with governance expectations.

Ongoing risk management requires tools that proactively detect and respond to threats such as bias, data drift, and performance degradation. Enterprises should conduct regular audits against established governance frameworks to verify that AI systems continue to meet compliance and risk management requirements.

For explainability and observability, transparency should be embedded using techniques such as execution graphs, confidence scores, and traceable reasoning chains. In regulated industries, both technical and non-technical stakeholders must be able to understand and assess AI-driven outputs to support accountability and trust.

A growing number of enterprises are recognizing the need for formalized, yet adaptive, governance frameworks to manage AI risk and maintain stakeholder trust.

This proportionate approach allows teams to apply safeguards based on potential impact, while still maintaining the pace of innovation.

ENSURING A NATIONAL POLICY FRAMEWORK FOR ARTIFICIAL INTELLIGENCE

By the authority vested in me as President by the Constitution and the laws of the United States of America, it is hereby ordered:

Section 1. Purpose. United States leadership in Artificial Intelligence (AI) will promote United States national and economic security and dominance across many domains. Pursuant to Executive Order 14179 of January 23, 2025 (Removing Barriers to American Leadership in Artificial Intelligence), I revoked my predecessor's attempt to paralyze this industry and directed my Administration to remove barriers to United States AI leadership. My Administration has already done tremendous work to advance that objective, including by updating existing Federal regulatory frameworks to remove barriers to and encourage adoption of AI applications across sectors. These efforts have already delivered tremendous benefits to the American people and led to trillions of dollars of investments across the country. But we remain in the earliest days of this technological revolution and are in a race with adversaries for supremacy within it.

To win, United States AI companies must be free to innovate without cumbersome regulation. But excessive State regulation thwarts this imperative. First, State-by-State regulation by definition creates a patchwork of 50 different regulatory regimes that makes compliance more challenging, particularly for start-ups. Second, State laws are increasingly responsible for requiring entities to embed ideological bias within models. For example, a new Colorado law banning "algorithmic discrimination" may even force AI models to produce false results in order to avoid a "differential treatment or impact" on protected groups. Third, State laws sometimes impermissibly regulate beyond State borders, impinging on interstate commerce.

My Administration must act with the Congress to ensure that there is a minimally burdensome national standard — not 50 discordant State ones. The resulting framework must forbid State laws that conflict with the policy set forth in this order. That framework should also ensure that children are protected, censorship is prevented, copyrights are respected, and communities are safeguarded. A carefully crafted national framework can ensure that the United States wins the AI race, as we must.

Until such a national standard exists, however, it is imperative that my Administration takes action to check the most onerous and excessive laws emerging from the States that threaten to stymie innovation.

Sec. 2. Policy. It is the policy of the United States to sustain and enhance the United States' global AI dominance through a minimally burdensome national policy framework for AI.

Sec. 3. AI Litigation Task Force. Within 30 days of the date of this order, the Attorney General shall establish an AI Litigation Task Force (Task Force) whose sole responsibility shall be to challenge State AI laws inconsistent with the policy set forth in section 2 of this order, including on grounds that such laws unconstitutionally regulate interstate commerce, are preempted by existing Federal regulations, or are otherwise unlawful in the Attorney General's judgment, including, if appropriate, those laws identified pursuant to section 4 of this order. The Task Force shall consult from time to time with the Special Advisor for AI and Crypto, the Assistant to the President for Science and Technology, the Assistant to the President for Economic Policy, and the Assistant to the President and Counsel to the President regarding the emergence of specific State AI laws that warrant challenge.

Sec. 4. Evaluation of State AI Laws. Within 90 days of the date of this order, the Secretary of Commerce, consistent with the Secretary's authorities under 47 U.S.C. 902(b), shall, in consultation with the Special Advisor for AI and Crypto, the Assistant to the President for Economic Policy, the Assistant to the President for Science and Technology, and the Assistant to the President and Counsel to the President, publish an evaluation of existing State AI laws that identifies onerous laws that conflict with the policy set forth in section 2 of this order, as well as laws that should be referred to the Task Force established pursuant to section 3 of this order. That evaluation of State AI laws shall, at a minimum, identify laws that require AI models to alter their truthful outputs, or that may compel AI developers or deployers to disclose or report information in a manner that would violate the First Amendment or any other provision of the Constitution. The evaluation may additionally identify State laws that promote AI innovation consistent with the policy set forth in section 2 of this order.

Sec. 5. Restrictions on State Funding. (a) Within 90 days of the date of this order, the Secretary of Commerce, through the Assistant Secretary of Commerce for Communications and Information, shall issue a Policy Notice specifying the conditions under which States may be eligible for remaining funding under the Broadband Equity Access and Deployment (BEAD) Program that was saved through my Administration's "Benefit of the Bargain" reforms, consistent with 47 U.S.C. 1702(e)-(f). That Policy Notice must provide that States with onerous AI laws identified pursuant to section 4 of this order are ineligible for non-deployment funds, to the maximum extent allowed by Federal law. The Policy Notice must also describe how a fragmented State regulatory landscape for AI threatens to undermine BEAD-funded deployments, the growth of AI applications reliant on high-speed networks, and BEAD's mission of delivering universal, high-speed connectivity.

(b) Executive departments and agencies (agencies) shall assess their discretionary grant programs in consultation with the Special Advisor for AI and Crypto and determine whether agencies may condition such grants on States either not enacting an AI law that conflicts with the policy of this order, including any AI law identified pursuant to section 4 or challenged pursuant to section 3 of this order, or, for those States that have enacted such laws, on those States entering into a binding agreement with the relevant agency not to enforce any such laws during the performance period in which it receives the discretionary funding.

Sec. 6. Federal Reporting and Disclosure Standard. Within 90 days of the publication of the identification specified in section 4 of this order, the Chairman of the Federal Communications Commission shall, in consultation with the Special Advisor for AI and Crypto, initiate a proceeding to determine whether to adopt a Federal reporting and disclosure standard for AI models that preempts conflicting State laws.

Sec. 7. Preemption of State Laws Mandating Deceptive Conduct in AI Models. Within 90 days of the date of this order, the Chairman of the Federal Trade Commission shall, in consultation with the Special Advisor for AI and Crypto, issue a policy statement on the application of the Federal Trade Commission Act's prohibition on unfair and deceptive acts or practices under 15 U.S.C. 45 to AI models. That policy statement must explain the circumstances under which State laws that require alterations to the truthful outputs of AI models are preempted by the Federal Trade Commission Act's prohibition on engaging in deceptive acts or practices affecting commerce.

Sec. 8. Legislation. (a) The Special Advisor for AI and Crypto and the Assistant to the President for Science and Technology shall jointly prepare a legislative recommendation establishing a uniform Federal policy framework for AI that preempts State AI laws that conflict with the policy set forth in this order.

(b) The legislative recommendation called for in subsection (a) of this section shall not propose preempting otherwise lawful State AI laws relating to:

(ii) the functions of the Director of the Office of Management and Budget relating to budgetary, administrative, or legislative proposals.

(c) This order is not intended to, and does not, create any right or benefit, substantive or procedural, enforceable at law or in equity by any party against the United States, its departments, agencies, or entities, its officers, employees, or agents, or any other person.

AI Governance at a Crossroads: America's AI Action Plan and its Impact on Businesses

This post is the first in a series of commentaries on recent developments in the sphere of AI governance and innovation. In this post, we examine the recent America's AI Action Plan and its shifting priority towards a philosophy of AI deregulation. The Plan advocates for innovation over caution, while this uncertain and ever-changing legislative landscape will demand a further shift towards the private sector to manage AI risks related to ethics, responsibility, and governance in the absence of legal mandates.

In the United States, the first significant federal action on AI governance came during President Trump's first administration, with the 2019 Executive Order on Maintaining American Leadership in Artificial Intelligence. The order focused on promoting U.S. competitiveness by directing federal agencies to prioritize AI R&D;, workforce training, and international cooperation. It framed AI largely as an engine for innovation and economic growth but offered little in the way of oversight or guardrails.

Building on this foundation, President Biden's 2023 Executive Order on Safe, Secure, and Trustworthy AI shifted the tone considerably. This order signaled that AI was not just an economic driver, but a matter of civil rights and national security, in need of guidelines. The order emphasized the importance of ensuring AI systems are safe and equitable while encouraging innovation and safeguarding America's competitive edge. The order directed agencies to adopt the NIST AI Risk Management framework as the standard for AI oversight, and agencies began developing principles and best practices for use of AI within their domains, covering areas such as data privacy, workplace monitoring, and protection against algorithmic discrimination.

At the same time, the absence of comprehensive federal legislation opened the door to a growing wave of state-level activity, with now over 480 enacted bills referencing "artificial intelligence". States such as New York, Colorado, and California introduced their own AI-related laws and proposals, ranging from hiring bias audits to algorithmic discrimination bills. These efforts, combined with a lack of clear federal law, created a fragmented and inconsistent regulatory landscape making it challenging for companies operating across jurisdictions to anticipate potential liabilities and comply with competing standards.

When Trump came into office for his second term, President Biden's Executive Order was lifted, with the intention of making unhindered innovation a priority. Subsequent efforts to streamline oversight, including a proposed moratorium on state-level AI regulation, did not gain traction. As a result, by 2025 the U.S. regulatory environment remained characterized by gaps, overlaps, and uncertainties. This unsettled foundation set the stage for the development of a more comprehensive national framework through the AI Action Plan.

On July 23, 2025, the White House released America's AI Action Plan, outlining the Trump administration's federal AI strategy in its effort to win the AI race against its closest competitor,

China. Alongside this Plan, the administration also released a series of three executive orders that align with this national promotion of AI innovation:

Preventing Woke AI in the Federal Government, Accelerating Federal Permitting of Data Center Infrastructure, and Promoting the Export of the American AI Technology Stack.

These executive orders reflect a broader shift in U.S. policy, where the central driver of AI governance focuses less on ethical safeguards and more on maintaining a competitive edge on the world stage. The Plan frames AI innovation as a domain of economic dominance and national security in which the AI race is presented not simply as a matter of technological progress, but as essential to protecting U.S. influence in international markets and strategic alliances.

At its core, this national vision adopts a pro-innovation approach to AI regulation, aiming to ensure that the United States remains the commanding force in global technology norms and standards. This approach involves prioritizing technological growth and development over precautionary oversight, while also sending a strong message to both allies and competitors alike about American intentions. More broadly, this shift in regulatory strategy underscores the volatility of the American regulatory landscape as shaped by changes in administration.

The Plan advocates for reduced federal regulation of AI systems and use cases, which, as we argue further, would place greater responsibility on corporate boards, officers, and senior management teams, to self-manage and mitigate AI risks within their own organizations. This deregulatory stance is clear from the outset of the document, as the first section of the first pillar is titled “Remove Red Tape and Onerous Regulation.” Its goal is to allow the private AI sector to flourish free of what it describes as the “bureaucratic red tape” imposed under the Biden administration. Central to this vision is rolling back existing policies and frameworks such as Biden’s Executive Order 14110 described previously. The Plan further recommends withholding federal funding from states with “burdensome” AI regulatory frameworks, while still allowing states to pass laws so long as they do not obstruct technological progress. This approach sets the stage for a fragmented and largely deregulated AI landscape across the country.

The conclusion of this section lists recommended policy actions to support deregulation, beginning with an investigation into current policies that impede AI innovation and adoption led by the Office of Science and Technology Policy. Additional recommendations task the Office of Management and Budget and the Federal Communications Commission with revising regulations that conflict with these priorities and aligning federal funding with states’ regulatory climates. Although many of these recommendations lack a clear timeline or funding plan, raising questions about the pace and extent of implementation, it is apparent that greater changes in the U.S. regulatory landscape both at the federal and state level should be expected soon. At the state level, legislative actions have the potential to be swayed by bipartisan alignment with either this new innovation-driven approach from the Trump administration or greater attention to values based principles set forth under the Biden administration.

Subsequent sections of the Plan re-emphasize the administration’s pro-innovation stance while also maintaining the role of the federal government to manage risks involving technological standards and model evaluations across a variety of government departments. According to a recent Stanford HAI report, this strategy can be thought of as a shift from prescriptive regulation towards evidence-based policymaking. While there are calls for greater security and risk management frameworks, the Plan notably lacks federal commitment to mitigating ethical concerns around bias, transparency, and data privacy that characterized the Biden Administration’s approach. As in the opening section, many of the following proposed actions also lack specificity and concrete pathways for implementation.

Notably, just three days after the Trump administration's Plan was published, the Chinese government released a responding Global AI Governance Action Plan, emphasizing fairness, cooperation, national sovereignty, and the service of humanity at the forefront of its AI goals. This back-and-forth of policy initiatives heightens the AI race between the United States and China that characterizes much of the Trump Administration's regulatory frameworks. As explored in one of our previous posts, these tensions have been building gradually as policymakers across the globe are trying to find the right way to balance innovation with the need to protect human values and guard against risks.

Within our research, we have established the immature and volatile AI regulatory landscape as a major motivator for corporate action. As the federal government shifts from prescriptive regulation to "light-touch" governance and state-level legislators are incentivized to move in a similar direction, the responsibility increasingly falls to the private sector to establish and maintain its own standards of governance as they are not able to rely on governments to provide effective risk-management frameworks. This raises a critical question about what motivates companies to spend the time and resources on ethical AI oversight systems without external compliance motivators.

National deregulation does not eliminate the risks that AI poses to businesses, including reputational, operational, financial, strategic, and data security risks that remain significant and that require ethical AI frameworks to mitigate them. As we've explored in a previous blogpost, there are numerous examples of misaligned AI systems and governance structures that have led to a variety of negative impacts across companies. Therefore, we believe that the implementation of ethical AI governance practices is more than just a moral obligation, but rather a necessity for companies to avoid a variety of risks, especially financial and reputational risks. Moreover, existing legal precedents, such as the Caremark progeny of cases that we discussed in a previous blogpost, maintain legal liability for boards and executives to provide sufficient oversight of risks tied to "mission critical" company operations that increasingly involve AI systems. This can be achieved by implementing necessary guardrails within organizations to ensure safe, secure, and trustworthy AI development and deployment.

Throughout our research we've developed an AI governance framework – the Boundaries of Tolerance Framework – that revolves around determining the boundaries of what risks, ethical considerations, and outcomes each specific organization will tolerate in their AI development and deployment. Our research uses AI and business ethics as mediums to achieve this goal amongst an evolving national regulatory landscape.

Ultimately, the shifting nature of U.S. national regulation places increasing pressure on corporations to self-manage their AI practices and to develop robust AI-management frameworks. Within the environment of evolving state-level legislative frameworks and federal executive branch policies it is critical for companies to remain informed and hold themselves accountable to manage not only potential legal compliance, but also to manage broader ethical, operational, and reputational risks.

Jeffrey Saviano: Business AI Ethics Initiative Leader; Edmond & Lily Safra Center for Ethics, Harvard University | jeffreysaviano@fas.harvard.edu

Jonathan Hack: Director of Content & Strategy; Edmond & Lily Safra Center for Ethics, Harvard University | jhack@fas.harvard.edu

The Business AI Ethics research team is part of the Edmond & Lily Safra Center's ongoing effort to promote the application of ethics in practice. Their research assists business leaders in examining the promise and challenges of AI technologies through an ethical lens. Views expressed in these posts are those of the author(s) and do not imply endorsement by ELSCE.

The 20 Biggest AI Governance Statistics and Trends of 2025

AI governance refers to the structures and controls used to manage risks and ensure accountability in GenAI deployments across enterprises.

Only 25% of organizations have fully implemented AI governance programs, revealing a sharp gap between awareness and execution.

Just 27% of boards have formally incorporated AI governance into committee charters, showing limited leadership integration despite growing risk.

A staggering 97% of AI-related breach victims have been shown to lack proper access controls, highlighting enforcement, not policy, as the most significant vulnerability.

Nearly all organizations (98%) expect AI governance budgets to rise, signaling a clear shift from reactive compliance to proactive operational investment.

In general, reliable statistics begin with verified sources and transparent validation steps. Each data point in this article is drawn primarily from published articles that document measurable enterprise trends. These include datasets on GenAI monitoring, oversharing prevention, and compliance automation, all derived from aggregated enterprise behavior observed through real-time controls. Supporting context comes from a limited set of authoritative frameworks, the NIST AI Risk Management Framework (AI RMF) and ISO/IEC 42001:2023, both recognized globally for defining auditability and performance metrics in AI systems.

The 20 selected AI governance statistics were chosen because they capture the most critical changes in GenAI governance for 2025, linking security incidents, adoption maturity, and privacy trends that directly shape enterprise readiness. Each figure highlights either a measurable risk or a governance opportunity aligned with emerging standards. Data sources include McKinsey & Company's State of AI 2025 report, Verizon's 2025 Data Breach Investigations Report, IBM's Cost of a Data Breach 2025, Cisco's Data Privacy Benchmark 2025, and the FBI's Internet Crime Report 2024. Naming these sources helps readers understand the analytical approach and reinforces credibility across sectors like strategy, compliance, and cybersecurity.

Next, it is vital to highlight that every statistic included here has been checked for traceability and recency, ensuring that it reflects developments through 2025. Also, terminology remains consistent across all sections. For instance, definitions of "traceability," "continuous monitoring," and "explainability" follow the usage in the National Institute of Standards and Technology (NIST) and International Organization for Standardization (ISO) documentation. Before publication, all metrics were reviewed for internal coherence, with cross-checks to ensure that none contradict established AI governance principles. The result is a rigorously validated dataset, with each metric cross-checked against NIST AI RMF traceability and explainability criteria to ensure methodological integrity. This approach connects industry-wide LLM governance standards to practical enterprise evidence.

Accurate AI governance statistics are essential for transforming complex risks into measurable, actionable insights that drive informed decision-making and enterprise resilience.

Sound statistics create a shared base of knowledge for decision-making. Leadership teams gain visibility into risk exposure, compliance maturity, and control performance across the enterprise. By quantifying governance performance across departments, organizations turn abstract discussions into targeted actions that strengthen executive trust and board accountability. This clarity helps justify resource allocation for remediation instead of relying on assumptions. Quantitative insights also connect governance to measurable business outcomes, such as fewer audit delays and noticeable improvements in incident detection and response times. Frameworks like NIST AI RMF and ISO/IEC 42001 recommend this evidence-based approach for continuous improvement.

Accurate statistics simplify audits by converting compliance into measurable evidence. Instead of collecting static documents, enterprises can present live metrics showing real-time AI monitoring, policy hit rates, and access violations. This shift moves audits from a reactive to an ongoing assurance approach. Reliable data also confirms that explainability and traceability are not theoretical, but are applied daily. In regulated sectors such as finance and healthcare, these figures become essential for certification under ISO/IEC 42001 and alignment with the EU AI Act. They also show auditors that oversight mechanisms are continuous and responsive.

Scalability depends on knowing where risk accumulates. Accurate statistics pinpoint which personas, tools, or datasets create the most governance friction. Such insights allow teams to tighten access control and monitoring without halting innovation. Real-time numbers expose oversharing trends early, enabling rapid mitigation before escalation. They also guide staged GenAI deployment, starting with high-risk environments and expanding as controls mature. Continuous visibility into metrics like leakage prevention and policy enforcement ensures that scaling GenAI systems remains safe and sustainable. Over time, these statistics fuel predictive governance, allowing users to identify weak spots before incidents occur. Data-informed oversight thus becomes the foundation of resilience, supporting both operational agility and long-term trust.

The information presented in this section is classified into five categories. It will demonstrate that widespread GenAI adoption across enterprises has outpaced governance maturity, exposing organizations to measurable risks and underscoring the urgent need for disciplined controls, clear KPIs, and continuous oversight.

The National Association of Corporate Directors' (NACD)'s 2025 survey reveals that board-level oversight of AI is increasing, but governance integration remains limited. While 62% of boards now hold regular AI discussions, only 27% have formally added AI governance to their committee charters. Most boards still focus on education and risk awareness rather than embedding AI oversight into core operations. NACD concludes that boards are at an inflection point, transitioning from awareness to strategic, structured governance of AI within corporate decision-making.

According to a Gartner 2025 poll of over 1,800 executive leaders, 55% of organizations reported having an AI board or dedicated oversight committee in place. Establishing these committees represents a tangible shift from informal sponsorship to formal governance structures. Firms with dedicated oversight bodies are more likely to integrate AI risk monitoring, stakeholder accountability, and continuous review into their operating model.

In McKinsey & Company's The state of AI survey, only 28% of organizations said the CEO takes direct responsibility for AI governance oversight, while just 17% report that their board does. This data indicates a governance gap at the highest leadership levels, which correlates with slower value creation from GenAI programs. High-performing firms often attribute this oversight to senior leadership that takes clear accountability for ethical, traceable, and safe enterprise AI deployment.

McKinsey finds that tracking explicit GenAI KPIs remains uncommon, even though it correlates most strongly with long-term business and compliance impact. Without KPI discipline, teams struggle to compare use cases, tune guardrails, or justify investment in security and governance.

Establishing measurable goals, such as policy hit rate, data-leakage rate, and audit latency, turns governance from reactive to proactive management.

According to McKinsey's report, revenue also increases across most business units using GenAI, ranging from 51% to 70%. Earnings before interest and taxes (EBIT) contribution remains modest overall, however. Estimates across multiple analyses suggest that roughly 15-20% of firms report any profit-level (or EBIT) impact from GenAI, confirming that adoption has outpaced monetization. The gap between high expectations and measured EBIT impact often reflects weak instrumentation, fragmented data governance, and limited policy enforcement. Strengthening governance mechanisms, particularly metrics, roles, and accountability, accelerates both time-to-value and regulatory confidence. A low EBIT yield underscores a governance execution gap, enterprises often deploy GenAI tools without integrated performance tracking, meaning financial outcomes are disconnected from governance oversight and control maturity.

Organizations that lag in practices, such as roadmaps, change management, training, and KPI tracking, see slower value and greater risk exposure. Governance maturity is a leading indicator for sustainable scale. McKinsey links program maturity to clear governance milestones, such as defined model-review cycles and documented accountability. Treating GenAI as an enterprise-wide governance program, not just a technical rollout, improves resilience and operational alignment.

A recent Gartner, Inc. survey reports that 45% of organisations with high AI maturity keep their AI initiatives live for at least three years, compared to only 20% among lower-maturity peers. One main differentiator is governance, and not just governance in name, but dedicated structures, leadership accountability, and lifecycle oversight. Projects that persist, indicate embedded governance, model versioning, monitoring, and change-control practices rather than one-off pilots. This reinforces that governance maturity correlates with sustainable AI value and risk management.

Trustmarque's 2025 AI Governance Report reveals that governance maturity remains minimal, with fewer than one in ten organizations integrating AI risk and compliance reviews directly into development pipelines.

According to IBM's 2025 Cost of a Data Breach Report, 13% of organizations reported breaches involving AI models or applications. Among these, 97% said they had no proper AI access controls in place. This statistic shows that the governance gap isn't just about policy; it's about enforcement, role clarity, and technical gate-keeping. Without effective controls over how and when AI systems access data, companies remain exposed, even if they have drafted policies.

By highlighting that nearly all organizations hit by AI-related incidents lacked access controls, the finding shifts the focus from having a policy to whether you are executing it. By contrast, industry leaders are moving toward policy-based access control and immutable audit trails, approaches that enforce contextual authorization and preserve verifiable evidence of every system interaction.

Trustmarque's report notes that essential governance infrastructure, such as dataset version control, documentation standards, and audit trails, remains underdeveloped across most industries.

According to the IBM Cost of a Data Breach Report, 63% of organizations experiencing a breach did not have a formal AI governance policy in place. Effective governance bridges compliance and resilience by ensuring consistent decision logging and accountability. This reinforces that governance investments directly contribute to measurable risk reduction and business continuity.

A 2025 AuditBoard research study, From blueprint to reality, found that only one in four organizations have fully operational AI governance, despite widespread awareness of new regulations. Most firms have drafted policies but struggle to turn them into daily practice. The barriers include unclear ownership, limited expertise, and resource constraints. The report

concludes that effective AI governance is now a test of execution, not just writing policy.

Incidents range from compliance exceptions to quality errors and data exposure events. According to the McKinsey state of AI report, nearly half of organizations encountered measurable governance or ethical lapses linked to GenAI projects. This figure shows that risk is no longer theoretical and that controls must extend beyond model outputs to include access, context, and purpose limits. Continuous AI monitoring and explainability reduce surprise events and speed remediation.

Clear rules create stable operating conditions and reduce ambiguity in AI deployments. The Cisco 2025 Data Privacy Benchmark Study, The Privacy Advantage, emphasizes that enterprises integrating privacy governance into AI oversight experience higher operational consistency and stakeholder trust. This sentiment suggests that governance maturity is increasingly a market differentiator. Aligning GenAI controls with privacy-by-design builds a durable advantage.

Privacy programs pay off in trust, sales enablement, and reduced friction with regulators. The Cisco study reports that nearly all organizations achieving privacy maturity also show advanced governance practices, such as continuous audits and automated policy enforcement. For GenAI, privacy posture and governance posture are converging disciplines. Investment here directly supports safer scale and transparent accountability.

The Verizon Business 2025 Data Breach Investigations Report highlights that structured reporting standards now extend to AI-related incidents. Organizations aligning their incident classification and governance metrics with frameworks such as ISO/IEC 42001 and the EU AI Act gain stronger visibility and accountability. Improved transparency enables executives to connect technical risks with policy outcomes. It also allows for continuous benchmarking, turning compliance from a reactive function into an operational strength.

The 2024 IAPP Governance Survey found that only 28% of organizations have formally defined oversight roles for AI governance, highlighting persistent uncertainty about who owns responsibility for compliance, ethics, and model accountability. Most companies still distribute AI governance tasks across compliance, IT, and legal teams without a unified structure. This fragmented approach limits visibility into AI risk and slows the creation of consistent guardrails. Organizations that centralize AI oversight show faster alignment with frameworks like ISO/IEC 42001 and the EU AI Act.

IAPP’s AI Governance Profession Report 2025 states that 77% of surveyed organizations say they are actively building or refining AI governance programs, and that percentage rises to nearly 90% for organizations already using AI. This high number suggests that governance is widely recognised as a strategic imperative rather than an after-thought. Even so, many of these programs are still in early stages, with firms grappling with staffing, metrics, and accountability. This underscores that adoption of governance frameworks often lags behind AI deployment.

OneTrust’s The 2025 AI-Ready Governance Report found that 98% of organisations expect budgets for AI governance technology and oversight to increase substantially in the near term. This almost universal expectation shows that businesses recognise the material nature of AI governance investments, not just as compliance cost centres, but as enablers of scalable, trusted AI operations. With rapidly evolving regulation, model risk, and ethical concerns, organisations are beginning to allocate resources proactively rather than reactively.

Pacific AI’s 2025 AI Governance Survey found that 75% of organizations have established AI usage policies, yet only 36% have adopted a formal governance framework. Having a policy is an important first step, but the absence of a broader framework means many organizations lack consistent roles, controls, monitoring, and enforcement. This gap highlights how governance

maturity still lags behind policy creation. Teams responsible for governance must move beyond “we have a policy” toward “we have structured oversight, roles, KPIs and continual review.”

Explore the Knostic LLM Data Governance Playbook to move from awareness to action. Map your governance blind spots in under 20 minutes using Knostic’s built-in readiness checklist to assess your organization’s current governance maturity, then apply the diagnostic scorecard to benchmark real-time monitoring, policy coverage, and compliance posture. Together, these tools provide a practical path toward safer, scalable, and audit-ready GenAI deployment.

Accurate statistics help leadership teams identify specific governance gaps, prioritize investments, and connect AI risks to measurable business outcomes like audit readiness and incident response speed.

Prompt injection is the top risk for GenAI apps, requiring enterprises to enforce strict input validation and continuous monitoring to prevent instruction hijacking.

While nearly 80% of companies use GenAI in at least one function, fewer than 20% track key performance indicators, and only 17% report a meaningful impact on EBIT.

How to implement AI governance best practices in 2025

Organizations worldwide are rapidly integrating artificial intelligence into their operations, fundamentally transforming how teams work and make decisions. This technological shift brings unprecedented opportunities for productivity and innovation, but also introduces complex challenges around ethics, compliance, and risk management.

As AI systems become more sophisticated and autonomous, the need for robust governance frameworks has never been more critical. Workers using generative AI reported saving only 5.4% of their work hours, translating to just a 1.1% increase in overall productivity when factoring in all workers. Without proper oversight, AI implementations can lead to biased outcomes, regulatory violations, data breaches, and erosion of stakeholder trust — consequences that can devastate both reputation and bottom line.

The most successful enterprises recognize that AI governance isn't just about risk mitigation; it's about creating sustainable competitive advantage through responsible innovation. By establishing clear principles, processes, and accountability structures, organizations can harness AI's transformative power while maintaining the trust of employees, customers, and regulators.

AI governance encompasses the comprehensive frameworks, policies, and practices that guide how organizations develop, deploy, and manage artificial intelligence systems. At its core, it establishes the rules of engagement for AI technologies — ensuring they operate ethically, transparently, and in alignment with both organizational values and regulatory requirements.

Think of AI governance as the operational blueprint that transforms high-level ethical principles into concrete actions. It addresses fundamental questions: Who has access to AI systems? How are decisions documented and explained? What safeguards prevent discriminatory outcomes? These frameworks create accountability structures that ensure human oversight remains central, even as AI systems handle increasingly complex tasks.

The scope of AI governance extends across the entire AI lifecycle, from initial design through deployment and ongoing monitoring. It encompasses several critical dimensions:

For enterprise teams in engineering, customer service, and HR, effective AI governance translates abstract principles into practical guardrails. It determines how a support team's AI assistant handles sensitive customer data, guides an engineering team's use of code generation tools, and shapes how HR departments deploy AI in recruitment processes. All AI hiring models award significantly higher scores to female candidates (0.45 points higher) while penalizing black male candidates (0.30 point penalty), even when qualifications are identical. Without these frameworks, organizations risk creating "black box" systems that operate without transparency or accountability — a recipe for legal liability, ethical breaches, and loss of stakeholder trust.

For enterprise teams in engineering, customer service, and HR, effective AI governance translates abstract principles into practical guardrails. It determines how a support team's AI assistant handles sensitive customer data, guides an engineering team's use of code generation tools, and shapes how HR departments deploy AI in recruitment processes. When developers are allowed to use AI tools, they take 19% longer to complete issues, a significant slowdown that goes against developer beliefs and expert forecasts. Without these frameworks, organizations risk creating "black box" systems that operate without transparency or accountability — a recipe for legal liability, ethical breaches, and loss of stakeholder trust.

Implementing AI governance best practices requires a strategic approach that balances innovation with responsibility. Start by identifying the critical elements of a governance structure. This involves establishing principles that align with your organization’s mission and focus on ethical AI deployment.

Building a robust AI governance framework is essential for guiding ethical AI deployment. This framework should articulate principles that are in harmony with your organization's mission, ensuring AI activities support strategic goals. By establishing these guidelines, you create a cohesive approach to managing the ethical and operational boundaries of AI systems.

Emphasize Information Security: Protecting sensitive information is fundamental to any governance framework. Implement comprehensive data security measures to ensure compliance with applicable laws. There was a 202% increase in phishing emails in the second half of 2024, with 82.6% of phishing emails now using AI technology and 78% of people opening AI-generated phishing emails. By emphasizing information security, you shield your organization from potential legal and reputational challenges.

Promote Clarity and Responsibility: Clarity involves making AI operations transparent and open to stakeholders. Implement structures that clearly outline decision-making processes and assign responsibility. This openness not only builds trust but also aids in meeting compliance obligations.

Emphasize Information Security: Protecting sensitive information is fundamental to any governance framework. Implement comprehensive data security measures to ensure compliance with applicable laws. By emphasizing information security, you shield your organization from potential legal and reputational challenges.

Ensuring compliance with AI governance standards demands meticulous attention to data handling and regulatory adherence. Implement a permissions-aware framework that prioritizes precise data management and access protocols. By defining clear guidelines for data interaction, organizations can safeguard against unauthorized access and maintain data integrity.

Implement Role-Based Access: Assign access rights based on specific user roles, ensuring that data is handled by authorized personnel only. Utilize advanced verification measures to enhance

security and reduce the risk of unauthorized access.

Conduct Regular Evaluations: Systematically assess access permissions to confirm alignment with organizational policies. Implement checks to ensure compliance and address any discrepancies swiftly, thereby strengthening security measures.

Seamless Integration: Align the permissions framework with existing infrastructure, allowing for consistent security application across all systems. This alignment enhances operational coherence and supports comprehensive compliance efforts.

Proactive Monitoring: Keep updated on regulatory developments to ensure ongoing compliance. This involves actively tracking changes from relevant authorities and incorporating them into governance practices. Anticipating regulatory shifts enables timely policy adjustments.

Dynamic Policy Review: Regularly revise governance policies to incorporate new regulatory requirements. This includes updating documentation and training team members on the latest standards, ensuring that practices remain compliant and effective.

Engage Compliance Experts: Work closely with legal and compliance teams to interpret regulatory changes accurately. Their insights ensure that governance practices not only meet but also exceed compliance expectations, safeguarding against potential risks.

By keeping governance frameworks adaptable and aligned with regulatory standards, organizations can leverage AI responsibly while maintaining trust and compliance.

Effective AI risk management involves a proactive approach to addressing potential ethical and security challenges in AI systems. Establishing a structured risk management process allows organizations to navigate these complexities confidently, ensuring systems are robust and reliable.

Map Out Risks: Begin by identifying areas where AI systems may face ethical dilemmas or security threats. Consider data handling, algorithmic fairness, and compliance with industry standards. This comprehensive mapping provides a clear overview of potential vulnerabilities.

Prioritize and Analyze: Assess the severity and frequency of each identified risk. This prioritization helps allocate resources effectively, focusing on the most pressing concerns that require immediate attention.

Craft Strategic Responses: Develop tailored strategies to manage each risk. This includes preventive measures and adaptive responses to minimize potential impacts, ensuring resilience in the face of challenges.

Structured Audits: Implement regular audits to verify that AI operations comply with ethical guidelines and accountability requirements. These audits serve as checkpoints to confirm that systems function as intended.

Dynamic Oversight: Establish mechanisms for continuous oversight, allowing for real-time detection and correction of deviations. This dynamic approach ensures ongoing compliance and ethical integrity.

Collaborative Input: Involve a diverse range of stakeholders in the audit process to gain comprehensive insights. Their perspectives can reveal additional areas for improvement and enhance risk management strategies.

Through a well-defined risk management framework, organizations can uphold ethical standards and fortify AI systems against potential threats, reinforcing stakeholder trust and system reliability.

Ensuring openness and responsibility in AI systems reinforces trust and integrity. Organizations must implement strategies that make AI operations and decisions clear to stakeholders, enhancing both understanding and confidence in these technologies.

Integrated Documentation: Develop a cohesive documentation strategy that captures AI processes and decision criteria. This integration ensures that all relevant information is systematically recorded and easily retrievable.

Transparent Decision Logs: Create detailed logs that track AI decision-making pathways, explaining the logic and data inputs involved. These logs should be concise and provide clarity on how outcomes are reached.

Information Accessibility: Make documentation available through user-friendly platforms, enabling quick access and facilitating informed discussions about AI functionalities.

Facilitated Discussions: Encourage open forums where team members can discuss AI systems and their implications. This fosters a collaborative environment that values honest communication and feedback.

Engaging Stakeholders: Involve diverse stakeholder groups in AI planning and evaluation, drawing on their insights to improve system transparency and effectiveness.

Foundation of Principles: Create a curriculum that emphasizes core AI values such as impartiality, responsibility, and openness. These values should steer all AI-related activities, ensuring that ethical considerations are integral to decision-making. This is increasingly urgent as 68% of companies face a moderate to extreme AI talent shortage, with a projected 700,000 AI job shortage by 2027 despite AI specialists receiving a 56% salary premium.

By adopting these practices, organizations can create an environment where AI transparency and accountability are prioritized, ensuring alignment with ethical standards and organizational goals.

Developing a strong AI ethics framework is crucial for guiding responsible technology use across an organization. This involves not only defining ethical guidelines but also ensuring that all employees understand and adhere to them. By integrating ethics into training programs, organizations foster a culture of responsibility and awareness.

Foundation of Principles: Create a curriculum that emphasizes core AI values such as impartiality, responsibility, and openness. These values should steer all AI-related activities, ensuring that ethical considerations are integral to decision-making.

Real-World Applications: Use examples and scenarios that employees might encounter. This practical approach helps them identify ethical challenges and apply principles effectively in their roles.

Engaging Formats: Use interactive sessions and workshops to actively involve employees. This approach encourages dialogue and reflection, deepening their understanding of AI ethics.

Certification Programs: Launch certification initiatives that confirm employees' grasp of AI ethics and governance. These programs offer recognition for competency and commitment to ethical practices.

Continuous Learning: Implement ongoing educational opportunities to keep employees informed about the latest in AI ethics and governance. Regular updates ensure alignment with evolving standards and practices.

Feedback Channels: Establish methods for employees to share feedback on training programs. This input helps refine and improve the curriculum to better meet organizational needs.

By embedding ethics into training and certification programs, organizations can nurture an environment where ethical AI use is a collective responsibility, enhancing trust and accountability across all levels.

Implementing effective AI governance requires a strategic framework that emphasizes adaptability, stakeholder involvement, and proactive adjustments. This approach ensures that AI initiatives are not only innovative but also aligned with evolving business landscapes.

Visionary Integration: Position AI governance as a core element of your organization's future vision. This ensures that AI projects contribute to long-term goals and are embedded within the company's strategic evolution. By integrating AI into the organizational vision, you create pathways for sustainable advancements.

Outcome Evaluation: Develop robust evaluation methods to measure AI’s impact on achieving organizational milestones. This approach provides valuable insights into how AI initiatives contribute to broader business success. Despite $30 to 40 billion in enterprise GenAI investment, a stunning 95% of organizations are achieving zero measurable return.

Prioritization Strategy: Focus efforts on AI projects that align with critical business opportunities and challenges. This strategic prioritization ensures that resources are directed towards initiatives with the highest potential for meaningful impact.

Unified Collaboration: Encourage collaboration across departments to integrate diverse insights into AI governance. By fostering a culture of shared knowledge, you enhance the effectiveness of governance practices and decision-making.

Role Clarity: Clearly define roles and responsibilities for everyone involved in AI projects. This clarity ensures coordinated efforts and seamless integration of AI into various functions.

Continuous Feedback Loops: Establish ongoing channels for feedback that enable continuous refinement of AI strategies. This dynamic interaction empowers stakeholders to contribute insights, fostering an environment of collective growth and accountability.

By adopting these strategies, organizations can create a dynamic AI governance framework that supports adaptability, innovation, and sustained progress, ensuring a cohesive approach to AI implementation.

As you build your AI governance framework, remember that successful implementation requires both the right principles and the right technology foundation. We understand that navigating AI governance can be complex, which is why we're committed to helping organizations deploy AI responsibly and effectively. Request a demo to explore how Glean and AI can transform your workplace and discover how we can support your journey toward secure, ethical, and impactful AI

adoption.

How the world can build a global AI governance framework

This fragmented development will have a significant effect on trust in AI and on whether it can benefit people around the world.

The world is standing at the dawn of the Intelligent Age. This transformation is as profound as the discovery of electricity or the rise of the internet, with over $15 trillion in new value from artificial intelligence (AI) expected by 2030. But recent analysis from the Continuum Institute, which researches AI governance and innovation, suggests that a significant part of that promise could already be at risk, throttled by governance fragmentation, technical incompatibility and a fundamental deficit of trust.

This means an intelligent world economy is being built on a fractured foundation. Algorithms now travel faster than the agreements that govern them and no shared venue exists to coordinate their norms. The result is a trust gap, a silent drag quietly taxing global growth and confidence.

As AI reshapes the global economy, mistrust has become its greatest tax. A new global architecture based on cooperation could transform AI governance from fragmentation into shared growth. This would enable nations to align on AI sovereignty, interoperability and inclusion to build trust in the Intelligent Age.

The EU codes ethics into law, while the US leans on innovation-first self-governance. China embeds alignment within its national strategy, while much of the Global South still fights for access to compute and inclusion.

The different approaches to AI governance carry merit, but by not working together they fragment and slow its development. Further, early research shows that a 2023 ban of a major AI service in Italy could have led exposed firms to underperform by around 9%, proving that inconsistent governance can negatively affect economic performance.

Without confidence in how systems are trained, tested and validated, innovation stalls, investors hesitate, regulators tighten rules and citizens lose trust. The bottleneck of the Intelligent Age is no longer just compute, it’s also credibility.

Between national sovereignty and global accountability lies an unclaimed space. This is the vacuum in which global AI governance should exist and where a global architecture for cooperation on AI could emerge.

During the 20th century, nations built economic institutions to stabilize currencies and rebuild trust after crises such as wars. Today, the world faces a parallel challenge; but not to stabilize currencies, instead to stabilize trust in AI – the new currency of the Intelligent Age. This global architecture must rest on three pillars:

A standing council for cooperative intelligence that aligns national strategies, private innovation and social safeguards.

Together, these elements form the blueprint for a global architecture for intelligent cooperation — a system designed not for hierarchy or hegemony, but for coexistence, accountability and collective progress.

Click to enlarge How cooperation between nations during the 20th century could inspire global AI governance in the Intelligent Age. Image: Continuum Institute, 2025

A constitutional framework could help to resolve the central paradox of AI governance – how to uphold national sovereignty while ensuring global accountability – by balancing three core imperatives so that technology empowers, rather than divides, people:

Click to enlarge The three pillars of a proposed constitutional framework for global AI governance. Image: Continuum Institute, 2025

Together, these pillars form a governance compass that can enable unity of purpose on foundational AI principles such as safety and human dignity, while respecting diversity in implementation and cultural expression. One example of this approach is the African Union's Continental AI Strategy, which seeks to build sovereign capacity rather than rely solely on foreign technologies.

Governance must also evolve from policy on paper into protocol in practice. A two-layer architecture such as the one below could help support this because it would be light enough for many countries to adopt, but strong enough to help them all align on AI governance:

Click to enlarge A two-layer operating system could help nations manage global AI governance. Image: Continuum Institute, 2025

The first layer, the constitutional core, would establish shared technical standards and governance elements that ensure transparency, safety and accountability across the global system. Then, the local overlay would form a second layer that would allow each jurisdiction to apply context-specific regulations for sectors such as health, labour or defence, while maintaining alignment with the global baseline.

To steward this system, governments, research labs and standards bodies could jointly establish a World Council for Cooperative Intelligence (WCCI). This would be a lean, global institution that ensures intelligent systems remain safe, compatible and credible. It would not create standards from scratch but would harmonize and validate those from leading bodies like the International Organization for Standardization (ISO), focusing specifically on cross-border interoperability and systemic risk.

The goal should not be control, but continuity. This would ensure that intelligent systems remain transparent, compatible and accountable across borders, while fully respecting national security prerogatives.

The lessons of 20th century economic architecture still hold: cooperation can compound prosperity. Good governance is not bureaucracy, it is the engine of sustainable growth. And when compliance becomes confidence, nations can unlock three self-reinforcing flywheels:

Click to enlarge Through AI governance, nations can boost trust in AI so it benefits everyone. Image: Continuum Institute, 2025

In the Intelligent Age, the world is redefining how people create value, exercise leadership and thrive in an era of constant disruption. To support these developments, global leaders must:

Launch a model passport initiative with a coalition of willing nations and companies to operationalize AI transparency and trust

Convene a task force to draft the charter of an organization like the WCCI – a modern institution for a cooperative, intelligent world.

The world needs a new kind of global leadership to steward it into the Intelligent Age, one that fuses governance with innovation, informed by intelligence, driven by purpose and accountable to all stakeholders, including future generations.

Introducing the Databricks AI Governance Framework

Today, we're introducing the Databricks AI Governance Framework (DAGF v1.0), a structured and practical approach to governing AI adoption across the enterprise.

As organizations embrace AI at scale, the need for formal governance grows. Enterprises must align AI development with business goals, meet legal obligations, and account for ethical risks. This framework is designed to support program development, deployment, and continuous improvement.

The DAGF complements the Databricks AI Security Framework, offering a complete view of governance that spans both security and operational integrity.

According to a 2024 global survey of 1,100 technology executives and engineers conducted by Economist Impact, 40% of respondents believed that their organization's AI Governance program was insufficient in ensuring the safety and compliance of their AI assets and use cases. In addition, data privacy and security breaches were the top concern for 53% of enterprise architects, while security and governance are the most challenging aspects of data engineering for engineers.

In addition, according to Gartner, AI trust, risk, and security management is the #1 top strategy trend in 2024 that will factor into business and technology decisions, and by 2026, AI models from organizations that operationalize AI transparency, trust, and security will achieve 50% increase in terms of adoption, business goals, and user acceptance.

While it is evident that the lack of enterprise-level AI governance programs is fast becoming a key blocker to realizing return on value from AI investments and AI adoption as a whole, we realized that there is not a single, comprehensive guidance framework that enterprises can leverage to build effective AI governance programs.

In this framework, we introduce 43 key considerations that are essential for every enterprise to understand (and implement as appropriate) to effectively govern their AI journeys.

These key considerations were then logically grouped across 5 foundational pillars, designed and sequenced to reflect typical enterprise org-structures and personas.

The AI Organization pillar embeds AI governance within the organization's broader governance strategy. It underscores the foundation for an effective AI program through best practices like clearly defined business objectives and integrating the appropriate governance practices that

oversee the organization's people, processes, technology, and data. It explains how organizations can establish the oversight required to achieve their strategic goals while reducing risk.

The Legal and Regulatory Compliance pillar helps organizations align AI initiatives with applicable laws and regulations. It guides managing legal risks, interpreting sector-specific requirements, and adapting compliance strategies in response to evolving regulatory landscapes. The outcome is AI programs are developed and deployed within a robust legal and regulatory framework.

The Ethics, Transparency, and Interpretability pillar supports organizations in building trustworthy and responsible AI systems. It emphasizes adherence to ethical principles such as fairness, accountability, and human oversight while promoting explainability and stakeholder engagement. This pillar provides methods to establish accountability and structure within organizational teams, helping to ensure that AI decisions are interpretable, aligned with evolving ethical standards, and fostering long-term trust and societal acceptance.

The Data, AI Operations (AIOps), and Infrastructure pillar defines the foundation that supports organizations in fully deploying and maintaining AI. It provides guidelines for creating a scalable and reliable AI infrastructure, managing the machine learning lifecycle, and ensuring data quality, security, and compliance. This pillar also emphasizes best practices for AI operations, including model training, evaluation, deployment, and monitoring, so AI systems are reliable, efficient, and aligned with business goals.

The AI Security pillar introduces the Databricks AI Security Framework (DASF), a comprehensive framework for understanding and mitigating security risks across the AI lifecycle. It covers critical areas such as data protection, model management, secure model serving, and the implementation of robust cybersecurity measures to protect AI assets.

For an additional overview of DAGF and for an example walkthrough of how an organization can leverage the framework to create clear ownership and alignment across the AI program lifecycle, please watch this presentation from the authors made during the 2025 Data + AI Summit.